

JProlog Reference Manual

Built-in predicates and operators, with examples

Version 4.6.0 - 24 September 2026

JProlog - a Prolog interpreter in Java

<https://github.com/DenzoSOFTHub/JProlog>

Generated from the source tree and the reference documentation.

Contents

Part I - Introduction	4
About this manual	4
Running JProlog	4
Part II - Language essentials	6
Terms	6
Clauses, programs and queries	6
Unification and control	7
Arithmetic	7
Definite clause grammars	7
Modules and tabling	7
Part III - Operators	8
Default operator table	8
Clause and directive operators	8
Control operators	9
Unification and comparison operators	10
Arithmetic operators (evaluable functors)	11
CLP(FD) operators	12
User-defined operators	12
Part IV - Built-in predicates	13
1. Type Checking Predicates	14
2. Term Manipulation	21
3. List Operations	28
4. Arithmetic Predicates	39
5. Control Flow	44
6. Meta-Predicates	50
7. Input/Output	57
8. Database Operations	70
9. Atom and String Operations	75
10. Character Processing	82
11. DCG (Grammar) Predicates	86
12. Exception Handling	91
13. System Predicates	94
14. Cryptographic Predicates	100
15. JSON Predicates	104
16. Date/Time Predicates	106
17. Filesystem Predicates	109
18. OS Predicates	113
19. Regex Predicates	116
20. XML Predicates	118
21. Threading Predicates	119
22. CSV Predicates	126
23. Logging Predicates	128
24. CLP(FD) Predicates	130
25. Tabling Predicates	139
26. HTTP Predicates	143

27. Persistence Predicates	146
28. Graph Algorithm Predicates	149
29. Java FFI Predicates	152
30. Concurrent Execution Predicates	155
31. Module System	157
32. String predicates	160
33. Stream predicates	163
34. Global variables	166
35. Attributed variables and coroutining	167
36. More list and pair predicates	169
37. Term inspection and modification	171
38. Clause inspection and operators	173
39. Debugging and profiling	175
40. System and memory	176
41. JDBC database access	178
42. TCP, UDP and DNS	180
43. DCG support predicates	182
Appendix A - Prolog flags	183
Appendix B - Engine configuration and embedding	184
Appendix C - ISO error terms	186
Appendix D - Standard order of terms and text conversion summary	187
Appendix E - Libraries not loaded by default	188
Index of predicates and operators	189

Part I - Introduction

About this manual

This manual is the complete reference to the **default** JProlog system: every operator in the initial operator table and every built-in predicate registered when a Prolog engine is created (the CLI, the IDE and the Java embedding API all start from the same registry). Each predicate entry gives its purpose, the argument modes, notes on behaviour that differs from other Prolog systems, and at least one runnable code example. Libraries that ship as source packages but are not loaded by default (expert-system, ASP, CHR, CLP(R), ... extensions) are listed in Appendix E and are not described here.

Conventions used throughout:

- `name/N` - a predicate with `N` arguments; `name//N` - a DCG non-terminal.
- Argument modes: `+Arg` must be instantiated, `-Arg` is an output (normally unbound), `?Arg` may be either, `@Arg` is not further instantiated by the call, `:Goal` is a callable meta-argument.
- Interactive examples show the query after the `?-` prompt followed by the answer as printed by the JProlog CLI: bindings such as `X = 3.`, `true.` for success without bindings, `false.` for failure, and `ERROR: ...` when an exception escapes to the top level. Alternative solutions are separated by `;`.
- Program examples are ordinary clauses you can put in a file and load with `consult/1` or `:consult`.
- Answers are written in the usual Prolog notation, which is what the JProlog console prints since v3.14.0 - except that this manual puts a space after a comma inside a compound term or a list (`[1, 2], f(a, b)`) for readability, where the console writes none (`[1, 2], f(a, b)`): bindings appear in quoted operator notation - `a-1`, `foo/1`, `(p, q)`, the atom `'42'` - and a variable that is still free prints as `_A`, `_B`, ... (a query variable that comes back unbound keeps its own name, so `X = f(Y)` prints as `X = f(Y)`). If an answer still carries constraints, they are listed after the bindings: `freeze(X, Goal), dif(X, a), X in 1..3`.

Running JProlog

JProlog is a Prolog interpreter written in Java (source level 1.8, runs on any modern JDK). It has no runtime dependencies beyond the JDK.

Command-line console - start it with `java -cp target/classes`

`it.denzosoft.jprolog.PrologCLI`, or build the runnable jar with `mvn package -DskipTests` and run `java -jar target/jprolog.jar`. Queries are typed after the `?-` prompt and end with a full stop.

Command line: `jprolog [options] [file.pl ...]` consults the files in order, then:

Option	Effect
<code>-g Goal</code>	run Goal once after loading (repeatable); failure or an uncaught error exits with status 1
<code>-t Goal</code>	run Goal as the toplevel instead of the interactive loop (<code>-t halt</code> ends after <code>-g</code>)
<code>--safe</code>	safe mode: no files, processes, network, threads, CSV files or log file
<code>--budget N</code>	inference budget per query
<code>--max-solutions N</code>	print at most N answers per query
<code>--demo</code>	load the demo facts (<code>father/2</code> , <code>mother/2</code> , <code>parent/2</code> , <code>likes/2</code> , <code>color/1</code>)
<code>--batch / -q</code>	print answers without prompting (<code>-q</code> also drops the banner)
<code>--interactive</code>	prompt for <code>;</code> even when stdin is not a terminal

`halt(N)` ends the process with status `N`, and a file containing `:- initialization(main, main).` runs `main` after loading and then exits (0 on success, 1 on failure or an uncaught error) - a Prolog script. The console also understands these commands:

Command	Effect
:consult <file> (:c)	Load facts and rules from a file
:compile <file> (:cc)	Compile a .pl file to the binary .jpc format
:listing (:l)	Show all loaded clauses
:save <file> (:s)	Save the current knowledge base to a file
:clear	Remove all user clauses
:trace [on off]	Toggle four-port call tracing (same as trace. / notrace.)
:help (:h)	Show the command list
:quit (:q)	Exit

Answers are computed one at a time (v4.5.0): the console prints an answer and, when more may follow, writes ; and waits - type ; and Enter to compute the next one, or Enter alone to stop. An answer that is certainly the last ends in .; a ; that finds nothing more prints false.. **When stdin is not a terminal** - a pipe, a here-document, a redirected file, a CI job - or when the console is started with --batch or -q, it prints the answers as they are found, separated by ; and terminated by ., and reads nothing back (an infinite enumeration therefore prints forever in constant memory; bound it with --max-solutions):

```
$ printf 'between(1, 3, X).\n:quit\n' | java -cp target/classes it.denzosoft.jprolog.Prolog
gCLI
?- X = 1 ;
X = 2 ;
X = 3.
```

IDE - java -cp target/classes it.denzosoft.jprolog.editor.PrologIDE (or mvn exec:java@run-ide) opens the Swing IDE: a project tree, a multi-tab editor with syntax highlighting, completion and a source formatter, a Build tab with line-accurate compile diagnostics, a Run console for queries with streaming/table results, and a four-port debugger with breakpoints, stepping, variable and watch views.

Embedding in Java - the engine is the class it.denzosoft.jprolog.core.engine.Prolog:

```
import it.denzosoft.jprolog.core.engine.Prolog;
import it.denzosoft.jprolog.core.terms.Term;
import java.util.List;
import java.util.Map;

Prolog prolog = new Prolog();
prolog.consult("parent(tom, bob).\nparent(bob, ann).\n"
    + "grandparent(X, Z) :- parent(X, Y), parent(Y, Z).");

List<Map<String, Term>> solutions = prolog.solve("grandparent(tom, who)");
for (Map<String, Term> s : solutions) {
    System.out.println(s.get("Who"));          // ann
}

// Lazy, cancellable enumeration: the sink returns false to stop early
prolog.solveStream("between(1, inf, N)", sol -> ((Number) sol.get("N")).intValue() < 5);

// Hardening for untrusted programs (see Appendix B)
prolog.enableSafeMode();                      // remove OS/file/network/JDBC/FFI/thread built-ins
// or: prolog.enableSafeMode(new SafeModeOptions().allowFileRead("data")); // read-only,
// one tree
prolog.setInferenceBudget(1_000_000);         // abort runaway queries
```

consultWithDiagnostics(source, fileName) compiles clause by clause and returns every syntax error with its line number instead of stopping at the first one; compileFile("prog.pl") produces prog.jpc, which consultSmart("prog.pl") loads when it is newer than the source.

Part II - Language essentials

This part is a compact tour of the Prolog language as implemented by JProlog, enough to read the reference entries. JProlog follows ISO/IEC 13211-1 for syntax, the standard order of terms, arithmetic and error terms, and adds the SWI-Prolog extensions that most programs rely on (strings, global variables, coroutining, CLP(FD), tabling, `format/2`, `must_be/2`, `pairs_*`, and the `apply` and `lists` libraries).

Terms

Every datum is a *term*:

- **Atoms** - symbolic constants: `hello`, `'New York'`, `[]`, `+`. Quote an atom when it does not start with a lowercase letter or contains characters other than letters, digits and `_`. Quoted atoms accept the ISO escapes `\n`, `\t`, `\\`, `\'`, `\xHH` and `\NNN` (octal).
- **Numbers** - integers of unbounded size (`42`, `-7`, `123456789012345678901234567890`, `0xFF`, `0o17`, `0b1010`, `0'a` for a character code) and IEEE double floats (`3.14`, `1.0e10`, `-2.5E-3`). An integer and a float are *different terms* even when numerically equal: `1 == 1.0` fails, `1 == 1.0` succeeds.
- **Strings** - `"text"` is a string object by default (`double_quotes` flag = `string`); set the flag to `codes`, `chars` or `atom` for the traditional readings. Back-quoted text `abc` is always a code list.
- **Variables** - start with an uppercase letter or `_`: `X`, `Result`, `_Tmp`. The anonymous variable `_` is a fresh variable at every occurrence.
- **Compound terms** - a functor with arguments: `point(1, 2)`, `person(Name, Age)`, `'-(a, b)`. Operators are syntactic sugar for compound terms: `1 + 2` is `+(1, 2)`, `X = Y` is `=(X, Y)`.
- **Lists** - `[1, 2, 3]` is sugar for `'.(1, '.(2, '.(3, []))'`; `[H|T]` splits head and tail; `[a, b|T]` is a partial list. Strings written in double quotes are not lists unless the `double_quotes` flag says so.
- **Curly terms** - `{a, b}` is `'{ }'((a, b))`, used by DCG and CLP(FD) syntax.

Comments run from `%` to the end of the line or between `/*` and `*/`.

Clauses, programs and queries

A program is a sequence of clauses terminated by `.` and whitespace:

```
% Facts: unconditionally true
parent(tom, bob).
parent(bob, ann).

% Rules: Head :- Body. The head is true when the body goals succeed in order.
grandparent(GP, GC) :-
    parent(GP, P),
    parent(P, GC).

% Directives are executed while loading
:- dynamic counter/1.
:- initialization(main).
```

Clauses for the same predicate (name/arity) should be contiguous; predicates that are modified at runtime with `assert/retract` are declared dynamic. A query is a goal (or conjunction of goals) evaluated against the program; Prolog searches clauses top to bottom, tries goals left to right, and *backtracks* to the most recent choice point on failure:

```
?- grandparent(tom, Who).
Who = ann.

?- parent(X, Y).
X = tom, Y = bob ;
X = bob, Y = ann.
```

Unification and control

- `=/2` unifies two terms, binding variables as needed; `\=/2` succeeds when they do not unify; `==/2` tests structural identity without binding.
- Conjunction `,`, disjunction `;`, if-then-else `(Cond -> Then ; Else)`, soft-cut `(Cond *-> Then ; Else)`, negation as failure `\+ Goal`, and the cut `!` (commits to the current clause and to the choices made so far in its body).
- `call/1..8` calls a goal built at runtime, adding extra arguments: `call(plus(1), 2, X)` gives `X = 3`. `findall/3`, `bagof/3`, `setof/3`, `forall/2` and `aggregate_all/3` collect or check all solutions of a goal. Lambda expressions of SWI-Prolog's library(`yall`) (`[X]>>Goal`) are **not** supported: pass a named auxiliary predicate to `maplist/2..5`, `include/3`, `foldl/4..6`,
- Exceptions: `throw(Ball)` unwinds to the nearest `catch(Goal, Catcher, Recovery)` whose `Catcher` unifies with the ball. Built-ins raise ISO error terms of the form `error(Formal, Context)`, e.g. `error(type_error(integer, abc), context(atom_length/2, _))` - see Appendix C.

```
max(X, Y, X) :- X >= Y, !.
max(_, Y, Y).

safe_div(X, Y, Z) :-
    catch(Z is X / Y, error(evaluation_error(zero_divisor), _), Z = undefined).

?- safe_div(1, 0, Z).
Z = undefined.
```

Arithmetic

Arithmetic is evaluated only by `is/2` and the comparison predicates `==`, `=\=`, `<`, `>`, `=<`, `>=`. Integers are exact and unbounded; `/` yields an integer when the division is exact and a float otherwise; `//`, `mod`, `rem`, `div` and the bitwise operators require integers. Every evaluable functor is listed in Part III (operators) and Chapter 4 (arithmetic predicates).

```
?- X is 2 ** 10, Y is 2 ^ 100, Z is 7 mod -2, W is max(3, 4.0).
X = 1024.0, Y = 1267650600228229401496703205376, Z = -1, W = 4.0.
```

Definite clause grammars

`Head --> Body` clauses are translated into ordinary predicates with two extra arguments (the difference list being parsed). Terminals are written as lists or strings, `{Goal}` embeds Prolog goals, `!`, `\+`, `->`, `;` and `call//N` are supported, and `phrase/2,3` runs a grammar:

```
greeting --> [hello], name.
name --> [world].
name --> [prolog].

?- phrase(greeting, [hello, prolog]).
true.
```

Modules and tabling

`:- module(Name, [pred/arity, ...])` starts a module; `:- use_module(Name)` imports its exports, and `Module:Goal` calls a predicate in a specific module. `:- table pred/arity` enables SLG tabling (memoisation with left-recursion support) for a predicate - see Chapter 25.

Part III - Operators

Operators let terms be written in infix, prefix or postfix notation. Each operator has a **priority** (1..1200; lower binds tighter) and a **type** that fixes its position and associativity: `xfx` (infix, non-associative), `xfy` (infix, right-associative), `yfx` (infix, left-associative), `fy/fx` (prefix), `yf/xf` (postfix). In a type, `x` marks an argument whose priority must be strictly lower than the operator's, `y` an argument whose priority may be equal. Parentheses always override precedence, and an operator written in canonical form - '+' (1, 2) - is an ordinary compound term. New operators are declared with `op/3`; `current_op/3` enumerates the active table.

Default operator table

Priori ty	Type	Operators	Meaning
1200	xfx	<code>:- --></code>	Clause, DCG rule
1200	fx	<code>:- ?-</code>	Directive, query
1150	fx	<code>dynamic</code> <code>discontiguous</code> <code>multifile</code> <code>table</code> <code>meta_predicate</code> <code>module_transparent</code>	Declarations
1100	xfy	<code>;</code>	Disjunction
1050	xfy	<code>-> *-></code>	If-then, soft-cut
1000	xfy	<code>,</code>	Conjunction
900	fy	<code>\+</code>	Negation as failure
760	yfx	<code>#<==></code>	CLP(FD) reified equivalence
750	xfy / yfx	<code>#==> / #<==</code>	CLP(FD) reified implication
740	yfx	<code>#\ /</code>	CLP(FD) reified disjunction
730	yfx	<code>#\</code>	CLP(FD) reified exclusive or
720	yfx	<code>#/\</code>	CLP(FD) reified conjunction
710	fy	<code>#\</code>	CLP(FD) reified negation
700	xfx	<code>= \= == \== @< @=< @> @>= =.. is =:= =\= < =< > >=</code>	Unification, comparison, arithmetic
700	xfx	<code>in ins #= #\= #< #> #=< #>=</code>	CLP(FD) constraints
600	xfy	<code>:</code>	Module qualification
500	yfx	<code>+ - /\ \ / xor</code>	Additive and bitwise
450	xfx	<code>..</code>	CLP(FD) domain range
400	yfx	<code>* / // rem mod div rdiv << >></code>	Multiplicative and shifts
200	xfx	<code>**</code>	Float exponentiation
200	xfy	<code>^</code>	Integer power / existential quantifier
200	fy	<code>- + \</code>	Unary minus, plus, bitwise complement

The list constructor `|` and the curly braces `{}` are part of the term syntax rather than operators.

Clause and directive operators

`:-/2`, `:-/1`, `?-/1`, `-->/2`

Head `:- Body` is a rule; `:- Goal` is a directive executed at load time; `?- Goal` is an interactive query (also accepted in files); Head `--> Body` is a DCG rule.

```
:- dynamic visited/1.
:- initialization((write('loaded'), nl)).

double(X, Y) :- Y is X * 2.
digits --> [D], { code_type(D, digit) }, digits.
digits --> [].
```

dynamic/1, discontinuous/1, multifile/1, table/1, meta_predicate/1, module_transparent/1

Prefix declaration operators; they take a predicate indicator or a comma/list of indicators.

```
:- dynamic stock/2, price/2.
:- discontinuous rule/2.
:- table fib/2.
:- meta_predicate twice(0).

twice(G) :- call(G), call(G).
```

Control operators

,/2 - conjunction

Both goals must succeed, left to right; bindings made by the left goal are visible to the right one.

```
?- X = 3, Y is X * 2.
X = 3, Y = 6.
```

;/2 - disjunction

Succeeds if either branch succeeds; on backtracking the second branch is tried.

```
?- (X = 1 ; X = 2).
X = 1 ;
X = 2.
```

->/2 and *->/2 - if-then(-else) and soft-cut

(C -> T ; E) runs T after the *first* solution of C, otherwise E; without an else branch it fails when C fails. (C *-> T ; E) runs T for every solution of C and E only if C has none. A cut inside C is local to the condition.

```
sign(X, S) :- ( X > 0 -> S = positive ; X < 0 -> S = negative ; S = zero ).

?- sign(-4, S).
S = negative.

?- (member(X, [1,2]) *-> true ; X = none).
X = 1 ;
X = 2.
```

\+/1 - negation as failure

\+ G succeeds iff G has no solution; it never binds variables.

```
?- \+ member(4, [1,2,3]).
true.
```

!/0 - cut

Discards the choice points created since the clause was entered, including the remaining clauses of the predicate. `call/1`, `\+`, `findall/3` and the condition of `->` are opaque to cut.

```
first_even([X|_], X) :- 0 is X mod 2, !.
first_even([_|T], X) :- first_even(T, X).
```

Unification and comparison operators

=/2 and \=/2

```
?- f(X, b) = f(a, Y).
X = a, Y = b.

?- f(X) \= g(X).
true.
```

==/2, \==/2, @</2, @<=/2, @>/2, @>=/2 - identity and standard order

Structural comparison without unification. The standard order of terms is `Var < Number < Atom < String < Compound`; numbers compare by value (a float precedes an equal integer), atoms and strings alphabetically, compound terms by arity, then name, then arguments.

```
?- X == Y.
false.

?- 1.0 @< 1, a @< b, f(a) @> z.
true.

?- msort([b, 2, a, 1.5, f(x), "s"], L).
L = [1.5, 2, a, b, "s", f(x)].
```

=../2 - univ

Converts between a term and the list `[Functor|Args]`.

```
?- foo(a, b) =.. L.
L = [foo, a, b].

?- T =.. [point, 1, 2].
T = point(1, 2).
```

is/2 and the arithmetic comparisons ==/2, \=/2, </2, <=/2, >/2, >=/2

`X is Expr` evaluates `Expr`; the comparisons evaluate both sides.

```
?- X is (3 + 4) * 2, 10 == 10.0, 2 \= 3, 1 < 2.
X = 14.
```

:/2 - module qualification

Calls a predicate defined in a named module (see *Modules and tabling*). Built-in predicates are global and cannot be qualified with a library name (`Lists:append(...)` fails).

```
:- module(shapes, [area/2]).
area(square(S), A) :- A is S * S.
area(circle(R), A) :- A is pi * R * R.

?- shapes:area(square(3), A).
A = 9.
```

Arithmetic operators (evaluable functors)

All of these are evaluated inside `is/2` and the arithmetic comparisons.

`+/2`, `-/2`, `*/2` - exact on integers, float otherwise

```
?- X is 2 + 3 * 4, Y is 2.5 * 2, Z is 10 - 12.
X = 14, Y = 5.0, Z = -2.
```

`//2`, `///2`, `div/2`, `mod/2`, `rem/2`, `rdiv/2`

`/` is exact division (integer result when it divides evenly, float otherwise); `//` truncates toward zero; `div` floors; `mod` takes the sign of the divisor, `rem` the sign of the dividend; `rdiv` builds a rational number.

```
?- A is 7 / 2, B is 8 / 2, C is -7 // 2, D is -7 div 2, E is -7 mod 2, F is -7 rem 2.
A = 3.5, B = 4, C = -3, D = -4, E = 1, F = -1.
```

`**/2` and `^/2` - exponentiation

`**` always returns a float; `^` is the integer power when both operands are integers (an error for a negative exponent unless the base is 1 or -1) and a float otherwise.

```
?- X is 2 ** 3, Y is 2 ^ 3, Z is 2 ^ 200, W is 2.0 ^ 0.5.
X = 8.0, Y = 8, Z = 1606938044258990275541962092341162602522202993782792835301376, W = 1.4142135623730951.
```

Unary `-/1`, `+/1` and `V/1`

```
?- X is - (3), Y is \ 5, Z is +(7).
X = -3, Y = -6, Z = 7.
```

`/2`, `V/2`, `xor/2`, `<</2`, `>>/2` - bitwise and shifts

```
?- A is 12 /\ 10, B is 12 \/ 10, C is 12 xor 10, D is 1 << 10, E is 1024 >> 3.
A = 8, B = 14, C = 6, D = 1024, E = 128.
```

`^/2` as existential quantifier

Inside `bagof/3` and `setof/3`, `Var^Goal` excludes `Var` from the free variables that split the result into groups. Elsewhere `V^G` simply calls `G`.

```
age(ann, 30). age(bob, 25). age(cid, 30).

?- setof(Name, Age^age(Name, Age), Names).
Names = [ann, bob, cid].

?- bagof(Name, age(Name, Age), Names).
Age = 25, Names = [bob] ;
```

```
Age = 30, Names = [ann, cid].
```

CLP(FD) operators

in/2, ins/2, ..1/2, #=/2, #\=/2, #</2, #>/2, #=</2, #>=/2

Constraint syntax for the finite-domain solver (Chapter 24): `X in 1..9` restricts a domain and the `#`-comparisons post arithmetic constraints that are propagated before and during `label/1`. `ins/2` is present in the operator table for compatibility but is not implemented as a predicate: restrict list elements with `in/2` (or `maplist/2` with an auxiliary predicate).

```
?- X in 1..10, Y in 1..10, X + Y #= 12, X - Y #= 4, label([X, Y]).
X = 8, Y = 4.

?- A in 0..1, B in 0..1, C in 0..1, A + B + C #= 2, A #\= B, label([A, B, C]).
A = 0, B = 1, C = 1 ;
A = 1, B = 0, C = 1.
```

User-defined operators

`op(+Priority, +Type, +Name)` adds (or, with priority 0, removes) an operator; `Name` may be a list. Definitions made in a file apply from that point on and are stored with compiled `.jpc` files. A query is parsed as a whole before it runs, so an operator must be declared in an earlier directive or query before it can be used in operator syntax.

```
:- op(700, xfx, ==>).
:- op(200, xfy, likes).

rule(rain ==> wet).
fact(mary likes wine).

?- rule(A ==> B), X = (john likes (mary likes wine)).
A = rain, B = wet, X = john likes mary likes wine.

?- op(650, xfx, is_bigger), current_op(P, T, is_bigger).
P = 650, T = xfx.
```

Since v3.14.0 there is a single operator store per engine, so `current_op/3` reports operators declared with `op/3` at the top level **and** those declared by `:- op/3` directives while consulting a file. An `op/3` executed while a module file is being consulted is local to that module for `current_op/3`; it is still honoured by the parser everywhere in the session.

Part IV - Built-in predicates

The chapters that follow describe every registered built-in predicate, grouped by function. The first thirteen chapters cover the ISO core and the classic library; the remaining chapters cover the extended libraries that JProlog loads by default (cryptography, JSON, date/time, filesystem, OS, regular expressions, XML, threads, CSV, logging, CLP(FD), tabling, HTTP, persistence, graphs, concurrency, Java FFI, the module system, JDBC, networking, strings, streams, global variables, coroutining and profiling).

1. Type Checking Predicates

Type checking predicates are used to determine the type of a term or to ensure a term is of a specific type. These are essential for writing robust Prolog programs that handle different kinds of data correctly.

Understanding Prolog Types

In Prolog, every piece of data is a **term**. Terms can be:

- **Variables:** Uninstantiated placeholders (e.g., `X`, `_Value`)
- **Atoms:** Constant symbols (e.g., `hello`, `'New York'`)
- **Numbers:** Integers or floats (e.g., `42`, `3.14`)
- **Compound Terms:** Structures with a functor and arguments (e.g., `person(john, 25)`)
- **Lists:** Special compound terms (e.g., `[1, 2, 3]`)

var/1

Purpose: Checks if a term is an uninstantiated variable.

When to use: Use this when you need to handle unbound variables differently from bound values, often in conditional logic or to prevent errors.

```
% Check if a variable is unbound before processing
process_value(X) :-
    var(X),
    !,
    write('Error: Variable must be instantiated'),
    fail.
process_value(X) :-
    write('Processing: '), write(X).

% Example usage:
?- process_value(X).
Error: Variable must be instantiated
false.

?- process_value(42).
Processing: 42
true.
```

nonvar/1

Purpose: Checks if a term is NOT a variable (i.e., it's instantiated).

When to use: Use this to ensure a value has been provided before performing operations on it.

```
% Safe division that checks arguments are instantiated
safe_divide(X, Y, Result) :-
    nonvar(X),
    nonvar(Y),
    Y \= 0,
    Result is X / Y.

% Example usage:
?- safe_divide(10, 2, R).
R = 5.0.

?- safe_divide(X, 2, R).
false. % Fails because X is not instantiated
```

atom/1

Purpose: Checks if a term is an atom (a constant symbol).

When to use: Use when you need to verify that input is a symbolic constant, not a number or structure.

```
% A predicate that processes color names
set_color(Color) :-
    atom(Color),
    member(Color, [red, green, blue, yellow]),
    !,
    write('Color set to: '), write(Color).
set_color(_) :-
    write('Invalid color - must be an atom: red, green, blue, or yellow').

% Example usage:
?- set_color(red).
Color set to: red
true.

?- set_color(123).
Invalid color - must be an atom: red, green, blue, or yellow
true.
```

number/1

Purpose: Checks if a term is a number (integer or float).

When to use: Use before arithmetic operations to ensure the term can be used in calculations.

```
% Calculate the square of a number with type checking
square(X, Result) :-
    number(X),
    Result is X * X.
square(X, _) :-
    \+ number(X),
    write('Error: '), write(X), write(' is not a number'), nl,
    fail.

% Example usage:
?- square(5, R).
R = 25.

?- square(hello, R).
Error: hello is not a number
false.
```

integer/1

Purpose: Checks if a term is specifically an integer (whole number).

When to use: Use when you need whole numbers only, such as for array indices or counting.

```
% Get the Nth element of a list (N must be an integer)
get_nth(N, List, Element) :-
    integer(N),
    N > 0,
    nth1(N, List, Element).
get_nth(N, _, _) :-
    \+ integer(N),
    write('Error: Position must be a positive integer'), nl,
    fail.
```

```
% Example usage:
?- get_nth(2, [a, b, c], E).
E = b.

?- get_nth(2.5, [a, b, c], E).
Error: Position must be a positive integer
false.
```

float/1

Purpose: Checks if a term is a floating-point number.

When to use: Use when you need to distinguish between exact integers and decimal numbers.

```
% Format a number for display based on its type
format_number(N, Formatted) :-
    integer(N),
    !,
    atom_number(Formatted, N).
format_number(N, Formatted) :-
    float(N),
    !,
    format(atom(Formatted), '~2f', [N]).

% Example usage:
?- format_number(42, F).
F = '42'.

?- format_number(3.14159, F).
F = '3.14'.
```

rational/1 and rational/3

Purpose: `rational(@X)` is true when `X` is an integer or a rational number (SWI: rationals include the integers); `rational(@X, -Numerator, -Denominator)` also gives its parts (an integer `N` is `N/1`); both fail on anything else. (v4.6, ISS-2025-0712)

Rationals are a number kind since 4.6 (SWI-Prolog 9 semantics, `prefer_rationals = false`): they come from `rdiv/2`, `rational/1`, `rationalize/1`, arithmetic on rationals and the literal `1r3`, and are always normalised (lowest terms, positive denominator, an integer when the denominator is 1). `number/1` and `atomic/1` are true for them, `integer/1` and `float/1` false. They compare exactly with integers and rationals (`1r3 < 1`, standard order by value); against a float they compare as floats, and the float sorts first on a tie.

```
?- X is 1 rdiv 3, rational(X), \+ float(X).
X = 1r3.

?- rational(3r4, N, D).
N = 3, D = 4.

?- X = 2r4.
X = 1r2.
```

atomic/1

Purpose: Checks if a term is atomic (atom or number).

When to use: Use when you need simple, non-compound values.

```
% Store simple values in a database
store_simple_value(Key, Value) :-
```

```

    atomic(Value),
    !,
    assertz(stored(Key, Value)).
store_simple_value(Key, Value) :-
    write('Cannot store complex term: '), write(Value), nl,
    fail.

% Example usage:
?- store_simple_value(age, 25).
true.

?- store_simple_value(name, 'John').
true.

?- store_simple_value(person, person(john, 25)).
Cannot store complex term: person(john, 25)
false.

```

string/1

Purpose: Checks if a term is a string (a "..." text object, distinct from atoms and character/code lists). (v3.6.0)

When to use: Use to distinguish string terms from atoms when handling text that may arrive in either representation. With the default flag `double_quotes=string`, "abc" denotes a string term.

```

?- string("hello").
true.

?- string(hello).
false.    % atom, not a string

?- string(123).
false.

?- string(X).
false.    % unbound variable

% Dispatch on the text representation
text_length(T, L) :- string(T), !, string_length(T, L).
text_length(T, L) :- atom(T), atom_length(T, L).

```

compound/1

Purpose: Checks if a term is a compound structure (has a functor and arguments).

When to use: Use to identify complex data structures.

```

% Process different types of data differently
process_data(Data) :-
    atomic(Data),
    !,
    write('Simple value: '), write(Data).
process_data(Data) :-
    compound(Data),
    Data =.. [Functor|Args],
    write('Complex structure: '), write(Functor),
    write(' with '), length(Args, N), write(N), write(' arguments').

% Example usage:
?- process_data(42).
Simple value: 42
true.

```

```
?- process_data(person(john, 25, london)).
Complex structure: person with 3 arguments
true.
```

callable/1

Purpose: Checks if a term can be called as a goal (atom or compound).

When to use: Use when you need to verify that a term can be executed as a Prolog goal.

```
% Execute a goal safely
safe_call(Goal) :-
    callable(Goal),
    !,
    catch(call(Goal), Error, (write('Error: '), write(Error), nl, fail)).
safe_call(Goal) :-
    write('Not callable: '), write(Goal), nl,
    fail.

% Example usage:
?- safe_call(append([1], [2], X)).
X = [1, 2].

?- safe_call(123).
Not callable: 123
false.
```

ground/1

Purpose: Checks if a term is fully instantiated (contains no variables).

When to use: Use to ensure all data is complete before processing or storing.

```
% Only store fully instantiated facts
add_fact(Fact) :-
    ground(Fact),
    !,
    assertz(Fact),
    write('Fact added: '), write(Fact).
add_fact(Fact) :-
    write('Cannot add fact with variables: '), write(Fact), nl,
    fail.

% Example usage:
?- add_fact(parent(john, mary)).
Fact added: parent(john, mary)
true.

?- add_fact(parent(X, mary)).
Cannot add fact with variables: parent(_G123, mary)
false.
```

is_list/1

Purpose: Checks if a term is a proper list.

When to use: Use before list operations to ensure the data structure is valid.

```
% Sum all numbers in a list
sum_list(List, Sum) :-
    is_list(List),
    sum_list_helper(List, 0, Sum).
```

```

sum_list(NotList, _) :-
    \+ is_list(NotList),
    write('Error: Expected a list, got: '), write(NotList), nl,
    fail.

sum_list_helper([], Acc, Acc).
sum_list_helper([H|T], Acc, Sum) :-
    number(H),
    NewAcc is Acc + H,
    sum_list_helper(T, NewAcc, Sum).

% Example usage:
?- sum_list([1, 2, 3, 4], Sum).
Sum = 10.

?- sum_list(not_a_list, Sum).
Error: Expected a list, got: not_a_list
false.

```

<!-- START_CHANGE: ISS-2025-0179 - Add acyclic_term/1 and proper_list/1 -->

acyclic_term/1

Purpose: Checks if a term is acyclic (contains no circular references).

When to use: Use to verify term safety before operations that would loop on cyclic terms.

v3.9.0 (ISS-2025-0441): on the opt-in **v4 engine** (-Djprolog.engine=v4) this is a real test over rational trees - $X = f(X)$, `acyclic_term(X)` fails and `cyclic_term(X)` succeeds - because v4 supports cyclic terms instead of raising `representation_error(cyclic_term)`. On the default engine a cyclic term cannot normally be built, so the predicate is effectively always true.

```

% Syntax: acyclic_term(+Term)
?- acyclic_term(f(a, b)).
true.

```

cyclic_term/1

Purpose: Succeeds when Term is a rational (cyclic) tree; the complement of `acyclic_term/1`.

When to use: Guard code that walks a term structurally before handing it to something that cannot represent cycles (a writer, a serialiser, an external API).

Availability: added with engine v4 (v3.9.0). The engines deleted in 4.0.0/4.1.0 could not construct cyclic terms and did not register this predicate.

```

% Syntax: cyclic_term(+Term)
?- X = f(X), cyclic_term(X).
true.

?- cyclic_term(f(a, g(b))).
false.

```

proper_list/1

Purpose: Checks if a term is a proper list (terminates with []).

When to use: Use to distinguish proper lists from partial lists or non-list terms.

```

% Syntax: proper_list(+Term)
?- proper_list([1, 2, 3]).
true.

```

```
?- proper_list([1|2]).  
false.
```

<!-- END_CHANGE: ISS-2025-0179 -->

2. Term Manipulation

Term manipulation predicates allow you to inspect and construct complex terms. These are essential for meta-programming and working with structured data.

functor/3

Purpose: Relates a compound term to its name and arity (number of arguments).

When to use: Use to inspect the structure of compound terms or create new terms with a specific structure.

```
% Example: Extract information about a term's structure
analyze_term(Term) :-
    functor(Term, Name, Arity),
    write('Functor: '), write(Name),
    write(', Arity: '), write(Arity), nl.

% Example usage:
?- analyze_term(person(john, 25, london)).
Functor: person, Arity: 3
true.

% Create a term with specific structure
?- functor(Term, employee, 4).
Term = employee(_G1, _G2, _G3, _G4).

% Practical example: Generic predicate to check term type
is_person_record(Term) :-
    functor(Term, person, 3).

?- is_person_record(person(john, 25, london)).
true.

?- is_person_record(employee(john, 25, london, 50000)).
false.
```

v4.2.0: a **non-ground compound is decomposed** (ISO 8.5.1): `functor(f(X, b), N, A)` gives $N = f$, $A = 2$. It used to raise `instantiation_error`, because the mode was chosen with a groundness test rather than a variable test. Construct mode makes fresh unnamed cells.

arg/3

Purpose: Extracts or checks a specific argument from a compound term.

When to use: Use to access fields in structured data without pattern matching.

v3.6.0: full ISO error clauses (8.5.2.3) - an unbound term raises `instantiation_error`, a non-integer index raises `type_error(integer, N)`, a negative index raises `domain_error(not_less_than_zero, N)`, a non-compound term raises `type_error(compound, T)`; an out-of-range index still just fails. Also works on non-ground compounds: `arg(1, f(X), A)` gives $A = X$ (previously failed).

v4.2.0: an **unbound index enumerates** the arguments (ISO 8.5.2), lazily - one per redo, so `once/1` stops it at the first. It used to raise `instantiation_error`.

```
?- arg(N, f(a, b), X).
N = 1, X = a ;
N = 2, X = b.
```

```
% Example: Database of person records
% person(Name, Age, City)
```

```
% Get the age of a person (2nd argument)
get_age(PersonRecord, Age) :-
    compound(PersonRecord),
    functor(PersonRecord, person, 3),
    arg(2, PersonRecord, Age).

% Example usage:
?- get_age(person(john, 25, london), Age).
Age = 25.

% Update a specific field in a record
update_age(OldPerson, NewAge, NewPerson) :-
    OldPerson = person(Name, _, City),
    NewPerson = person(Name, NewAge, City).

% Or using arg/3 for generic field access:
get_field(Record, FieldNum, Value) :-
    compound(Record),
    arg(FieldNum, Record, Value).

?- get_field(person(john, 25, london), 3, City).
City = london.
```

`=..`/2 (univ)

Purpose: Converts between a compound term and a list (name followed by arguments).

When to use: Use for generic term manipulation, creating terms dynamically, or converting between representations.

ISO §8.5.3 - handles atoms, numbers, and compound terms. Numbers are 0-ary atomic terms: `42 =.. [42]` (v2.8.2). Construction with numeric functor and arity > 0 throws `type_error(atom, _)`.

v3.6.0: full ISO error clauses (8.5.3.3) - both sides unbound raises `instantiation_error`, a non-list (or improper-tail) list side raises `type_error(list, L)`, `X =.. []` raises `domain_error(non_empty_list, [])`, a compound head with extra arguments raises `type_error(atom, Head)`, `X =.. [f(a)]` raises `type_error(atomic, f(a))`. Works on non-ground terms in both directions: `f(Q) =.. L` decomposes and `X =.. [f, Y]` constructs with unbound arguments (both previously raised `instantiation_error`).

```
% Example: Convert term to list and back
?- person(john, 25, london) =.. List.
List = [person, john, 25, london].

?- 42 =.. L.
L = [42].

?- X =.. [42, a, b]. % Error: numeric functor not allowed with args
% throws type_error(atom, 42)

?- Term =.. [student, mary, 20, 'computer science'].
Term = student(mary, 20, 'computer science').

% Practical example: Add one more argument to any term
add_argument(OldTerm, NewArg, NewTerm) :-
    OldTerm =.. [Functor|OldArgs],
    append(OldArgs, [NewArg], NewArgs),
    NewTerm =.. [Functor|NewArgs].

?- add_argument(point(3, 4), 5, Result).
Result = point(3, 4, 5).

% Generic term builder
```

```
build_term(Functor, Args, Term) :-
    Term =.. [Functor|Args].

?- build_term(employee, [john, manager, 75000], T).
T = employee(john, manager, 75000).
```

copy_term/2

Purpose: Creates a copy of a term with fresh (renamed) variables.

When to use: Use when you need an independent copy of a term with its own variables, often in meta-programming or template processing.

```
% Example: Template-based fact generation
person_template(person(Name, Age, _City)) :-
    atom(Name),
    number(Age),
    Age > 0.

create_person_facts :-
    Template = person(_, _, london),
    % Create multiple facts from template
    copy_term(Template, person(john, 25, london)),
    copy_term(Template, person(mary, 30, london)),
    copy_term(Template, person(bob, 35, london)).

% Example of variable independence:
?- Template = foo(X, X), copy_term(Template, Copy).
Template = foo(X, X),
Copy = foo(_G1, _G1). % New variables but same binding pattern

% Practical use: Avoid variable conflicts in meta-predicates
apply_template(Template, Values, Result) :-
    copy_term(Template, Result),
    Result =.. [Functor|Values].
```

Attributed variables and coroutining

v3.11.0, engine v4 (ISS-2025-0457..0462): these predicates are native (put_attr/3, get_attr/3, del_attr/2, attvar/1, term_attvars/2) or prelude Prolog (freeze/2, frozen/2, when/2, dif/2, ?=/2) over a real wake queue - binding an attributed variable runs its suspended goals on the machine's goal stack, so **the bindings a woken goal makes propagate**, it is traced through the four ports, the inference budget and Stop can abort it, and an exception it throws reaches the enclosing catch/3. (The older Java implementations, which lost a when/2-woken goal's bindings - ISS-2025-0336 - belonged to the engine deleted in 4.1.0 and are gone with it.)

One more deliberate behaviour: **a query's variables die with the query**. A goal still suspended when a query ends never fires in a later one (before 4.0.0 it did - session-scoped attributed variables, v2.9.4).

Predicate	Purpose
put_attr(-Var, +Module, +Value)	Attach/replace Module's attribute on an unbound variable. Backtrackable. <code>type_error(variable, T)</code> on a non-variable.
get_attr(+Var, +Module, ?Value)	Unify Value with the attribute; fails when there is none.
del_attr(+Var, +Module)	Remove the attribute; succeeds when there is none. Backtrackable.
attvar(@Term)	True when Term is an unbound variable carrying at least one attribute.
term_attvars(+Term, -AttVars)	The attributed variables of Term, in depth-first order.
copy_term(+Term, -Copy, -Goals)	Copy with the attributes stripped ; Goals are the residual goals that would restore them, expressed over Copy's variables.

<code>unifiable(@X, @Y, -Unifier)</code>	The bindings <code>X = Y</code> would make, as a list of <code>Var = Value</code> , without making them . Fails when the terms do not unify.
<code>freeze(?Var, :Goal)</code>	Run <code>Goal</code> as soon as <code>Var</code> is bound (immediately when it already is). Several frozen goals on one variable aggregate into a conjunction.
<code>frozen(@Var, -Goal)</code>	The goal (or conjunction) delayed on <code>Var</code> , or true.
<code>when(+Condition, :Goal)</code>	Run <code>Goal</code> once <code>Condition</code> holds. Conditions: <code>nonvar/1</code> , <code>ground/1</code> , <code>?=/2</code> , <code>(C1, C2)</code> , <code>(C1 ; C2)</code> . <code>instantiation_error</code> for an unbound condition, <code>domain_error(when_condition, C)</code> for anything else. A disjunctive condition fires the goal exactly once.
<code>dif(@X, @Y)</code>	<code>X</code> and <code>Y</code> can never become identical. Decided immediately when they are identical (fail) or cannot unify (succeed); otherwise it suspends on the variables of the remaining unifier and is re-checked as they are bound.
<code>?=(@X, @Y)</code>	True when <code>X</code> and <code>Y</code> are already identical or already cannot unify - i.e. their (dis)equality is decided.

``Module:attr_unify_hook(AttValue, Other)`` - user-definable in Prolog. The machine calls it, through the normal goal stack, for every attribute of a variable that has just been bound; `AttValue` is that module's attribute and `Other` is what the variable was bound to (possibly another variable). Failing the hook fails the unification. An attribute of a module with no hook is inert data.

```
% A one-module constraint library: the variable may only ever be bound to an even integer.
even:attr_unify_hook(_, Other) :-
    (   var(Other) -> put_attr(Other, even, true)
    ;   integer(Other), 0 is Other mod 2
    ).

even(X) :- put_attr(X, even, true).

?- even(X), X = 4.
X = 4.

?- even(X), X = 5.
false.

% Coroutining
?- freeze(X, format("X became ~w~n", [X])), X = hello.
X became hello
X = hello.

?- when(ground(X-Y), Z is X + Y), X = 1, Y = 2.
X = 1, Y = 2, Z = 3.

?- dif(f(X), f(Y)), X = 1, Y = 1.
false.

?- dif(f(X), f(Y)), X = 1, Y = 2.
X = 1, Y = 2.

% Residual goals of a copy
?- freeze(X, foo(X)), copy_term(X, Y, Goals).
Goals = [freeze(Y, foo(Y))].

?- unifiable(f(X, b), f(a, Y), U).
U = [X = a, Y = b].
```

From Java, `Prolog.residualGoals(solution)` returns the same residual goals for the variables of an answer (`freeze/2`, `when/2`, `dif/2`, `CLP(FD)` `in/2`, `put_attr/3`). Nothing prints them yet - the CLI and the IDE start showing them in wave W7 of the v4 design.

compare/3

Purpose: Three-way comparison of terms using standard ordering.

When to use: Use for sorting, searching, or implementing ordered data structures.

v3.6.0: a pre-bound order argument is validated (ISO 8.4.2.3) - a non-atom raises `type_error(atom, Order)` and an atom other than `<`, `=`, `>` raises `domain_error(order, Order)`; valid pre-bound orders still verify by unification.

```
% Standard term ordering: variables < numbers < atoms < compound terms
% Within each category, there's a specific ordering

% Example: Compare any two terms
?- compare(Order, 3, 5).
Order = (<). % 3 is less than 5

?- compare(Order, apple, banana).
Order = (<). % Alphabetical ordering for atoms

?- compare(Order, foo(1), foo(2)).
Order = (<). % Compound terms compared by arguments

% Practical example: Insert into sorted list
insert_sorted(X, [], [X]).
insert_sorted(X, [H|T], [X,H|T]) :-
    compare(<, X, H), !.
insert_sorted(X, [H|T], [H|Rest]) :-
    insert_sorted(X, T, Rest).

?- insert_sorted(3, [1, 2, 4, 5], Result).
Result = [1, 2, 3, 4, 5].

% Binary search tree insertion
insert_bst(X, nil, tree(X, nil, nil)).
insert_bst(X, tree(V, L, R), tree(V, NewL, R)) :-
    compare(<, X, V),
    insert_bst(X, L, NewL).
insert_bst(X, tree(V, L, R), tree(V, L, NewR)) :-
    compare(>, X, V),
    insert_bst(X, R, NewR).
insert_bst(X, tree(X, L, R), tree(X, L, R)). % Already exists
```

term_variables/2

Purpose: Extracts all variables from a term into a list.

When to use: Use to find all uninstantiated variables in complex terms.

```
% Example: Find all variables in a term
?- term_variables(foo(X, bar(Y, X, 3), Z), Vars).
Vars = [X, Y, Z].

% Practical example: Check if a rule has unbound variables
check_rule_complete(Rule) :-
    term_variables(Rule, Vars),
    ( Vars = []
    -> write('Rule is fully instantiated')
    ; length(Vars, N),
      write('Rule has '), write(N), write(' unbound variables')
    ).

?- check_rule_complete(parent(john, mary)).
Rule is fully instantiated
```

```
true.

?- check_rule_complete(parent(X, mary)).
Rule has 1 unbound variables
true.
```

subsumes_term/2

Purpose: Checks if one term is more general than another.

When to use: Use for pattern matching and generalization checking.

```
% A term T1 subsumes T2 if T1 can be made identical to T2 by binding variables

% Examples:
?- subsumes_term(foo(X, Y), foo(a, b)).
true. % foo(X,Y) is more general than foo(a,b)

?- subsumes_term(foo(X, X), foo(a, b)).
false. % Would need X=a and X=b simultaneously

?- subsumes_term(foo(a, b), foo(X, Y)).
false. % foo(a,b) is more specific, not more general

% Practical example: Find matching templates
find_matching_template(Data, Templates, Match) :-
    member(Match, Templates),
    subsumes_term(Match, Data).

% With templates:
templates([
    person(_, _, _),           % Any person
    person(john, _, _),        % John specifically
    person(_, Age, _) :- Age > 18 % Adults
]).

?- find_matching_template(person(john, 25, london),
    [person(_, _, paris), person(john, _, _)],
    Match).
Match = person(john, _, _).
```

term_to_atom/2

Purpose: Convert between a term and its atom representation. Bidirectional.

v2.8.2: term->atomdirection uses operator-aware formatter for proper roundtrip (1+2 is written as 1+2, not +(1,2)).

v3.6.0: works on non-ground terms - term_to_atom(foo(X, bar), A) formats the variable (A = 'foo(_G1, bar)') instead of failing; a bound atom side keeps parse-and-unify semantics (term_to_atom(foo(Z), 'foo(bar)') binds Z = bar).

v4.4.0 (ISS-2025-0506): both arguments unbound is instantiation_error.

```
?- term_to_atom(f(a, b), X).
X = 'f(a, b)'.

?- term_to_atom(T, 'f(a, b)').
T = f(a, b).

?- term_to_atom(1+2*3, A).
A = '1+2*3'. % operator notation preserved
```

v4.5.0 (ISS-2025-0568, wave P3.4): the atom is parsed by the **v2 reader** with the engine's operator table, so

user operators are visible (`op(700, xfx, ==>)`), `term_to_atom(T, 'a ==> b')`, `writeq/1` output always reads back, and a syntax error `RAISES error(syntax_error(Description), _)` (it used to fail silently). The variables of the parsed term are fresh cells.

term_string/2

Purpose: Convert between a term and its **string** representation. The SWI string twin of `term_to_atom/2`, bidirectional and with the same rules. (*added v4.2.0*)

```
?- term_string(f(x), S).
S = "f(x)".

?- term_string(T, "foo(a, B)").
T = foo(a, B).

?- term_string(1+2*3, S).
S = "1+2*3". % operator notation preserved
```

The string side also accepts an atom, so `term_string(T, 'f(a)')` parses; the term side is written with `quoted(true)`, exactly as `term_to_atom/2` writes it.

atom_to_term/3

Purpose: Parse an atom as a Prolog term, returning the term plus a list of variable bindings (`Name=Var` pairs). (*v2.8.2+*)

v4.4.0 (ISS-2025-0507): a real `error/2` `term - instantiation_error, type_error(atom, A), syntax_error(Message)`. It used to throw the `FORMAL` as a bare atom, so `catch(..., error(type_error(atom, _), _), ...)` never matched.

```
?- atom_to_term('foo(X, Y)', T, B).
T = foo(_42, _43),
B = ['X'=_42, 'Y'=_43].
```

Use case: reading and post-processing user input that may contain variables.

v4.5.0 (ISS-2025-0568): parsed by the `v2` reader; the bindings list names every named variable in order of first occurrence, `_Y`-style ones included (SWI), `_` excluded.

numbervars/3

Purpose: Number unbound variables in a term with `$VAR(N)` terms.

```
?- numbervars(f(X, Y, X), 0, End).
X = '$VAR'(0), Y = '$VAR'(1), End = 2.
```

3. List Operations

Lists are fundamental data structures in Prolog. These predicates provide essential list manipulation capabilities.

Understanding Prolog Lists

A list in Prolog is either:

- Empty: `[]`
- Non-empty: `[Head|Tail]` where Head is an element and Tail is a list

Examples:

- `[1, 2, 3]` is syntactic sugar for `[1|[2|[3|[]]]]`
- `[a, b|Rest]` matches a list starting with a, b, with Rest as the remaining list

append/3

Purpose: Concatenates two lists or splits a list into parts.

When to use: Use for joining lists, finding prefixes/suffixes, or generating list partitions.

v3.10.0, engine v4 (ISS-2025-0453): `member/2`, `memberchk/2`, `append/3`, `select/3`, `selectchk/3`, `nth0/3`, `nth1/3`, `last/2`, `reverse/2`, `length/2`, `msort/2`, `sort/2`, `sum_list/2`, `numlist/3`, `copy_term/2` and `clause/2` are native to the v4 machine, and the nondeterministic ones are **lazy generators**: one alternative per redo, O(1) memory, and the enumeration stops the moment the caller cuts, instead of every solution being materialised before the first is used. Modes, solution order and ISO error terms are unchanged - including two deliberate parity points: `append(X, Y, Z)` with all three arguments open still yields only the single standard solution `X = [], Z = Y`, and `member(X, PartialList)` does not extend the open tail.

```
% Mode 1: Concatenate two lists (inputs: +List1, +List2, output: -List3)
?- append([1, 2], [3, 4], Result).
Result = [1, 2, 3, 4].

% Mode 2: Find all ways to split a list (input: +List3, outputs: -List1, -List2)
?- append(Left, Right, [a, b, c]).
Left = [], Right = [a, b, c] ;
Left = [a], Right = [b, c] ;
Left = [a, b], Right = [c] ;
Left = [a, b, c], Right = [].

% Practical example: Insert element in middle of list
insert_after(Element, After, List, NewList) :-
    append(Prefix, [After|Suffix], List),
    append(Prefix, [After, Element|Suffix], NewList).

?- insert_after(new, b, [a, b, c, d], Result).
Result = [a, b, new, c, d].

% Find if one list is contained in another
sublist(Sub, List) :-
    append(_, Rest, List),
    append(Sub, _, Rest).

?- sublist([b, c], [a, b, c, d]).
true.
```

member/2

Purpose: Checks membership or generates elements of a list.

When to use: Use to test if an element is in a list or to iterate through list elements.

```
% Mode 1: Check if element is in list
?- member(b, [a, b, c]).
true.

?- member(x, [a, b, c]).
false.

% Mode 2: Generate all elements of a list
?- member(X, [red, green, blue]).
X = red ;
X = green ;
X = blue.

% Practical example: Find common elements in two lists
common_elements(List1, List2, Common) :-
    findall(X, (member(X, List1), member(X, List2)), CommonList),
    list_to_set(CommonList, Common). % Remove duplicates

?- common_elements([1, 2, 3, 4], [3, 4, 5, 6], Common).
Common = [3, 4].

% Validate input from a set of options
validate_option(Option, ValidOptions) :-
    member(Option, ValidOptions),
    !.
validate_option(Option, ValidOptions) :-
    write('Invalid option: '), write(Option),
    write('. Valid options are: '), write(ValidOptions), nl,
    fail.
```

length/2

Purpose: Relates a list to its length.

When to use: Use to count elements, create lists of specific length, or constrain list size.

v3.5.0: accepts proper lists containing unbound elements (e.g. `length([A, B, C], N)` gives `N = 3`) - only the list skeleton must be proper.

v3.6.1 (ISS-2025-0425): the **generative** mode works. When the length is unbound and the list's spine ends in an unbound tail, `length/2` enumerates `N = Prefix, Prefix+1, ..` on backtracking instead of failing, so `length(L, N)`, `N >= 3`, `!` gives `N = 3`, `L = [_ , _ , _]` and `length([a|T], N)` enumerates `N = 1, 2, 3, ..` Like SWI-Prolog, an **unguarded** `length(L, N)` with both arguments unbound is therefore a non-terminating generator - bound it with a cut, a comparison, or a known length.

v4.4.0 (ISS-2025-0509): the argument contract raises instead of failing silently - `length(foo, N)` and `length([a|b], N)` are `type_error(list, L)`, `length([a], a)` is `type_error(integer, a)` and `length(L, -1)` is `domain_error(not_less_than_zero, -1)`.

```
% Mode 1: Find length of a list
?- length([a, b, c, d], Len).
Len = 4.

% Mode 2: Create a list of specific length with unbound variables
?- length(List, 3).
List = [_G1, _G2, _G3].

% Mode 3: Check if list has specific length
```

```
?- length([a, b, c], 3).
true.

% Practical example: Pad a list to specific length
pad_list(List, TargetLen, PadValue, PaddedList) :-
    length(List, CurrentLen),
    PadCount is TargetLen - CurrentLen,
    (   PadCount <= 0
    -> PaddedList = List
    ;   length(Padding, PadCount),
        maplist(=(PadValue), Padding),
        append(List, Padding, PaddedList)
    ).

?- pad_list([a, b], 5, x, Result).
Result = [a, b, x, x, x].

% Generate lists of increasing length
generate_lists(Max) :-
    between(0, Max, N),
    length(List, N),
    write('List of length '), write(N), write(': '), write(List), nl,
    fail.
generate_lists(_).
```

reverse/2

Purpose: Reverses the order of elements in a list.

When to use: Use for reversing sequences, implementing stacks, or palindrome checking.

v3.5.0: the inverse mode `reverse(-List, +Reversed)` works, and lists may contain unbound elements (`reverse([A, B], R)` gives `R = [B, A]`).

```
% Basic usage
?- reverse([1, 2, 3, 4], Rev).
Rev = [4, 3, 2, 1].

% Inverse mode (v3.5.0+)
?- reverse(X, [1, 2, 3]).
X = [3, 2, 1].

% Check if list is palindrome
is_palindrome(List) :-
    reverse(List, List).

?- is_palindrome([a, b, c, b, a]).
true.

?- is_palindrome([a, b, c]).
false.

% Practical example: Process list in reverse order without modifying original
process_reverse(List) :-
    reverse(List, Reversed),
    process_items(Reversed).

% Build a string in reverse (useful for efficiency)
build_string_reverse([], Acc, Result) :-
    reverse(Acc, Result).
build_string_reverse([H|T], Acc, Result) :-
    process_char(H, Processed),
    build_string_reverse(T, [Processed|Acc], Result).
```

nth0/3 and nth1/3

v4.5.0 (ISS-2025-0603): on a partial list an integer index EXTENDS the list (`nth0(1, L, x)` gives `L = [_ , x | _]`) and an unbound index enumerates past the end; a negative index fails (SWI `nth0/3`). **Purpose:** Access or find the position of an element in a list (0-based or 1-based indexing).

When to use: Use for indexed access to list elements or finding positions.

```
% nth0/3 - Zero-based indexing (like arrays in most languages)
?- nth0(0, [a, b, c, d], Element).
Element = a.

?- nth0(2, [a, b, c, d], Element).
Element = c.

% nth1/3 - One-based indexing (more natural for humans)
?- nth1(1, [a, b, c, d], Element).
Element = a.

?- nth1(3, [a, b, c, d], Element).
Element = c.

% Find position of element
?- nth0(Position, [apple, banana, cherry], banana).
Position = 1.

?- nth1(Position, [apple, banana, cherry], banana).
Position = 2.

% Practical example: Replace element at position
replace_nth(N, List, NewElement, NewList) :-
    nth0(N, List, OldElement, Rest),
    nth0(N, NewList, NewElement, Rest).

?- replace_nth(1, [a, b, c, d], x, Result).
Result = [a, x, c, d].

% Get multiple elements by indices
get_elements_at([], _, []).
get_elements_at([Index|Indices], List, [Element|Elements]) :-
    nth0(Index, List, Element),
    get_elements_at(Indices, List, Elements).

?- get_elements_at([0, 2, 3], [a, b, c, d, e], Elements).
Elements = [a, c, d].
```

permutation/2

v4.5.0 (ISS-2025-0603): SWI's `select/3`-based clauses (`prelude/lists.pl`): lazy, in SWI's order (`[1, 2, 3]`, `[1, 3, 2]`, `[2, 1, 3]`, ...); either argument may be the partial one. **Purpose:** True when one list is a permutation (reordering) of the other.

When to use: Use to enumerate all orderings of a list or to check that two lists contain the same elements.

```
% Enumerate permutations
?- permutation([a, b, c], P).
P = [a, b, c] ;
P = [a, c, b] ;
P = [b, a, c] ;
P = [b, c, a] ;
P = [c, a, b] ;
P = [c, b, a].
```

```
% Check a specific reordering
?- permutation([1, 2, 3], [3, 1, 2]).
true.

% Inverse mode (v3.5.0+)
?- permutation(P, [1, 2]).
P = [1, 2] ;
P = [2, 1].
```

v3.5.0: the inverse mode `permutation(-List, +Permutation)` works, and lists may contain unbound elements.

select/3

Purpose: Selects an element from a list, returning the element and the rest.

When to use: Use for removing elements, permutations, or non-deterministic selection.

v3.5.0: the insertion mode `select(+Elem, -List, +Rest)` (shown below) works, and lists may contain unbound elements.

```
% Remove an element from a list
?- select(b, [a, b, c, d], Rest).
Rest = [a, c, d].

% Non-deterministic selection - try all possibilities
?- select(X, [1, 2, 3], Rest).
X = 1, Rest = [2, 3] ;
X = 2, Rest = [1, 3] ;
X = 3, Rest = [1, 2].

% Insert an element (using select/3 backwards)
?- select(x, Result, [a, b, c]).
Result = [x, a, b, c] ;
Result = [a, x, b, c] ;
Result = [a, b, x, c] ;
Result = [a, b, c, x].

% Practical example: Remove all occurrences of an element
remove_all(_, [], []).
remove_all(X, List, Result) :-
    select(X, List, Rest),
    !,
    remove_all(X, Rest, Result).
remove_all(_, List, List).

?- remove_all(a, [a, b, a, c, a, d], Result).
Result = [b, c, d].

% Generate permutations
permutation([], []).
permutation(List, [H|Perm]) :-
    select(H, List, Rest),
    permutation(Rest, Perm).

?- permutation([a, b, c], P).
P = [a, b, c] ;
P = [a, c, b] ;
P = [b, a, c] ;
P = [b, c, a] ;
P = [c, a, b] ;
P = [c, b, a].
```

sort/2, sort/4 and msort/2

Purpose: Sort lists in standard order.

- sort/2 removes duplicates
- msort/2 keeps duplicates
- sort/4 (v2.8.1+) accepts key index + order operator: sort(+Key, +Order, +List, -Sorted)

```
% sort/4 with custom order
?- sort(0, @>, [3, 1, 2, 1], L).      % descending, dedup
L = [3, 2, 1].

?- sort(0, @=<=, [3, 1, 2, 1], L).    % ascending, KEEP dups
L = [1, 1, 2, 3].

% sort by 2nd arg of compound
?- sort(2, @<, [pair(a, 30), pair(b, 10), pair(c, 20)], L).
L = [pair(b, 10), pair(c, 20), pair(a, 30)].
```

Key = 0 means compare whole terms; Key = N (N>=1) extracts N-th argument of compound. Order operators: @<, @=<=, @>, @>= (the =</>= variants keep duplicates).

v3.5.0: ISO error handling - an unbound first argument or a partial list (e.g. [a|_]) raises `instantiation_error`, a non-list raises `type_error(list, Culprit)` (previously these failed silently). Lists may contain unbound variables, which sort first in the standard order.

v3.6.0: sort/4 validates keys and options with ISO error terms - with Key > 0 every element must be a compound of arity >= Key (`type_error(compound, Elem)`, `domain_error(argument_index, Key)` otherwise); a bad Key raises `type_error(integer, K)` / `domain_error(not_less_than_zero, K)` and a bad Order raises `type_error(atom, 0)` / `domain_error(order, 0)`. `predsort/3` accepts non-ground lists (`predsort(compare, [X, Y], L)` works), fails when the comparison predicate fails or yields a non-order, and raises `instantiation_error` / `type_error(callable, Pred)` on a bad predicate.

When to use: Use for ordering data, removing duplicates, or preparing data for efficient searching.

```
% sort/2 - Sort and remove duplicates
?- sort([3, 1, 4, 1, 5, 9, 2, 6], Sorted).
Sorted = [1, 2, 3, 4, 5, 6, 9]. % Note: only one 1

% msort/2 - Sort but keep duplicates
?- msort([3, 1, 4, 1, 5, 9, 2, 6], Sorted).
Sorted = [1, 1, 2, 3, 4, 5, 6, 9]. % Both 1s kept

% Sort complex terms (sorted by standard term ordering)
?- sort([person(john, 30), person(alice, 25), person(bob, 25)], Sorted).
Sorted = [person(alice, 25), person(bob, 25), person(john, 30)].

% Practical example: Find unique elements
unique_elements(List, Unique) :-
    sort(List, Unique).

?- unique_elements([a, b, a, c, b, d], Unique).
Unique = [a, b, c, d].

% Count occurrences after sorting
count_occurrences(List, Counts) :-
    msort(List, Sorted),
    count_consecutive(Sorted, Counts).

count_consecutive([], []).
count_consecutive([H|T], [H-Count|Rest]) :-
    count_same(H, T, 1, Count, Remaining),
    count_consecutive(Remaining, Rest).
```

```
count_same(X, [X|T], Acc, Count, Rest) :-
    !,
    Acc1 is Acc + 1,
    count_same(X, T, Acc1, Count, Rest).
count_same(_, List, Count, Count, List).
```

keysort/2

Purpose: Sorts a list of Key-Value pairs by key in standard order, keeping duplicates (stable sort).

When to use: Use to order tagged data by key while preserving the relative order of values with equal keys.

```
?- keysort([b-2, a-1, c-3, a-0], Sorted).
Sorted = [a-1, a-0, b-2, c-3].    % stable: a-1 stays before a-0
```

v3.5.0: accepts non-ground pairs (`keysort([K-V, b-2], S)` works - its primary use case), compares keys only (values are never compared), and raises ISO errors - `instantiation_error` on an unbound argument or partial list, `type_error(list, _)` on a non-list, `type_error(pair, _)` on an element that is not a Key-Value pair.

last/2

v4.5.0 (ISS-2025-0603): on a partial list it ENUMERATES (SWI): `last([a|T], X)` gives `T = []`, `X = a`, then `T = [X]`, `T = [_ , X]`, ... without end - bound the search. **Purpose:** True if Elem is the last element of List.

```
?- last([1, 2, 3], X).
X = 3.
```

flatten/2

v4.5.0 (ISS-2025-0602): iterative (no stack limit on the length or depth); a variable element or tail is kept as an element; a cyclic term is `type_error(acyclic_term, L)`. **Purpose:** Flatten a nested list structure into a single flat list.

```
?- flatten([1, [2, [3, 4]], 5], X).
X = [1, 2, 3, 4, 5].
```

numlist/3

Purpose: Generate a list of consecutive integers from Low to High.

```
?- numlist(1, 5, X).
X = [1, 2, 3, 4, 5].
```

v3.5.0: raises ISO errors on bad bounds - `instantiation_error` if Low or High is unbound, `type_error(integer, Culprit)` if a bound is not an integer (previously failed silently).

sum_list/2, sumlist/2

v4.5.0 (ISS-2025-0603): SWI semantics - the elements are evaluated (`Sum is Sum0 + X`), so a non-number is `type_error(evaluable, F/N)`; integer sums are exact past 2^{63} ; a partial list is `instantiation_error`. **Purpose:** Sum all numeric elements of a list.

```
?- sum_list([1, 2, 3, 4], X).
X = 10.
```

max_list/2, min_list/2

v4.5.0 (ISS-2025-0603): the elements are evaluated ($\max(\text{Max0}, X)$); a non-number is `type_error(evaluable, F/N)`; big integers compare exactly; an empty list fails. **Purpose:** Find the maximum/minimum numeric element in a list.

```
?- max_list([3, 1, 4, 1, 5], X).
X = 5.

?- min_list([3, 1, 4, 1, 5], X).
X = 1.
```

delete/3

Purpose: Remove all occurrences of an element from a list.

```
?- delete([1, 2, 1, 3, 1], 1, X).
X = [2, 3].
```

subtract/3

Purpose: Remove from Set all elements present in Delete.

```
?- subtract([1, 2, 3, 4], [2, 4], X).
X = [1, 3].
```

intersection/3

v4.5.0 (ISS-2025-0603): SWI's definition - duplicates of the first list are KEPT: `intersection([1,1,2], [1,2], X)` gives `X = [1,1,2]`. **Purpose:** Elements present in both sets.

```
?- intersection([1, 2, 3], [2, 3, 4], X).
X = [2, 3].
```

union/3

v4.5.0 (ISS-2025-0603): SWI's definition - duplicates of the first list are KEPT: `union([1,1,2], [2,3], X)` gives `X = [1,1,2,3]`. `subtract/3` likewise. **Purpose:** Set1 plus elements from Set2 not already in Set1.

```
?- union([1, 2, 3], [2, 3, 4], X).
X = [1, 2, 3, 4].
```

append/2

v4.6.0 (ISS-2025-0790, wave Q7): SWI-Prolog's `library(lists)` definition. **Purpose:** `append(+ListOfLists, ?List)` - concatenate a list of lists. `ListOfLists` must be a list (`must_be(list, ...)`); a non-list is `type_error(list, T)`; its elements may be partial, so the predicate also splits.

```
?- append([[a], [b, c], []], L).
L = [a, b, c].

?- append([X, Y], [a]).
X = [], Y = [a] ;
X = [a], Y = [].
```

nextto/3

v4.6.0 (ISS-2025-0790): **Purpose:** `nextto(?X, ?Y, ?List)` - Y directly follows X in List.

```
?- findall(X-Y, nextto(X, Y, [1, 2, 3]), L).
L = [1-2, 2-3].
```

max_member/2, min_member/2, max_member/3, min_member/3

v4.6.0 (ISS-2025-0790): SWI's definitions. **Purpose:** the largest / smallest element by the standard order of terms (/2), or by the order predicate Pred (/3: `max_member(:Pred, -Max, +List)`, `call(Pred, A, B)` means A is "not greater than" B). Unlike `max_list/2` nothing is evaluated; an empty list fails.

```
?- max_member(M, [3, 1, f(x), a]).
M = f(x).

?- min_member(@=<, M, [3, 1, 4]).
M = 1.
```

list_to_set/2

v4.6.0 (ISS-2025-0790): SWI's definition. **Purpose:** `list_to_set(+List, -Set)` removes duplicates, keeping the FIRST occurrence of each element; elements are compared with `==` (so two distinct variables stay). $O(n \log n)$.

```
?- list_to_set([a, b, a, X, c, X, b, Y], S).
S = [a, b, X, c, Y].
```

proper_length/2

v4.6.0 (ISS-2025-0790): **Purpose:** `proper_length(@List, -Length)` - the length of a proper list; fails on a partial or cyclic list (unlike `length/2`, which extends a partial one).

```
?- proper_length([a, b], N).
N = 2.

?- proper_length([a|_], N).
false.
```

library(ordsets): ord_union/2, ord_union/3, ord_subtract/3, ord_intersection/2, ord_intersection/3, ord_memberchk/2, ord_add_element/3, ord_insert/3, ord_del_element/3, ord_selectchk/3, ord_subset/2, ord_intersect/2, ord_disjoint/2, ord_seteq/2, ord_symdiff/3, ord_empty/1, is_ordset/1, list_to_ord_set/2

v4.6.0 (ISS-2025-0790, wave Q7): the prelude module `ordsets` (autoloaded on first use, `use_module(library(ordsets))` accepted), with SWI-Prolog 9's definitions. An ordered set is a list sorted by the standard order of terms without duplicates (`list_to_ord_set/2` is `sort/2`). The merge-based predicates are linear and deterministic. `ord_insert/3` is the DEC-10/YAP name of `ord_add_element/3`; `ord_union/2` and `ord_intersection/2` take a list of sets (`ord_intersection([], [])`).

```
?- ord_union([a, c, e], [b, c, d], U).
U = [a, b, c, d, e].

?- ord_subtract([a, b, c, d], [b, d, e], D).
D = [a, c].
```

```
?- ord_intersection([[a, b, c], [b, c], [c, d]], I).
I = [c].

?- ord_memberchk(c, [a, b, c]), ord_subset([b], [a, b]), \+ ord_intersect([a], [b]).
true.
```

maplist/2, maplist/3, maplist/4, maplist/5

Purpose: Apply a goal to each element of a list (or each tuple across multiple lists). Higher-order predicate.

```
?- maplist(atom, [a, b, c]).
true.

?- maplist(succ, [1, 2, 3], Result).
Result = [2, 3, 4].

?- maplist(plus, [1,2,3], [10,20,30], R).
R = [11, 22, 33].
```

Added in v2.8.1: maplist/5 arity for call(Goal, E1, E2, E3, E4) across 4 lists.

v3.10.0, engine v4 only (-Djprolog.engine=v4, ISS-2025-0454): maplist/2..7 are **Prolog clauses** loaded from the v4 prelude (prelude/apply.pl) instead of a Java built-in. They are linear rather than quadratic with an output list, lazy (re-satisfiable through the mapped goal), traceable through the four ports and interruptible; maplist/6 and maplist/7 exist only there. **A program that defines its own `maplist/3` overrides the library definition** - the prelude is consulted only when the knowledge base has no clause for that indicator.

include/3, exclude/3

Purpose: Filter a list by keeping (include) or removing (exclude) elements where Goal succeeds.

```
?- include(atom, [a, 1, b, 2], X).
X = [a, b].

?- exclude(atom, [a, 1, b, 2], X).
X = [1, 2].
```

partition/4, partition/5

Purpose: Split a list in one pass. partition(:Pred, +List, ?Included, ?Excluded) puts each element where call(Pred, X) succeeds into Included and the rest into Excluded; partition(:Pred, +List, ?Less, ?Equal, ?Greater) uses call(Pred, X, Order) with Order one of <, =, >.

Availability: since v3.10.0 (ISS-2025-0454), as prelude Prolog clauses; before that a call raised existence_error(procedure, partition/4) and programs defined their own. That still works: **a user definition of `partition/4` overrides the library one** (the quicksort in examples/test_16_sorting.pl relies on this, and its partition/4 takes a pivot rather than a goal).

```
?- partition([X]>>(X > 2), [1,2,3,4], Big, Small).
Big = [3, 4], Small = [1, 2].
```

foldl/4, foldl/5, foldl/6, foldl/7

Purpose: Left fold over a list with an accumulator.

```
add(X, Y, Z) :- Z is X + Y.
?- foldl(add, [1, 2, 3], 0, Sum).
Sum = 6.
```

v3.10.0, the default (v4) engine only: `foldl/4..7` are prelude Prolog clauses (see `maplist`), so they fold over four parallel lists and are linear and interruptible.

Lambda expressions - `library(yall)`

Purpose: Write an anonymous predicate inline instead of naming a helper.

Availability: since v3.10.0 (ISS-2025-0455); before that a lambda raised `existence_error(procedure, >>/4)`.

Form	Meaning
<code>Params>>Body</code>	<code>Params</code> is a list of formal parameters, e.g. <code>[X,Y]>>(Y is X*2)</code>
<code>Free/Params>>Body</code>	the variables of <code>Free</code> are shared with the caller instead of renamed apart
<code>\X1^...^Xn^Body</code>	the <code>library(lambda)</code> spelling of the same thing
<code>Free/\X^Body</code>	... with shared free variables

The lambda is copied before **every** call, so one lambda serves every element of a `maplist`. A variable that is already bound outside is copied as its value, so it needs no `/`; use `Free/` only to share an *unbound* variable.

```
?- maplist([X,Y]>>(Y is X*2), [1,2,3], L).
L = [2, 4, 6].

?- foldl([X,A0,A]>>(A is A0+X), [1,2,3,4], 0, S).
S = 10.

?- N = 10, maplist(N/[X,Y]>>(Y is X*N), [1,2], L).
L = [10, 20].
```

4. Arithmetic Predicates

Arithmetic predicates handle mathematical calculations and comparisons. Prolog uses special evaluation rules for arithmetic.

Understanding Arithmetic in Prolog

Prolog treats arithmetic expressions differently from other terms:

- `X = 2 + 3` unifies `X` with the term `+(2, 3)`, NOT with 5
- `X is 2 + 3` evaluates the expression and unifies `X` with 5
- Arithmetic comparisons (`<`, `>`, etc.) automatically evaluate their arguments

is/2

Purpose: Evaluates an arithmetic expression and unifies the result.

When to use: Use for all arithmetic calculations where you need the computed value.

Available evaluable functors

Arithmetic (+, -, *, /, // integer div, mod, rem, div, rdiv (*rational*), ** (ISO float power - ALWAYS float, e.g. `2*3 := 8.0`), `^` integer power (v2.8.1), unary `-`, `+`, `abs`, `sign`, `min`, `max`, `gcd` (v2.8.0*))

Float-only (sqrt, sin, cos, tan, asin, acos, atan, atan2, exp, log, log/2 (v2.8.1), cot (v2.8.1), acot (v2.8.1), cbrt (v2.8.1), sinh, cosh, tanh, asinh, acosh, atanh (*all* v2.8.1))

Float manipulation (truncate, round, floor, ceiling, float, integer (v2.8.1), float_integer_part, float_fractional_part, rational / rationalize (v2.8.1))

Rational (v4.6, ISS-2025-0712: rdiv (exact division), numerator, denominator, rational (the EXACT value of a float), rationalize (the simplest rational that converts back to the same float); addition, subtraction, multiplication, division, min, max and an integer power stay exact on integers and rationals; the flag `prefer_rationals`, false by default, makes a non-exact integer division a rational)

Bitwise integer (/ \, \ /, xor, \ (NOT), <<, >>, msb, lsb, popcount)

Constants (pi, e, inf, nan, epsilon (v2.8.1), max_tagged_integer (v2.8.1), min_tagged_integer (v2.8.1))

v3.5.0: ISO error behavior - a computed float overflow raises `evaluation_error(float_overflow)` and a NaN result raises `evaluation_error(undefined)` (the `inf/nan` constants and their propagation still work); `0 ^ -1` raises `evaluation_error(zero_divisor)`; huge `^<</>>` operands raise a catchable ISO error instead of an unhandled Java exception.

v3.6.0: `float_integer_part` and `float_fractional_part` are correct beyond $\pm 2^{63}$; `float_integer_part(1.0e20)` gives `1.0e20` (results no longer saturate at the 64-bit integer range); truncate-toward-zero semantics for small and negative values are unchanged.

v4.5.0 (ISS-2025-0590..0594, 0613): `integer/1` ROUNDS (half away from zero, like `round/1`: `integer(2.5) := 3`; it truncated); `round/1` is exact (`round(0.49999999999999994) := 0`, `round(-2.5) := -3`); an integer argument of `integer`, `round`, `truncate`, `floor`, `ceiling` is returned unchanged (big integers stay exact); / of two integers past 2^{53} is computed exactly then rounded (`(10^400+1)/10^400 := 1.0`); a negative shift count shifts the other way (`1 << -1 := 0`, `8 >> -2 := 32`, SWI); arbitrarily deep expressions are evaluated without a stack overflow; `[X]` evaluates `X` and a one-character string its code (SWI), a longer list is `type_error(evaluable, '[]'/2)`; errors of `:=`, `<`, ... name that comparison (`':=/2'`), not `is/2`; `cot/1`, `acot/1`, `lsb/1`, `popcount/1` exist (they were documented but missing).

```
% Basic arithmetic
?- X is 2 + 3.
```

```

X = 5.

?- Y is 10 * (3 + 4).
Y = 70.

% Using variables (right side must be instantiated)
?- A = 5, B = 3, C is A * B.
A = 5, B = 3, C = 15.

% Common arithmetic functions
?- X is sqrt(16).
X = 4.0.

?- Y is sin(0).
Y = 0.0.

?- Z is max(5, 3).
Z = 5.

% Integer power
?- X is 2 ^ 10.
X = 1024.

% Truncating cast
?- X is integer(3.7).
X = 3.

?- X is integer(-3.7).
X = -3.

% Hyperbolic
?- X is sinh(0.0).
X = 0.0.

% Base-N logarithm
?- X is log(10, 100).
X = 2.0.

% GCD
?- X is gcd(12, 18).
X = 6.

% Epsilon
?- X is epsilon.
X = 2.220446049250313e-16.

% Practical example: Calculate compound interest
compound_interest(Principal, Rate, Time, Amount) :-
    Amount is Principal * (1 + Rate/100) ** Time.

?- compound_interest(1000, 5, 3, Amount).
Amount = 1157.625.

% Factorial calculation
factorial(0, 1) :- !.
factorial(N, Result) :-
    N > 0,
    N1 is N - 1,
    factorial(N1, R1),
    Result is N * R1.

?- factorial(5, R).
R = 120.

```

Arithmetic Comparison Operators

These operators automatically evaluate arithmetic expressions before comparing.

```
% Less than: <
?- 3 < 5.
true.

?- 2 + 3 < 4 + 2. % Evaluates to: 5 < 6
true.

% Less than or equal: <= (Note: not <= like other languages!)
?- 5 <= 5.
true.

% Greater than: >
?- 10 > 3.
true.

% Greater than or equal: >=
?- 7 >= 7.
true.

% Arithmetic equality: ==
?- 2 + 3 == 5.
true.

?- 2 + 3 == 4 + 1.
true.

% Arithmetic inequality: \=
?- 2 + 3 \= 6.
true.

% Practical example: Temperature converter with validation
celsius_to_fahrenheit(C, F) :-
    C >= -273.15, % Absolute zero check
    F is C * 9/5 + 32.

?- celsius_to_fahrenheit(100, F).
F = 212.0.

?- celsius_to_fahrenheit(-300, F).
false. % Below absolute zero

% Finding values in range
in_range(X, Min, Max) :-
    X >= Min,
    X <= Max.

?- in_range(5, 1, 10).
true.

% Maximum of three numbers
max_of_three(A, B, C, Max) :-
    ( A >= B, A >= C -> Max = A
    ; B >= A, B >= C -> Max = B
    ; Max = C
    ).
```

between/3

Purpose: Generates or tests integers within a range.

v4.5.0 (ISS-2025-0521): every mode is native and lazy - it stops exactly at 9223372036854775807 (it used to wrap), continues into big integers for `inf`, and enumerates big-integer bounds lazily (the old eager fallback materialised the whole range). With the third argument bound it is a plain range test.

When to use: Use for generating sequences, validating ranges, or iteration.

v3.7.0 (ISS-2025-0432): the generative mode is **lazy** on the default engine - one integer per backtrack instead of materialising the whole range up front, so `between(1, 2000000, X), X >= 2000000, !` runs in constant memory. `between(Low, inf, X)` / `between(Low, infinite, X)` genuinely enumerate without an upper bound (they were silently capped at a million solutions).

v4.4.0 (ISS-2025-0509): an unbound bound is `instantiation_error` and a non-integer argument is `type_error(integer, N)`; `between(1, 2, a)` used to fail.

```
% Mode 1: Check if number is in range
?- between(1, 10, 5).
true.

?- between(1, 10, 15).
false.

% Mode 2: Generate all integers in range
?- between(1, 5, X).
X = 1 ;
X = 2 ;
X = 3 ;
X = 4 ;
X = 5.

% Practical example: Generate multiplication table
multiplication_table(N) :-
    between(1, 10, I),
    Result is N * I,
    format('~w x ~w = ~w~n', [N, I, Result]),
    fail.
multiplication_table(_).

?- multiplication_table(7).
7 x 1 = 7
7 x 2 = 14
7 x 3 = 21
...
7 x 10 = 70
true.

% Find perfect squares in range
perfect_squares(Min, Max, Squares) :-
    findall(X, (between(Min, Max, X),
                Sqrt is sqrt(X),
                Sqrt := floor(Sqrt)),
            Squares).

?- perfect_squares(1, 20, S).
S = [1, 4, 9, 16].
```

succ/2

Purpose: Relates consecutive integers (successor relation).

When to use: Use for increment/decrement operations or defining natural number sequences.

v4.4.0 (ISS-2025-0507): ISO error terms - `succ(, )` is `instantiation_error`, `succ(a, X)` is `type_error(integer, a)` and `succ(-1, X)` is `type_error(not_less_than_zero, -1)`. `succ(X, 0)` still fails (0 has no predecessor).

```
% Find successor
?- succ(5, X).
X = 6.

% Find predecessor
?- succ(X, 10).
X = 9.

% Check consecutive numbers
?- succ(7, 8).
true.

% Practical example: Count down
countdown(0) :- write('Blast off!'), nl, !.
countdown(N) :-
    N > 0,
    write(N), nl,
    succ(N1, N),
    countdown(N1).

?- countdown(3).
3
2
1
Blast off!
true.
```

plus/3

v4.5.0 (ISS-2025-0593): integers only (SWI) - `plus(1.5, 2.5, X)` is `type_error(integer, 1.5)`; the arithmetic is exact (big integers). **Purpose:** Relates three integers where the third is the sum of the first two.

When to use: Use for addition with multiple unknown values.

v4.4.0 (ISS-2025-0507): more than one unbound argument is `instantiation_error` and a non-integer bound argument is `type_error(integer, N)`.

```
% Normal addition
?- plus(3, 4, X).
X = 7.

% Find what to add
?- plus(3, X, 10).
X = 7.

% Find first addend
?- plus(X, 5, 12).
X = 7.

% Practical example: Calculate remaining budget
remaining_budget(Total, Spent, Remaining) :-
    plus(Spent, Remaining, Total).

?- remaining_budget(1000, 750, R).
R = 250.

?- remaining_budget(1000, S, 400).
S = 600.
```

5. Control Flow

Control flow predicates determine how Prolog executes programs, including conditionals, loops, and backtracking control.

! (cut)

Purpose: Prevents backtracking past this point.

When to use: Use to commit to a choice, improve efficiency, or implement deterministic predicates.

```
% Example: Deterministic max/3 using cut
max(X, Y, X) :- X >= Y, !.
max(X, Y, Y).

% Without cut, max/3 might give multiple solutions
% With cut, once X >= Y succeeds, we commit to that answer

?- max(5, 3, M).
M = 5. % Only one solution

% Practical example: Grade classification
grade(Score, Grade) :-
    Score >= 90, !,
    Grade = 'A'.
grade(Score, Grade) :-
    Score >= 80, !,
    Grade = 'B'.
grade(Score, Grade) :-
    Score >= 70, !,
    Grade = 'C'.
grade(Score, Grade) :-
    Score >= 60, !,
    Grade = 'D'.
grade(_, 'F').

?- grade(85, G).
G = 'B'. % Only one answer, no backtracking

% Red cut vs Green cut
% Green cut: Doesn't change meaning, just efficiency
% Red cut: Changes program meaning

% Green cut example:
member_check(X, [X|_]) :- !. % Found it, no need to search more
member_check(X, [_|T]) :- member_check(X, T).

% Red cut example (dangerous):
min(X, Y, X) :- X < Y, !.
min(_, Y, Y). % Assumes X >= Y without checking
```

-> (if-then) and ; (else)

Purpose: Implements conditional execution (if-then-else).

v4.6.0 (ISS-2025-0734, wave Q3.5): the bar is no longer read as ;. (a | b) reads as the term '|'(a, b) (infix operator |, priority 1105, type xfy - SWI-Prolog 9); as a GOAL '|'/2 is the same as ;/2 (cut-transparent, -> inside it an if-then-else), and in a DCG body it is alternation. X = (a | b), X =.. L gives L = ['|', a, b]; writeq prints '|'(a,b).

When to use: Use for conditional logic where you need different actions based on conditions.

```
% Basic if-then syntax: (Condition -> ThenPart)
test_sign(X) :-
  ( X > 0
  -> write('Positive')
  ).

% If-then-else syntax: (Condition -> ThenPart ; ElsePart)
sign(X, Sign) :-
  ( X > 0
  -> Sign = positive
  ; X < 0
  -> Sign = negative
  ; Sign = zero
  ).

?- sign(5, S).
S = positive.

?- sign(-3, S).
S = negative.

?- sign(0, S).
S = zero.

% Practical example: Safe division
safe_divide(X, Y, Result) :-
  ( Y == 0
  -> write('Error: Division by zero'), nl,
    fail
  ; Result is X / Y
  ).

?- safe_divide(10, 2, R).
R = 5.0.

?- safe_divide(10, 0, R).
Error: Division by zero
false.

% Nested conditions
classify_age(Age, Category) :-
  ( Age < 0
  -> Category = invalid
  ; Age < 13
  -> Category = child
  ; Age < 20
  -> Category = teenager
  ; Age < 60
  -> Category = adult
  ; Category = senior
  ).
```

\+ (negation as failure)

Purpose: Succeeds if the goal fails (negation by failure).

When to use: Use to test that something is NOT true.

```
% Basic negation
?- \+ member(x, [a, b, c]).
true. % x is not a member

?- \+ member(b, [a, b, c]).
false. % b is a member, so negation fails
```

```
% Practical example: Find elements not in a list
not_in_list([], _).
not_in_list([H|T], Exclude) :-
    \+ member(H, Exclude),
    not_in_list(T, Exclude).

?- not_in_list([a, b, c], [b, d, e]).
false. % b is in the exclude list

?- not_in_list([a, c], [b, d, e]).
true. % neither a nor c is in exclude list

% Check uniqueness
all_different([]).
all_different([H|T]) :-
    \+ member(H, T),
    all_different(T).

?- all_different([a, b, c]).
true.

?- all_different([a, b, a]).
false.

% Important: Negation by failure has limitations
% \+ X = 5 fails if X is unbound (can't prove X is not 5)
% Use with ground terms for predictable behavior
```

once/1

Purpose: Succeeds at most once (finds first solution only).

When to use: Use to make non-deterministic predicates deterministic.

v3.7.0 (ISS-2025-0431): `once/1`, `ignore/1` and `forall/2` run natively on the default engine (as `(G -> true), (G -> true ; true)` and `\+ (C, \+ A)`) and every other meta-call built-in - `aggregate_all/3`, `bagof/3`, `setof/3`, `setup_call_cleanup/3`, `with_output_to/2`, `maplist/2..5`, `foldl/4..6`, `include/3`, `exclude/3`, `partition/4`, `predsort/3` - runs its sub-goal on the iterative machine instead of the (since 4.0.0 deleted) recursive solver. Visible effects: they are one to two orders of magnitude faster, they no longer hit a recursion limit on long lists, and (importantly for embedders) the inference budget set by `Prolog.setInferenceBudget` and the `Stop` interrupt now apply **inside** them. Semantics - `cut` opacity, determinism, ISO error terms - are unchanged.

v3.6.0: a non-callable goal raises `type_error(callable, Goal)` (`once(1)` used to fail silently); an unbound goal raises `instantiation_error`.

```
% Without once/1 - multiple solutions
?- member(X, [a, b, c]).
X = a ;
X = b ;
X = c.

% With once/1 - first solution only
?- once(member(X, [a, b, c])).
X = a.

% Practical example: Find first matching record
find_first_adult(People, Adult) :-
    once((member(person(Name, Age), People), Age >= 18)),
    Adult = Name.

?- find_first_adult([person(tom, 12), person(jane, 20), person(bob, 25)], A).
A = jane. % Only returns first adult, not all
```

```
% Optimization: Prevent unnecessary backtracking
expensive_check(X) :-
    once(complex_computation(X, Result)),
    Result > threshold.

% Ensure deterministic behavior
get_default(Key, Value) :-
    once((lookup(Key, Value) ; Value = default)).
```

repeat/0

Purpose: Always succeeds and provides infinite choice points.

When to use: Use to create loops that retry on failure.

v3.6.1 (ISS-2025-0423): genuinely **infinite** on the default engine. It previously produced exactly 1000 solutions, so a `repeat, ... Done, !` driver loop that needed more than 1000 iterations failed silently. (The 1000-solution bound of the pre-3.6.1 recursive engine is history: that engine was deleted in 4.0.0.)

```
% repeat/0 always succeeds and creates a choice point
% On backtracking, it succeeds again (infinitely)

% Interactive menu example
menu :-
    repeat,
    write('1. Option A'), nl,
    write('2. Option B'), nl,
    write('3. Quit'), nl,
    write('Enter choice: '),
    read(Choice),
    process_choice(Choice),
    Choice = 3, % Exit condition
    !.

process_choice(1) :- write('You chose A'), nl.
process_choice(2) :- write('You chose B'), nl.
process_choice(3) :- write('Goodbye'), nl.
process_choice(_) :- write('Invalid choice'), nl.

% Read until valid input
get_positive_number(N) :-
    repeat,
    write('Enter a positive number: '),
    read(N),
    ( number(N), N > 0
    -> ! % Cut to stop repeating
    ; write('Invalid input, try again'), nl,
      fail % Force backtrack to repeat
    ).
```

forall/2

Purpose: Succeeds if Action succeeds for all solutions of Condition.

When to use: Use to verify universal conditions or perform actions on all solutions.

v3.6.0: a non-callable Condition or Action raises `type_error(callable, G)` (`forall(1, true)` and `forall(true, 1)` used to succeed silently); an unbound goal raises `instantiation_error`.

```
% Syntax: forall(Condition, Action)
% Succeeds if Action succeeds for every solution of Condition

% Check if all elements satisfy a condition
?- forall(member(X, [2, 4, 6, 8]), X mod 2 =:= 0).
```

```

true. % All are even

?- forall(member(X, [2, 4, 5, 8]), X mod 2 == 0).
false. % 5 is not even

% Practical example: Validate all fields in a form
validate_form(Form) :-
    forall(member(field(Name, Value), Form),
            validate_field(Name, Value)).

validate_field(age, Value) :- number(Value), Value > 0, Value < 150.
validate_field(name, Value) :- atom(Value), Value \= ''.
validate_field(email, Value) :- atom(Value), sub_atom(Value, _, _, _, '@').

% Print all solutions
print_all_solutions(Goal) :-
    forall(Goal, (write(Goal), nl)).

?- print_all_solutions(between(1, 5, X)).
between(1, 5, 1)
between(1, 5, 2)
between(1, 5, 3)
between(1, 5, 4)
between(1, 5, 5)
true.

% Check database consistency
check_parent_child_consistency :-
    forall(parent(P, C), child(C, P)).

```

ignore/1

Purpose: Always succeeds, whether the goal succeeds or fails.

When to use: Use for optional operations that shouldn't stop execution if they fail.

v3.6.0: a non-callable goal raises `type_error(callable, Goal)` (`ignore(1)` used to succeed silently); an unbound goal raises `instantiation_error`.

```

% ignore/1 tries to execute the goal but always succeeds

% Example: Optional logging
process_data(Data) :-
    validate(Data),
    ignore(log_to_file(Data)), % Don't fail if logging fails
    compute_result(Data, Result),
    display(Result).

% Optional cleanup
cleanup_resources :-
    ignore(close_file(F1)),
    ignore(close_connection(C1)),
    ignore(free_memory(M1)).

% Try to set optional configuration
initialize :-
    set_required_config,
    ignore(set_optional_feature(feature1)),
    ignore(set_optional_feature(feature2)),
    start_system.

```

setup_call_cleanup/3

Purpose: `setup_call_cleanup(:Setup, :Goal, :Cleanup)` - runs `Setup` once, then `Goal`, and runs `Cleanup` exactly once when `Goal` finishes (all solutions exhausted, failure, or an exception). If `Setup` fails or raises, `Cleanup` is not run.

When to use: Guaranteed resource cleanup (closing files/streams/connections) regardless of how the goal terminates.

v4.4.0 (ISS-2025-0509): `Setup`, `Goal` and `Cleanup` are checked BEFORE `Setup` runs - an unbound one is `instantiation_error`, a non-callable one `type_error(callable, G)`.

v4.4.0 (ISS-2025-0513): **an exception thrown by Cleanup now reaches the enclosing `catch/3`**, whatever `Goal` did. It used to escape to the Java embedder whenever `Goal` had also thrown, because the cleanup ran after the matching catch frame had been consumed. The cleanup's ball **replaces** the goal's (SWI's semantics), so it is matched against the catchers that enclose the `setup_call_cleanup/3`, not against the one that matched the goal's ball:

```
?- catch(call_cleanup(throw(a), throw(b)), E, true).
E = b.

?- catch(catch(call_cleanup(throw(a), throw(b)), a, r1), E2, true).
E2 = b.                % the inner catcher matches only a, so it does not swallow
                        b

?- catch(setup_call_cleanup(true, setup_call_cleanup(true, throw(a), throw(b)), throw(c)),
E, true).
E = c.                % every cleanup runs; the outermost ball survives
```

A cleanup reached by an unwinding ball also runs with the bindings `Goal` made before it threw, so `catch(setup_call_cleanup(true, (X = bound, throw(a)), throw(saw(X))), E, true)` gives `E = saw(bound)`. Resource-limit and cancellation aborts are unaffected: an inference-budget abort or an IDE Stop inside a cleanup is not a `PrologException` and stays uncatchable by `catch/3`.

```
% Cleanup runs whether the goal succeeds, fails, or raises:
read_first_line(File, Line) :-
    setup_call_cleanup(
        open(File, read, S),      % Setup: acquire the stream
        read_line_to_string(S, Line), % Goal
        close(S)).               % Cleanup: always closes S

?- setup_call_cleanup(true, member(X, [1,2]), writeln(done)).
done
X = 1 ;
X = 2.
```

Note: in this engine's eager-solution model, `Cleanup` runs after all of `Goal`'s solutions have been produced (or as soon as `Goal` fails/raises).

call_cleanup/2

Purpose: `call_cleanup(:Goal, :Cleanup)` - equivalent to `setup_call_cleanup(true, Goal, Cleanup)`.

```
?- call_cleanup(member(X, [a,b]), writeln(cleaned)).
cleaned
X = a ;
X = b.
```

6. Meta-Predicates

Meta-predicates operate on other predicates, enabling powerful programming techniques like finding all solutions, applying predicates to collections, and dynamic execution.

call/1

Purpose: Executes a goal constructed at runtime.

v4.5.0 (ISS-2025-0518): the goal is converted to a body **before** any of it runs (ISO 7.6.2): a number anywhere in its `,/;-> /*->/\+` structure raises `type_error(callable, Goal)` for the whole goal, so `call(write(a), 1)` prints nothing and `call((fail, 1))` raises (SWI agrees). A goal written as a variable is `call(G)`, so a `!` it is bound to is local (ISS-2025-0517).

When to use: Use for dynamic predicate execution, higher-order programming, or when the predicate to execute is determined at runtime.

```
% Basic usage: Execute a goal
?- call(append([1], [2], X)).
X = [1, 2].

% Call with a variable goal
?- Goal = member(X, [a, b, c]), call(Goal).
Goal = member(a, [a, b, c]), X = a ;
Goal = member(b, [a, b, c]), X = b ;
Goal = member(c, [a, b, c]), X = c.

% Practical example: Generic filter predicate
filter([], _, []).
filter([H|T], Predicate, [H|Filtered]) :-
    call(Predicate, H),
    !,
    filter(T, Predicate, Filtered).
filter(_|T, Predicate, Filtered) :-
    filter(T, Predicate, Filtered).

% Define some test predicates
positive(X) :- number(X), X > 0.
even(X) :- number(X), 0 is X mod 2.

?- filter([1, -2, 3, -4, 5], positive, Result).
Result = [1, 3, 5].

?- filter([1, 2, 3, 4, 5, 6], even, Result).
Result = [2, 4, 6].

% Higher-order predicate: apply operation to list
map([], _, []).
map([H|T], Operation, [R|Results]) :-
    call(Operation, H, R),
    map(T, Operation, Results).

double(X, Y) :- Y is X * 2.
square(X, Y) :- Y is X * X.

?- map([1, 2, 3], double, Result).
Result = [2, 4, 6].

?- map([1, 2, 3, 4], square, Result).
Result = [1, 4, 9, 16].
```

$\wedge/2$ (existential quantification / goal call)

Purpose: $\wedge$ Goal is callable as an ordinary goal, equivalent to `call(Goal)` - the $\wedge$ prefix is simply ignored outside `bagof/3/setof/3` (SWI/SICStus/YAP-compatible). (v3.6.0)

When to use: Normally written inside `bagof/3/setof/3` to existentially quantify variables; direct calls matter when such a goal is built dynamically or passed around as data and eventually invoked.

```
% As an ordinary goal:  $\wedge$ Goal = call(Goal)
?-  $\wedge$ member(X, [a, b]).
X = a ;
X = b.

% Unchanged inside bagof/setof:  $\wedge$  marks variables to ignore for grouping
?- bagof(Child, Parent $\wedge$ parent(Parent, Child), Children).
Children = [bob, liz, ann, pat, jim].

% The arithmetic  $\wedge$  (integer power) is a separate evaluable functor:
?- X is 2  $\wedge$  3.
X = 8.
```

An unbound `Goal` raises `instantiation_error`; a non-callable one raises `type_error(callable, Goal)`. Like `call/1`, the called goal is opaque to `cut`.

findall/3

Purpose: Collects all solutions to a goal into a list.

When to use: Use when you need all possible solutions collected together.

v3.6.0: the result argument is type-checked (ISO 8.10.1.3) - `findall(X, fail, a)` raises `type_error(list, a)`; variables, partial lists and proper lists remain legal.

```
% Syntax: findall(Template, Goal, List)
% Finds all solutions where Goal succeeds and collects Template values

% Basic example
?- findall(X, member(X, [a, b, c]), List).
List = [a, b, c].

% Collect specific information
?- findall(Name, person(Name, Age, _), Names).
Names = [john, mary, bob]. % Assuming those facts exist

% Practical example: Database queries
% Assume we have facts: employee(Name, Department, Salary)
employee(john, sales, 50000).
employee(mary, it, 60000).
employee(bob, sales, 55000).
employee(alice, it, 65000).

% Find all employees in IT
it_employees(Employees) :-
    findall(Name, employee(Name, it, _), Employees).

?- it_employees(E).
E = [mary, alice].

% Collect complex terms
high_earners(Earners) :-
    findall(
        emp(Name, Salary),
        (employee(Name, _, Salary), Salary > 55000),
        Earners
```

```

    ).

?- high_earners(E).
E = [emp(mary, 60000), emp(alice, 65000)].

% Calculate statistics
average_salary(Avg) :-
    findall(Salary, employee(_, _, Salary), Salaries),
    sum_list(Salaries, Total),
    length(Salaries, Count),
    Avg is Total / Count.

% Important: findall/3 returns [] if no solutions found
?- findall(X, member(X, []), List).
List = []. % Empty list, not failure

```

findall/4

Purpose: `findall(+Template, :Goal, -List, +Tail)` - like `findall/3`, but the collected list ends in `Tail` instead of `[]` (a difference list). (*added v4.2.0*)

When to use: to append the solutions of several goals without a second `append/3` pass.

```

?- findall(X, member(X, [1, 2]), L, t).
L = [1, 2|t].

?- findall(X, member(X, [1, 2]), L, [3]).
L = [1, 2, 3].

?- findall(X, fail, L, t).
L = t.

% Concatenating two collections in one pass
?- findall(X, member(X, [a, b]), L, T), findall(Y, member(Y, [c, d]), T, []).
L = [a, b, c, d].

```

An unbound `Goal` raises `instantiation_error`; a non-callable one raises `type_error(callable, Goal)`.

bagof/3

Purpose: Collects solutions like `findall/3` but respects variable bindings.

v4.5.0 (ISS-2025-0520): solutions are grouped by sorting variant keys of the witness - $O(n \log n)$ in the number of distinct witnesses (40 000 witnesses: about 0.2 s, was over 7 minutes).

When to use: Use when you want solutions grouped by free variables.

v3.5.0: each collected solution is a renamed-apart fresh copy (result lists no longer alias caller variables); an unbound goal raises `instantiation_error` and a non-callable goal raises `type_error(callable, Goal)`.

v3.10.0, engine v4 (ISS-2025-0452): `bagof/3` and `setof/3` are native - the goal runs on the machine and the witness groups are handed out lazily, one per redo. The $\wedge$ handling, the variant-witness grouping and `setof/3`'s standard-order group enumeration are unchanged.

v3.6.0: witness grouping follows ISO 8.10.2.1 - solutions whose witness tuples are variants of each other merge into a single group (e.g. fresh clause variables in the witness no longer split groups), with the member tuples unified against the witness variables on emission.

```

% bagof/3 is like findall/3 but treats free variables differently

% With these facts:

```

```

parent(tom, bob).
parent(tom, liz).
parent(bob, ann).
parent(bob, pat).
parent(pat, jim).

% findall/3 collects all solutions
?- findall(Child, parent(_, Child), Children).
Children = [bob, liz, ann, pat, jim].

% bagof/3 groups by free variables
?- bagof(Child, parent(Parent, Child), Children).
Parent = bob, Children = [ann, pat] ;
Parent = pat, Children = [jim] ;
Parent = tom, Children = [bob, liz].

% Use ^ to existentially quantify variables (ignore them)
?- bagof(Child, Parent^parent(Parent, Child), Children).
Children = [bob, liz, ann, pat, jim]. % Like findall now

% Practical example: Group products by category
product(electronics, laptop, 1000).
product(electronics, phone, 500).
product(clothing, shirt, 30).
product(clothing, pants, 50).
product(food, apple, 2).
product(food, bread, 3).

products_by_category(Category, Products) :-
    bagof(Product, product(Category, Product, _), Products).

?- products_by_category(Cat, Prods).
Cat = electronics, Prods = [laptop, phone] ;
Cat = clothing, Prods = [shirt, pants] ;
Cat = food, Prods = [apple, bread].

```

setof/3

Purpose: Like bagof/3 but removes duplicates and sorts results.

v4.5.0 (ISS-2025-0519): each group is sorted **after** its witnesses are unified (ISO 8.10.3.4), so setof(*X*, member(*X*, [*Y*, *Y*]), *L*) gives *L* = [*Y*].

When to use: Use when you want unique, sorted solutions.

v3.5.0: each collected solution is a renamed-apart fresh copy (result lists no longer alias caller variables); an unbound goal raises `instantiation_error` and a non-callable goal raises `type_error(callable, Goal)`.

v3.6.0: witness groups are enumerated in the standard order of terms (not in textual order) - setof(*X*, member(*X*-*Y*, [*a*-10, *b*-2]), *L*) yields *Y* = 2, *L* = [*b*] before *Y* = 10, *L* = [*a*] - and variant witnesses merge into a single group (ISO 8.10.2.1, also in bagof/3).

```

% setof/3 = bagof/3 + sort + remove duplicates

% With duplicate data:
likes(mary, food).
likes(mary, wine).
likes(john, wine).
likes(john, wine). % Duplicate
likes(bob, food).

% bagof/3 keeps duplicates
?- bagof(X, likes(john, X), Things).
Things = [wine, wine].

```

```
% setof/3 removes duplicates and sorts
?- setof(X, likes(john, X), Things).
Things = [wine].

% Collect all unique items liked
?- setof(Item, Person^likes(Person, Item), AllItems).
AllItems = [food, wine]. % Sorted and unique

% Practical example: Find unique skills
has_skill(john, python).
has_skill(john, java).
has_skill(mary, python).
has_skill(mary, python). % Duplicate
has_skill(bob, javascript).
has_skill(bob, java).

unique_skills(Skills) :-
    setof(Skill, Person^has_skill(Person, Skill), Skills).

?- unique_skills(S).
S = [java, javascript, python]. % Alphabetically sorted

% Find people with common skills
people_with_skill(Skill, People) :-
    setof(Person, has_skill(Person, Skill), People).

?- people_with_skill(python, P).
P = [john, mary]. % Sorted list of people
```

aggregate_all/3

Purpose: Aggregates values from all solutions to a goal using a specified aggregate template.

When to use: Use when you need to compute aggregate statistics (sum, count, max, min, bag, set) over all solutions to a goal in a single call.

v3.5.0: ISO error balls raised by Goal propagate unchanged (no longer wrapped or swallowed); an unbound or non-callable goal raises `instantiation_error` / `type_error(callable, _)`.

v3.10.0, *engine v4* (ISS-2025-0452): native over the machine's `findall`. New template forms `max(Value-Witness)` and `min(Value-Witness)`, which compare the numeric left-hand side and answer with the winning pair: `aggregate_all(max(V-W), member(V-W, [1-a, 3-b]), 3-b)`.

v4.5.0 (ISS-2025-0522): the SWI-Prolog library(`aggregate`) contract. Specs: `count`, `count(T)`, `sum(Expr)`, `max(Expr)`, `min(Expr)`, `max(Expr, Witness)`, `min(Expr, Witness)`, `bag(T)`, `set(T)`. `Expr` is **evaluated** (`sum(X*2)`, `max(X+1)`), so a non-number is `type_error(evaluatable, Name/0)`; `max(Expr, Witness)` answers `max(Value, Witness)` (`aggregate_all(max(X, W), member(X-W, [1-a, 3-b]), max(3, b))`) - the v3.10 `max(V-W)` pair form is gone (`V-W` is now an expression). `count`, `sum`, `max` and `min` accumulate while the goal runs (constant memory). An unbound spec raises `instantiation_error`, an unknown one `domain_error(aggregate_spec, Spec)`. Goal's control structure is checked before it runs (ISS-2025-0518).

v3.6.0: `max(Expr)/min(Expr)` fail when the goal has no solutions and raise `type_error(number, T)` on a non-numeric solution (previously skipped silently); `sum(Expr)` accumulates integers exactly (BigInteger - no 64-bit overflow), with float contagion producing a genuine float (`sum over [1.5, 2.5]` is `4.0`); the empty sum is the integer `0`.

```
% Syntax: aggregate_all(Template, Goal, Result)
% Template can be: count, sum(Expr), max(Expr), min(Expr), bag(Expr), set(Expr)

% Count solutions
?- aggregate_all(count, member(_, [a, b, c]), Count).
```

```

Count = 3.

% Sum values
employee(john, 50000).
employee(mary, 60000).
employee(bob, 55000).

?- aggregate_all(sum(S), employee(_, S), Total).
Total = 165000.

% Collect into a sorted set
?- aggregate_all(set(X), member(X, [b, a, c, a, b]), Sorted).
Sorted = [a, b, c].

% Find maximum
?- aggregate_all(max(S), employee(_, S), Highest).
Highest = 60000.

% Practical example: Generate summary statistics
salary_report(Count, Total, Max, Min) :-
    aggregate_all(count, employee(_, _), Count),
    aggregate_all(sum(S), employee(_, S), Total),
    aggregate_all(max(S), employee(_, S), Max),
    aggregate_all(min(S), employee(_, S), Min).

?- salary_report(Count, Total, Max, Min).
Count = 3, Total = 165000, Max = 60000, Min = 50000.

```

aggregate/3, aggregate/4 and aggregate_all/4

Purpose: library(aggregate)'s bagof-style forms. (v4.6, ISS-2025-0716) `aggregate(+Template, :Goal, -Result)` is bagof/3 followed by the aggregation: it GROUPS by the free variables of Goal (use $V^{\wedge}Goal$ to quantify one away) and fails when Goal has no solution. `aggregate(+Template, +Discriminator, :Goal, -Result)` is the setof/3 form: only solutions with distinct Discriminator-Template values count. `aggregate_all(+Template, +Discriminator, :Goal, -Result)` is the findall/3 form of the latter (no grouping, an empty count is 0). Templates are those of `aggregate_all/3` - count, count(T), sum(E), max(E), min(E), max(E, W), min(E, W), bag(T), set(T) - or a compound of them (`r(count, sum(X))` answers `r(C, S)`).

```

?- aggregate(count, member(X, [a,b]), C).
X = a, C = 1 ;
X = b, C = 1.

?- aggregate(count, X^member(X, [a,b]), C).
C = 2.

?- aggregate(sum(V), member(K-V, [a-1, b-2, a-3]), S).
K = a, S = 4 ;
K = b, S = 2.

?- aggregate(r(count, max(X)), member(X, [3,1,2]), R).
R = r(3, 3).

?- aggregate_all(count, P, member(P-, [a-1, a-2, b-3]), C).
C = 2.

```

limit/2, offset/2, call_nth/2, distinct/1, distinct/2, order_by/2

Purpose: SWI-Prolog's library(solution_sequences). (v4.6, ISS-2025-0710)

- `limit(+Count, :Goal)` - at most Count solutions. Lazy: at the Count-th solution the goal's remaining choice points are cut (its `setup_call_cleanup/3` cleanups run then), so an infinite generator is fine.

- `limit(0, G)` fails, `limit(infinite, G)` is `call(G)`.
- `offset(+Count, :Goal)` - skip the first `Count` solutions (a negative `Count` is `domain_error(not_less_than_zero, C)`).
- `call_nth(:Goal, ?Nth)` - with `Nth` unbound, number the solutions 1, 2, ...; with `Nth` bound, only the `Nth` solution, and the goal is cut after it (`call_nth(G, 0)` fails).
- `distinct(:Goal)`, `distinct(?Witness, :Goal)` - only the first solution for each distinct `Witness` (Goal itself for `/1`); witnesses are compared as `VARIANTS`, and the goal stays lazy.
- `order_by(+Specs, :Goal)` - the solutions sorted by the `asc(Key)/desc(Key)` list (the first spec is the primary key; stable, duplicates kept). Necessarily eager.

The goal is called as `call/1` (a cut in it is local) in the caller's module. The counters are not undone on backtracking (as SWI's `nb_setarg/3` state).

```
?- findall(X, limit(3, between(1, inf, X)), L).
L = [1, 2, 3].

?- findall(X, limit(2, offset(1, member(X, [a,b,c,d]))), L).
L = [b, c].

?- call_nth(member(X, [a,b,c]), N).
X = a, N = 1 ;
X = b, N = 2 ;
X = c, N = 3.

?- findall(X, distinct(X, member(X, [a,b,a,c,b])), L).
L = [a, b, c].

?- findall(N-A, order_by([asc(A), desc(N)], member(N-A, [bob-30, al-25, cy-30])), L).
L = [al-25, cy-30, bob-30].
```

7. Input/Output

I/O predicates handle reading from and writing to files and streams.

Basic Output

write/1-2 and writeln/1-2

Purpose: Output terms to the current output stream.

When to use: Use for displaying results, debugging, or user interaction.

v2.8.2: Output is now **operator-aware** - consults the operator table for infix/prefix/postfix notation, list notation, curly braces, and '\$VAR'(N) rendering when `numbervars(true)`. Output of `write(1+2)` is `1+2` (was `+(1,2)`); lists print as `[a,b,c]`; precedence-aware parens added when needed.

v3.5.0: stream-argument forms `write(Stream, Term)` and `writeln(Stream, Term)` added; `write/1` and `writeln/1` imply `numbervars(true)`, so '\$VAR'(0) prints as A (ISO 8.14.2); floats print in ISO syntax - lowercase exponent (`1.0e10`) and `inf/-inf/nan` spellings - so output re-reads as the same term.

```
% write/1 - Output without newline
?- write('Hello'), write(' '), write('World').
Hello World
true.

% writeln/1 - Output with newline
?- writeln('Hello World').
Hello World
true.

% Writing different types of terms
?- write(42), nl, write([a, b, c]), nl, write(person(john, 25)).
42
[a, b, c]
person(john, 25)
true.

% Practical example: Formatted output
display_person(person(Name, Age, City)) :-
    write('Name: '), writeln(Name),
    write('Age: '), writeln(Age),
    write('City: '), writeln(City).

?- display_person(person(john, 25, london)).
Name: john
Age: 25
City: london
true.

% Building formatted messages
error_message(Code, Message) :-
    write('ERROR '), write(Code), write(': '), writeln(Message).

?- error_message(404, 'File not found').
ERROR 404: File not found
true.
```

nl/0-1

Purpose: Outputs a newline character.

When to use: Use to control line breaks in output.

v3.5.0: the stream form `nl(Stream)` writes the newline to the given stream.

```
% Basic usage
?- write('Line 1'), nl, write('Line 2'), nl.
Line 1
Line 2
true.

% Creating formatted reports
print_header :-
    writeln('=' * 40), % Note: This would need special handling
    writeln('      REPORT TITLE'),
    writeln('=' * 40),
    nl.

% Spacing output
print_list([]).
print_list([H|T]) :-
    write(' - '), writeln(H),
    print_list(T).

?- print_list([apple, banana, cherry]).
- apple
- banana
- cherry
true.
```

print/1-2

v4.5.0 (ISS-2025-0600): `print/1,2` is `portray + writeq + numbervars(true)` (SWI): `print('A b')` prints 'A b'. **Purpose:** Output a term with `write` semantics (unquoted, operators honoured) and `numbervars(true)`. (added v3.5.0)

When to use: Use as the conventional "user-friendly output" predicate; equivalent to `write_term(Term, [numbervars(true)])`. The `portray/1` hook is not supported.

```
?- print(hello), nl.
hello
true.

?- print('$VAR'(0)), nl.
A
true.

% print/2 writes to a specific stream
?- open('out.txt', write, S), print(S, foo(1)), close(S).
true.
```

writeq/1-2

Purpose: Output a term with quoting - atoms that need quotes are quoted so the output can be read back.

```
?- writeq('hello world').
'hello world'
true.

?- writeq(Stream, Term). % stream form
```

v4.2.0: `writeq/2` resolves its stream exactly like `write/2` - it is captured by `with_output_to/2` and by the IDE console (it used to write straight to the process stdout through a static stream map), and it reports the ISO stream errors (`instantiation_error`, `domain_error(stream_or_alias, S)`, `existence_error(stream, S)`) instead of a bare evaluation error.

v3.5.0: quoting/spacing fixes - `writeln(-1)` prints - 1, which re-reads as the same compound (it used to print -1, a number); `' , ' , ' . '` and comment-opening symbolic atoms such as `'/*'` are quoted; `numbervars(true)` is implied (`'$VAR'(0)` prints as A); floats print in ISO syntax (`1.0e10`, `inf`, `-inf`, `nan`).

v4.5.0 (ISS-2025-0562/0563/0564, wave P3.9): the output always reads back as the same term - `-(1)`, `-(2^2)`, `(-)-(-)`, `1-(-)`, `'[]'(a,b)`, `'_x'`, `'\177\'`; floats in SWI's shortest layout (`10000000.0`, `1.0e15`); `1.0Inf/-1.0Inf/1.5NaN` for the special floats. See "Reading and writing operator terms" in `BUILTIN_OPERATORS_REFERENCE.md`. The escape character is written `'\e'` (ISS-2025-0673; it was the octal `'\33\'`, which still reads back).

format/2 and format/3

v4.5.0 (ISS-2025-0595): rewritten after SWI-Prolog. Integers are exact for `~d`, `~D`, `~r`, `~R`, `~f`; `~e/~f/~g` follow C `printf` on the exact binary value (`~g` of 0.1 is 0.1); column stops take a fill character (`~-t~30|`, `~48t` for 0), `~+` defaults to 8 columns; `~Nn` prints N newlines; `~Nc`, `~*c`, `~i`, `~W` (term + write options), `~k` (`write_canonical`), `~p` (`print = portray + writeln`) are complete; `~@` runs its goal once. Every argument fault (`~a` of a compound, `~d` of a non-integer, `~c` of a non-code, `~e/~f/~g` of a non-number, `~s` of a non-text, `~r` without a radix), a missing and a SURPLUS argument raise `error(format(Message),_)`. `format/3` sinks: `atom(A)`, `string(S)`, `codes(C)`, `codes(C, Tail)`, `chars(C)`, `chars(C, Tail)`, or a stream. `~Nw/~Nq` right-align in N columns (JProlog extension).

v4.6 (ISS-2025-0714): column stops (`~-t~20|`, `~+`) count from the column the output STREAM is at (SWI's line position), not from the start of the call: `write(abc)`, `format("~-t~w~10|", [x])` writes `abc x`. The console, files and `with_output_to/2` track their column (`line_position(user_output, P)` is now exact); a memory sink (`atom(A)`, ...) starts at 0. **Purpose:** Formatted output driven by a directive string (`~w`, `~a`, `~d`, `~q`, `~n`, `~2f`, ...), similar to C's `printf`.

When to use: Use for readable, formatted output instead of chains of `write/1` calls.

```
?- format("Hello ~w, you are ~w years old~n", [john, 25]).
Hello john, you are 25 years old
true.

% format/3 with an output stream
?- open('out.txt', write, S), format(S, "~w~n", [data]), close(S).
true.

% format/3 with a capture sink
?- format(atom(A), '~2f', [3.14159]).
A = '3.14'.
```

v3.5.0: `format/2,3` now succeed as goals (output used to be produced with the goal then failing, killing any conjunction containing it); the format string may be an atom, a double-quoted string (the spelling produced by the default `double_quotes=string` flag), or a code/char list; `format/3` honours its first argument - a stream alias/handle, or a capture sink `atom(A)` / `string(S)` / `codes(C)` / `chars(C)`. `~w/~q` imply `numbervars(true)`.

v3.10.0 (ISS-2025-0452): new directive `~@` - the argument is a goal; it is called and everything it writes is spliced in at that point. `~p` (`portray`) likewise calls `portray/1`. Both sub-goals run on the resolution engine, so on the v4 engine they honour the inference budget and the Stop interrupt.

```
?- format("[~@]~n", [write(inner)]).
[inner]
true.
```

v3.6.0: argument mismatches raise errors instead of being papered over - too few arguments for the directives raise a format error, `~d` with a non-integer raises `type_error(integer, Arg)`, and an unknown directive raises an error (previously echoed literally). The atom `[]` in the argument position is the

empty argument list ([[]] passes the atom [] as a single argument; a non-list term still counts as one argument, SWI-style).

Basic Input

read/1, read/2

Purpose: Reads a Prolog term from input (must end with period).

- `read/1` reads from `current_input` (stdin by default)
- `read/2` reads from given stream: `read(+Stream, -Term)` (*added v2.8.0*)

```
% From file stream
?- open('data.pl', read, S), read(S, T), close(S).
T = foo(1, 2).
```

On end-of-file, both bind the term to the atom `end_of_file`.

v3.6.0: input is consumed up to the ISO **end token** (. followed by layout) rather than one line at a time - terms spanning several lines, several terms on one line, leading % and /* */ comments, and dots inside quotes/escapes/floats/graphic tokens are all handled; the stream position is preserved between calls, so the next read resumes right after the end token. Resolves LIM-029.

When to use: Use for reading structured Prolog data.

```
% Interactive reading
get_user_fact :-
    write('Enter a fact (end with .): '),
    read(Fact),
    assertz(Fact),
    write('Fact added: '), writeln(Fact).

% Example interaction:
?- get_user_fact.
Enter a fact (end with .): likes(john, pizza).
Fact added: likes(john, pizza)
true.

% Read and process commands
command_loop :-
    repeat,
    write('Command> '),
    read(Cmd),
    process_command(Cmd),
    Cmd = quit,
    !.

process_command(quit) :- writeln('Goodbye!').
process_command(help) :- writeln('Available commands: help, list, quit').
process_command(list) :- listing.
process_command(_) :- writeln('Unknown command. Type help for assistance.').
```

v4.5.0 (ISS-2025-0566/0567, wave P3.4): the reader is the **v2 parser** with the engine's operators and flags (it was the legacy parser): 'a' 'b', f(-), [-], -(-), "a""b", 0' ', ab (a code list), {}, f(==>), 1.0Inf all read as written. Exactly one clause is consumed per call, so a syntax error raises `error(syntax_error(D), _)` and the next read starts on the next term. The variables of a term read are fresh cells (two reads never share one). A non-stream argument raises the ISO errors: `instantiation_error`, `domain_error(stream_or_alias, S)`, `existence_error(stream, S)`, `permission_error(input, stream, S)` (it was a message atom).

read_term/2-3

Purpose: Read a term with control options.

- `read_term(-Term, +Options)` reads from the current input
- `read_term(+Stream, -Term)` reads from a stream with default options
- `read_term(+Stream, -Term, +Options)` - the primary ISO 8.14.1 form (*added v3.5.0*)

Supported options: `variables(Vars)` (all variables of the term), `variable_names(Pairs)` (Name=Var pairs for the named variables), `singletons(Pairs)` (Name=Var pairs for singleton variables), `syntax_errors(error|fail|quiet)`, `term_position(Pos)` (v3.14.0: the '\$stream_position'(CharCount, LineCount, LinePosition, ByteCount) term for the FIRST character of the term just read - feed it to `stream_position_data/3`).

```
?- read_term(T, [variable_names(Vs)]).
|: foo(X, Y, X).
T = foo(_A, _B, _A),
Vs = ['X'=_A, 'Y'=_B].

% From a stream (v3.5.0+)
?- open('data.pl', read, S), read_term(S, T, []), close(S).
T = foo(1, 2).
```

v3.5.0: `read_term/2, 3` succeed as goals (previously the read happened but the goal failed). On end-of-file, `Term` is bound to the atom `end_of_file`.

v3.6.0: like `read/1, 2`, reads up to the ISO end token instead of one line - multi-line terms, several terms per line, and leading comments all work, with the stream position preserved between calls.

v3.14.0: reads through the stream's own decoder, so `read_term/2, 3`, `get_char/2` and `seek/4` all agree about where the stream is (before this each of `read/1, 2`, `read_term/2, 3` and `get_char/2` kept its own buffer). Input nested more than 1000 levels deep raises `error(resource_error(parser_nesting), _)` rather than blowing the Java stack or failing silently.

v4.5.0 (ISS-2025-0566): on the v2 parser. Options: `variables/1`, `variable_names/1`, `singletons/1` (named variables occurring once, `_X` names excluded), `term_position/1`, `double_quotes/1` (overrides the flag for this read), `syntax_errors(error|fail|quiet|dec10)`, `comments/1`; `module/1`, `backquoted_string/1`, `cycles/1`, `dotlists/1`, `var_prefix/1` are accepted and ignored; anything else is `domain_error(read_option, 0)`. An unknown stream raises `existence_error(stream, S)` (it used to read stdin). Nesting is bounded by a real limit of 200 000 levels (`resource_error(parser_nesting)`), not 1 000.

v4.6.0 (ISS-2025-0738, wave Q3.7): `subterm_positions(P)` gives SWI's layout term, character offsets in the stream: `From-To` for an atom, number or variable, `string_position(From, To)`, `brace_term_position(From, To, Arg)`, `list_position(From, To, Elms, Tail)` (`Tail = none` without `|`), `term_position(From, To, FFrom, FTo, Args)` for a compound (the functor or operator span), and `parentheses_term_position(From, To, Inner)`. `comments(C)` gives the comments read with the term as '\$stream_position'(Char, Line, LinePos, Byte)-"text" pairs.

```
?- open_string("foo(X, bar).", S), read_term(S, T, [subterm_positions(P)]).
P = term_position(0, 11, 0, 3, [4-5, 7-10]).
```

read_term_from_atom/3

Purpose: `read_term_from_atom(+Atom, -Term, +Options)` parses the text of an atom (or string) as one term, with the `read_term/2` options. (*added v4.5.0, ISS-2025-0566*)

```
?- read_term_from_atom('g(A, B, A)', T, [variable_names(V)]).
T = g(_A, _B, _A),
V = ['A'=_A, 'B'=_B].
```

open_string/2

Purpose: `open_string(+Text, -Stream)` opens an input stream over an atom, string, number, code list or char list, so every stream reader (`read/2`, `read_term/3`, `get_char/2`, `read_line_to_string/2`, ...) can read from memory. Close it with `close/1`. (added v4.5.0, ISS-2025-0569)

```
?- open_string("foo(1). bar(2).", S), read(S, A), read(S, B), close(S).
A = foo(1), B = bar(2).
```

with_input_from/2

Purpose: `with_input_from(+Source, :Goal)` runs `Goal` once with the current input redirected to `atom(A)`, `string(S)`, `codes(Cs)` or `chars(Cs)`; the input is restored afterwards. Any other source is `domain_error(input_source, S)`. (added v4.5.0, ISS-2025-0569)

```
?- with_input_from(atom('t(1). t(2).'), (read(X), read(Y))).
X = t(1), Y = t(2).
```

read_line_to_string/2, read_line_to_codes/2,3

Purpose: `read_line(SWI library(readutil)). read_line_to_string(+S, -Str)` gives the line without its newline (a trailing CR is dropped), or `end_of_file`; `read_line_to_codes(+S, -Codes)` gives a code list, or `-1` at end of file; `read_line_to_codes(+S, -Codes, ?Tail)` is the difference-list form that KEEPS the newline and closes the list with `[]` at end of file. (added v4.5.0, ISS-2025-0569)

read_string/3, read_string/5

Purpose: `read_string(+S, ?Length, -String)` reads `Length` characters (all that remain when `Length` is unbound, which is then bound to the count). `read_string(+S, +SepChars, +PadChars, -Sep, -String)` reads up to the first character of `SepChars` (bound to `Sep` as a code, `-1` at end of file) and strips the characters of `PadChars` from both ends. (added v4.5.0, ISS-2025-0569)

```
?- open_string(" a , b", S), read_string(S, ",", " ", Sep, Str).
Sep = 44, Str = "a".
```

Character I/O

get_char/1-2 and put_char/1-2

Purpose: Read or write single characters.

When to use: Use for character-by-character processing.

v3.5.0: the stream form `put_char(Stream, Char)` writes the character to the given stream; `get_char/1` honours the current input set by `set_input/1`.

```
% Read a single character
?- get_char(C).
% User types: a
C = 'a'.

% Write a single character
?- put_char('X').
X
true.

% Read password with echo off (conceptual example)
```

```

read_password(Password) :-
    read_password_chars([], Password).

read_password_chars(Acc, Password) :-
    get_char(C),
    (   C = '\n'
    -> reverse(Acc, Password)
    ;   put_char('*'), % Echo asterisk instead
        read_password_chars([C|Acc], Password)
    ).

```

get_code/1-2, peek_char/1-2 and peek_code/1-2

Purpose: Read the next character code (get_code), or look ahead at the next character (peek_char) / character code (peek_code) without consuming it.

```

?- get_code(C).           % from current input; C is the character code
?- get_code(Stream, C).   % from a specific stream (v3.5.0+)

?- peek_char(C).          % C is the next character (an atom), not consumed
?- peek_char(Stream, C).

?- peek_code(C).          % C is the next character code, not consumed
?- peek_code(Stream, C).

```

On end-of-file, get_code/peek_code unify the code with -1 and peek_char unifies the character with end_of_file. (*stream-argument forms added v3.5.0*)

Byte I/O

get_byte/1-2 and put_byte/1-2

Purpose: Read or write single bytes for binary stream operations.

```

% Read a byte from current input
?- get_byte(B).

% Read a byte from a specific stream
?- get_byte(Stream, B).

% Write a byte to current output
?- put_byte(65). % Writes 'A'

% Write a byte to a specific stream
?- put_byte(Stream, 65).

```

peek_byte/1-2

Purpose: Non-consuming byte lookahead.

```

?- peek_byte(B).          % Peek from current input
?- peek_byte(Stream, B). % Peek from specific stream

```

at_end_of_stream/0-1

Purpose: Test whether end of stream has been reached.

```

?- at_end_of_stream.      % Current input
?- at_end_of_stream(Stream). % Specific stream

```

write_canonical/1-2

Purpose: Write term in canonical (functor) notation that can be read back.

```
?- write_canonical(1+2).
% Outputs: '+'(1,2)

?- write_canonical(Stream, Term).
```

write_term/2-3

Purpose: Write a term under explicit control options.

- `write_term(+Term, +Options)` writes to the current output
- `write_term(+Stream, +Term)` writes to a stream with default options
- `write_term(+Stream, +Term, +Options)` - the primary ISO 8.14.2 form (*added v3.5.0*)

Supported options (*complete since v3.14.0*):

Option	Meaning
<code>quoted(Bool)</code>	quote atoms and strings so the output re-reads as the same term
<code>ignore_ops(Bool)</code>	never use operator notation: <code>1-2</code> prints as <code>-(1,2)</code>
<code>numbervars(Bool)</code>	render '\$VAR'(N) as A, B, ..., Z, A1, ... (default false for <code>write_term</code> , true for <code>write/1</code> and <code>print/1</code>)
<code>max_depth(N)</code>	print at most N nesting levels; deeper structure prints as ... and a longer list as <code>[E1,E2 ...]</code> . 0 means unlimited
<code>portray(Bool)</code>	call the user's <code>portray/1</code> on every subterm; what it writes replaces the subterm
<code>cycles(Bool)</code>	render a rational tree as <code>@(Template, Substitutions) - X = f(X)</code> prints as <code>@(_S1, [_S1=f(_S1)])</code> . With <code>cycles(false)</code> a back edge prints as ..., so a cyclic term always terminates
<code>variable_names(Pairs)</code>	a list of Name=Var pairs; each listed variable prints as its name
<code>spacing(standard next_argument)</code>	<code>next_argument</code> puts a space after every argument separator: <code>f(a, b, c)</code>

An unrecognised option raises error(`domain_error(write_option, Option), _`).

```
?- write_term('hello world', [quoted(true)]).
'hello world'
true.

?- write_term('$VAR'(0), [numbervars(true)]).
A
true.

?- write_term(Stream, 1+2, [ignore_ops(true)]).
% Writes: +(1,2)
```

v3.5.0: `write_term/2` succeeds as a goal (previously it printed and then failed); the three-argument stream form was added.

v3.14.0: the whole family (`write/1,2`, `writeln/1,2`, `writeq/1,2`, `print/1,2`, `write_canonical/1,2`, `write_term/2,3`, `portray_clause/1,2`, `format ~w/~q/~p`) renders through one writer, `core.engine.v4.Writer`. It is iterative (a 1 000 000-element list and a 1 000 000-deep structure print at the default JVM stack), cycle-safe, and reads the **engine's own** operator table, so an operator declared by `:- op/3` in a consulted file is honoured when the term is written back.

```
?- write_term([1,2,3,4,5], [max_depth(3)]).
[1,2|...]
```

```
?- write_term(f(X,Y), [variable_names(['Foo'=X, 'Bar'=Y])]).
f(Foo,Bar)

?- X = f(X), write_term(X, [cycles(true)]).
@(_S1,[_S1=f(_S1)])
```

portray_clause/1-2

Purpose: Write a clause the way listing/1 does - quoted, variables numbered A, B, ..., one body goal per indented line, terminated by a full stop and a newline. (*added v3.14.0*)

```
?- portray_clause((p(X,Y) :- q(X), r(Y))).
p(A,B) :-
    q(A),
    r(B).
true.

?- portray_clause(fact).
fact.
true.
```

`portray_clause(+Stream, +Clause)` writes to `Stream`.

v4.5.0 (ISS-2025-0570, wave P3.7): SWI's layout - quoted, operators, A, B argument spacing, variables A, B, ... in order of first occurrence and `_` for a singleton, one body goal per line, and if-then-else / disjunction as an indented block:

```
?- portray_clause((p(X, Y) :- (X > 0 -> Y = pos ; Y = neg), q(Z))).
p(A, B) :-
    (
        A>0
        -> B=pos
        ; B=neg
    ),
    q(_).
```

print_message/2

v4.5.0 (ISS-2025-0607): `native.format(Format, Args)` is formatted; `error(Format, context(PI, Msg))` is rendered SWI-style (`PI: Type error: ...`); any other term prints `Unknown message: T`. Kinds `error/warning` go to `user_error` with `ERROR: /Warning:` on every line, the others to the current output with `%`; `silent` prints nothing. **Purpose:** `print_message(+Kind, +Message)` - report a message. (*added v3.14.0*)

An ISO `error(Format, Context)` ball is rendered readably; anything else is written with `quoted(true)`.

v4.6 (ISS-2025-0713): the SWI extension points. A user prolog:`message//1` rule translates a message term into lines (`Format-Args, Format, nl, ansi(Attrs, Format, Args), url(L), flush`); `user:message_hook(+Message, +Kind, +Lines)` is called first (also for `silent`) and, when it succeeds, the message counts as printed. EVERY kind is printed on `user_error` (SWI) with `ERROR: , Warning: or %` on each line; the flag `verbose = silent` suppresses the informational and banner kinds. A hook or rule that itself calls `print_message/2` is not re-entered; one that throws is ignored.

```
?- print_message(error, error(type_error(integer, abc), foo/1)).
ERROR: Type error: `integer' expected, found `abc' (foo/1)
true.

?- print_message(info, hello).
% Unknown message: hello
true.
```

```

:- multifile prolog:message//1.
prolog:message(greet(Who)) --> ['Hello ~w'-[Who], nl, 'nice to see you'].

?- print_message(informational, greet(bob)).
% Hello bob
% nice to see you
true.

message_hook(quiet(_), _, _).          % swallow quiet(...) messages

?- print_message(warning, quiet(x)).
true.

```

Stream properties and positions

stream_property/2

Purpose: `stream_property(?Stream, ?Property)` relates an open stream of **this engine** to its properties. Both arguments may be unbound; an unbound `Stream` enumerates every open stream.

Properties (*complete since v3.14.0*): `file_name(F)`, `mode(read|write|append)`, `input`, `output`, `alias(A)` (one solution per alias), `position(P)`, `end_of_stream(not|at|past)`, `eof_action(error|eof_code|reset)`, `reposition(true|false)`, `type(text|binary)`, `encoding(E)`, `line_count(N)`.

`P` is the opaque term `'$stream_position'` (`CharCount`, `LineCount`, `LinePosition`, `ByteCount`); take it apart with `stream_position_data/3` and feed it back to `set_stream_position/2`.

```

?- open('data.txt', read, S, [alias(input)]), stream_property(input, mode(M)).
M = read.

```

set_stream/2

Purpose: `set_stream(+Stream, +Property)` changes a property of an open stream. (*added v3.14.0*)

Accepted: `alias(A)` (adds another name for the stream), `type(text|binary)`, `eof_action(error|eof_code|reset)`, `encoding(E)`. Anything else raises `domain_error(stream_property, P)`.

stream_position_data/3

Purpose: `stream_position_data(+Field, +Position, ?Data)` extracts one field of a `'$stream_position'/4` term: `char_count`, `line_count`, `line_position` or `byte_count`. (*added v3.14.0*)

character_count/2, line_count/2, line_position/2

Purpose: The three counters of an open stream, without going through `stream_property/2`. (*added v3.14.0*)

`line_count/2` is 1-based (a freshly opened stream is on line 1); `line_position/2` is the 0-based column; `character_count/2` counts characters read or written so far. All three are exact on **text** streams, because the stream decodes through its own buffer.

```

?- open('two_lines.txt', read, S), get_char(S,_), get_char(S,_), get_char(S,_),
  get_char(S,_), get_char(S,_), get_char(S,_), get_char(S,C),
  line_count(S, L), line_position(S, LP).
C = w, L = 2, LP = 1.

```

current_stream/3

Purpose: `current_stream(?File, ?Mode, ?Stream)` enumerates the open **file** streams of this engine. (added v3.14.0)

seek/4, stream_position/2 and set_stream_position/2

`seek(+Stream, +Offset, +Method, -NewLocation)` with Method one of `bof`, `current`, `eof`; `stream_position(+Stream, -ByteOffset)`; `set_stream_position(+Stream, +Position)` where Position is a byte offset or a '\$stream_position'/4 term.

v3.14.0: **repositioning now works on text streams.** Before this release the stream handed out a `PushbackReader` over the raw file with its own 8 KB buffer, so a seek moved the file channel but not the reader - `get_char(S,C1)`, `seek(S,0,bof,_)`, `get_char(S,C2)` answered `C2 = e` after `C1 = h`. A reposition now flushes the stream's decode buffer and resets the decoder, and it recomputes the character/line counters.

Streams are per engine

v3.14.0: the stream table, its aliases and `current_input/current_output` belong to the Prolog instance (and, for the current streams, to the calling thread). A stream opened by one engine is invisible to another, and `set_output/1` on one thread does not redirect another thread's output. `open/3,4` unifies Stream with the canonical term '\$stream'(N); every stream argument also accepts an atom alias, the `stream_<id>` handle, and the reserved `user_input/user_output/user_error/current_input/current_output`.

char_conversion/2 and current_char_conversion/2

Purpose: Manage character conversion table used during term reading.

```
% Define conversion: 'A' reads as 'a'
?- char_conversion('A', a).

% Query current conversion
?- current_char_conversion('A', X).
X = a.

% Remove conversion (convert to self)
?- char_conversion('A', 'A').
```

File I/O

open/3 and close/1

Purpose: Open and close file streams.

When to use: Use for file-based I/O operations.

v3.5.0: failures raise `ISO error/2` terms instead of plain-atom exceptions - opening a missing file raises `existence_error(source_sink, File)`, an invalid stream or alias raises `domain_error(stream_or_alias, S) / existence_error(stream, S)`. Also applies to `open/4`, which takes an options list (e.g. `alias(Name)`).

v4.5.0 (ISS-2025-0625): `Prolog.enableSafeMode()` removes `open/3,4`; with `SafeModeOptions.allowFileRead(dir)` they stay for `read` mode on files inside `dir` only.

```
% Open file for reading
read_file_terms(Filename, Terms) :-
```

```

open(Filename, read, Stream),
read_all_terms(Stream, Terms),
close(Stream).

read_all_terms(Stream, []) :-
    at_end_of_stream(Stream), !.
read_all_terms(Stream, [Term|Terms]) :-
    read(Stream, Term),
    read_all_terms(Stream, Terms).

% Open file for writing
write_to_file(Filename, Data) :-
    open(Filename, write, Stream),
    write(Stream, Data),
    write(Stream, '.'), % Add period for Prolog term
    nl(Stream),
    close(Stream).

% Append to file
append_to_log(Message) :-
    open('log.txt', append, Stream),
    get_time(Time),
    write(Stream, Time), write(Stream, ': '),
    writeln(Stream, Message),
    close(Stream).

```

set_input/1 and set_output/1

Purpose: Make a stream the current input / current output.

```

% Redirect input: read/1, get_char/1, get_code/1, ... now read from the stream
?- open('data.pl', read, S), set_input(S), read(T).
T = foo(1, 2).

% Redirect output: write/1, nl/0, format/2, ... now write to the stream
?- open('out.txt', write, S), set_output(S), writeln(hello).
true.

```

v3.5.0: these predicates actually redirect - the arity-1 input predicates honour the current input, and the output built-ins honour the current output (previously the setting was recorded but ignored).

tab/1-2

v4.5.0 (ISS-2025-0604): the count is an arithmetic expression (`tab(1+1)` prints two spaces); a non-evaluable is `type_error(evaluable, F/N)`, a float `type_error(integer, F)`; a negative count prints nothing (SWI). **Purpose:** Write N space characters to current output (`tab/1`) or to a given stream (`tab(Stream, N)`), *added v3.5.0*.

```

?- write(hello), tab(5), write(world).
hello      world

```

with_output_to/2

v4.5.0 (ISS-2025-0606): the capture replaces the current output as it was on entry; a `set_output/1` to another stream inside the goal (or a write to an explicit stream) is NOT captured. **Purpose:** Execute Goal once, capturing everything it writes into Sink.

Sinks: `atom(A)` on every engine; `string(S)`, `codes(C)` and `chars(C)` in addition on the v4 engine (v3.10.0, ISS-2025-0452), where a target that is none of those raises `domain_error(output_sink, T)`. Goal runs as `once/1`: only the first solution's output is captured, and `with_output_to/2` fails if Goal fails.

```
?- with_output_to(atom(X), write(hello)).  
X = hello.  
  
?- with_output_to(codes(C), write(ab)).      % v4  
C = [97, 98].
```

v3.10.0, engine v4: the capture uses the thread-local output stream and restores whatever was installed (the IDE installs one per background solve), instead of swapping `System.out` and clearing it afterwards. See LIM-025.

8. Database Operations

Database operations allow dynamic modification of the Prolog knowledge base at runtime.

Understanding the Dynamic Database

Prolog's database can be modified during program execution:

- Add new facts and rules with `assert` family
- Remove facts and rules with `retract` family
- Query what's currently in the database with `listing`

`assert/1`, `asserta/1`, `assertz/1`

Purpose: Add facts or rules to the database.

- `assert/1` and `assertz/1` add at the end
- `asserta/1` adds at the beginning

When to use: Use to store runtime data, learn new information, or build dynamic knowledge bases.

v3.5.0: ISO validation - asserting an unbound term raises `instantiation_error`; a non-callable clause (`1, (1 :- true), (foo :- 7)`) raises `type_error(callable, _)`; asserting over a built-in raises `permission_error(modify, static_procedure, Name/Arity)`. Asserting a predicate implies declaring it dynamic (see `dynamic/1`).

```
% Add a simple fact
?- assertz(likes(john, pizza)).
true.

?- likes(john, X).
X = pizza.

% Add a rule
?- assertz((parent(X, Y) :- father(X, Y))).
true.

% Difference between asserta and assertz
?- assertz(color(red)).
true.
?- assertz(color(blue)).
true.
?- asserta(color(green)). % Added at beginning
true.

?- color(X).
X = green ; % First (added with asserta)
X = red ; % Second
X = blue. % Third

% Practical example: Learning system
learn_fact :-
    write('What did you learn? '),
    read(Fact),
    ( \+ Fact % Check if not already known
    -> assertz(Fact),
        writeln('I learned something new!')
    ; writeln('I already knew that!')
    ).

% Cache computation results
fibonacci_cached(N, Result) :-
    ( fib_cache(N, Result) % Check cache first
```

```
-> true
; calculate_fibonacci(N, Result),
  assertz(fib_cache(N, Result)) % Cache the result
).
```

retract/1 and retractall/1

Purpose: Remove facts or rules from the database.

- `retract/1` removes first matching clause
- `retractall/1` removes all matching clauses

When to use: Use to remove outdated information, clean up temporary data, or implement undo functionality.

v3.5.0: ISO validation - `retract/1` validates its argument (`instantiation_error` for an unbound term, `type_error(callable, _)` for a non-callable one, instead of an internal error); retracting clauses of a built-in raises `permission_error(modify, static_procedure, Name/Arity)`; `retractall/1` validates its argument the same way and implies declaring the predicate dynamic.

v3.6.0: `retract/1` is **re-executable on backtracking** on the v4 and v2 engines (ISO 8.9.3) - each redo retracts the next matching clause, so `findall(X, retract(p(X)), L)` drains the predicate one clause per solution; retractions of earlier solutions persist across backtracking. (The recursive engine, which retracted all matching clauses eagerly on the first call, was deleted in 4.0.0.)

```
% Remove a specific fact
?- assertz(temp(1)), assertz(temp(2)), assertz(temp(3)).
true.

?- retract(temp(2)).
true.

?- temp(X).
X = 1 ;
X = 3.

% Remove all matching facts
?- retractall(temp(_)).
true.

?- temp(X).
false. % All removed

% Practical example: Session management
login(User) :-
    retractall(current_user(_)), % Remove any existing login
    assertz(current_user(User)),
    write('Logged in as: '), writeln(User).

logout :-
    retractall(current_user(_)),
    writeln('Logged out').

% Update a fact (retract old, assert new)
update_score(Player, NewScore) :-
    retractall(score(Player, _)),
    assertz(score(Player, NewScore)).

% Temporary facts with cleanup
with_temp_fact(Fact, Goal) :-
    assertz(Fact),
    call(Goal),
    retract(Fact).
```

abolish/1

v4.5.0 (ISS-2025-0610): removes the predicate, not just its clauses: the dynamic declaration goes too, so a later call raises `existence_error(procedure, F/A)`. A float arity is `type_error(integer, A)`, an arity past the largest representable one `representation_error(max_arity)`. **Purpose:** Removes all clauses of a predicate.

When to use: Use to completely remove a predicate definition.

v3.5.0: abolishing a built-in raises `permission_error(modify, static_procedure, Name/Arity)`.

v4.4.0 (ISS-2025-0508): an unbound half of the indicator (`abolish(a/A)`, `abolish(A/1)`) is `instantiation_error` - it used to be a `type_error` whose culprit was a fresh variable.

```
% Remove entire predicate
?- assertz(test(1)), assertz(test(2)), assertz((test(X) :- X > 10)).
true.

?- abolish(test/1). % Remove all clauses of test/1
true.

?- test(X).
ERROR: Undefined predicate: test/1

% Practical example: Reset game state
reset_game :-
    abolish(player_position/2),
    abolish(player_score/2),
    abolish(game_object/3),
    initialize_game.

% Clear all temporary predicates
cleanup_temp :-
    abolish(temp/1),
    abolish(cache/2),
    abolish(session/1).
```

dynamic/1

v4.5.0 (ISS-2025-0610): `predicate_property(H, dynamic)` and `number_of_clauses(0)` hold for a declared dynamic predicate that has no clauses; `abolish/1` removes the declaration. **Purpose:** Declares procedures dynamic, so calling them while they have no clauses fails silently instead of raising `existence_error(procedure, Name/Arity)`. (*callable as a goal since v3.5.0*)

When to use: Declare every predicate you plan to assert/retract at runtime, so querying it before any clause exists fails instead of raising an error (with the `unknown` flag at its default error).

```
% Directive form (in a consulted file)
:- dynamic(counter/1).

% Goal form (v3.5.0+): a predicate indicator ...
?- dynamic(score/2).
true.

?- score(X, Y).
false. % no existence_error: score/2 is dynamic

% ... a ','-sequence of indicators, or a list of indicators
?- dynamic((foo/1, bar/2)).
true.

?- dynamic([baz/0, quux/3]).
true.
```

Note: `assert/retractall` on a predicate imply declaring it dynamic; the `:- dynamic` directive is honoured during `consult` (it was a no-op before v3.5.0).

listing/0 and listing/1

Purpose: Display current database contents.

When to use: Use for debugging, inspecting dynamic predicates, or showing current state.

v4.2.0: **listing/1** **works**. Every earlier release registered one implementation for the name `listing`, the arity-0 one, which rejected any argument - `listing(foo/1)` raised `listing/0` takes no arguments despite being documented here since v2. Both arities are now separate v4 natives; a bare name lists every arity, `Name/Arity` lists one, and both print through the current output (so `with_output_to/2` and the IDE console capture them) with one full stop per clause instead of two.

```
% List everything
?- listing.
% Shows all user-defined predicates

% List specific predicate
?- listing(likes/2).
likes(john, pizza).
likes(mary, wine).

% List predicates by name (all arities)
?- listing(parent).
parent(tom, bob).
parent(bob, ann).
parent(X, Y) :- father(X, Y).
parent(X, Y) :- mother(X, Y).

% Practical example: Show current configuration
show_config :-
    writeln('Current Configuration:'),
    writeln('====='),
    listing(config/2),
    listing(option/1).
```

v4.5.0 (ISS-2025-0570, wave P3.7): **listing output is re-readable**. Every clause goes through `portray_clause/1` (it printed `Rule.toString(): lst(_G27,_G26) :- ;(,(>(...,unquoted lst(A b, it's))`), predicates are separated by a blank line and a dynamic predicate starts with `:- dynamic Name/Arity.`; consulting the output gives back variant clauses. The `% Listing for p/1: header line` is gone (SWI prints none).

current_predicate/1

v4.5.0 (ISS-2025-0610): a declared dynamic predicate without clauses is current (SWI). v4.6.0 (ISS-2025-0732): `current_predicate(M:Name/Arity)` checks or enumerates module `M`'s own predicates (`M` bound or not): its clauses, its `M:H` clauses and its `:- multifile M:PI` declarations; `user:` is the flat knowledge base. It raised `type_error` before. **Purpose:** Check or enumerate defined predicates.

When to use: Use to check if predicates exist or list available predicates.

```
% Check if predicate exists
?- current_predicate(member/2).
true.

?- current_predicate(nonexistent/3).
false.

% Enumerate predicates
?- current_predicate(Name/2).
Name = append ;
```

```

Name = member ;
Name = select ;
...

% Find all predicates with specific arity
?- findall(P, current_predicate(P/3), Predicates).
Predicates = [append/3, select/3, nth0/3, ...].

% Practical example: Safe predicate call
safe_call_predicate(Name, Arity, Args) :-
    current_predicate(Name/Arity),
    length(Args, Arity),
    Goal =.. [Name|Args],
    call(Goal).
safe_call_predicate(Name, Arity, _) :-
    \+ current_predicate(Name/Arity),
    write('Undefined predicate: '), write(Name/Arity), nl,
    fail.

```

recorda/2, recorda/3, recordz/2, recordz/3, recorded/2, recorded/3, erase/1, current_key/1

v4.6.0 (ISS-2025-0791, wave Q7): SWI-Prolog's recorded database, as v4 natives. **Purpose:** a key-indexed store of terms outside the clause database.

- `recorda(+Key, +Term)` / `recordz(+Key, +Term)` add a COPY of Term first / last under Key; the /3 forms return a reference '\$record'(N) (the third argument must be unbound).
- `recorded(?Key, ?Term)` / `recorded(?Key, ?Term, ?Ref)` enumerate the records (a fresh copy each time), in order; an unbound Key enumerates every key; a bound Ref looks the record up directly. The logical update view applies: the records are taken when the call starts.
- `erase(+Ref)` removes a record; erasing an already erased record fails.
- `current_key(?Key)` enumerates the keys that have records.
- A key is an atom, an integer or a compound term, of which only the name and arity count (`foo(1, 2)` and `foo(a, b)` are the same key). Anything else is `type_error(key, K)`, an unbound key of `recorda/recordz` is `instantiation_error`, a bad reference `type_error(db_reference, R)`.
- The store belongs to the Prolog instance and is shared by its threads.

```

?- recordz(color, red), recordz(color, green), recorda(color, blue, R).
R = '$record'(3).

?- findall(C, recorded(color, C), L).
L = [blue, red, green].

?- recorded(color, red, R), erase(R), findall(C, recorded(color, C), L).
L = [blue, green].

```

flag/3

v4.6.0 (ISS-2025-0791): **Purpose:** `flag(+Key, -Old, +New)` - a global counter per key (SWI). Old is unified with the current value (0 for a new key), then New - an arithmetic expression evaluated after that unification, or an atom - becomes the value. Keys as for `recorda/2`; the update is atomic with respect to other threads of the engine.

```

?- flag(hits, N, N+1), flag(hits, M, M+1).
N = 0, M = 1.

```

9. Atom and String Operations

These predicates manipulate atoms (symbolic constants) and strings.

v3.6.0 (SWI-style text interop): the `atom_*` predicates accept strings (`atom_length("abc", 3)`, `atom_codes(X, "abc")` gives the atom `abc`), and `string_length/2`, `string_chars/2`, `string_concat/3` accept atoms (`string_concat(a, b, S)` gives the string `S = "ab"` - the result stays a string).

atom_length/2

Purpose: Determines the length of an atom.

When to use: Use for validation, formatting, or string processing.

v3.6.0: ISO error terms - an unbound atom raises `instantiation_error`, a non-atom raises `type_error(atom, A)`, a non-integer length raises `type_error(integer, L)` and a negative length raises `domain_error(not_less_than_zero, L)`. Also accepts a string first argument (SWI interop).

```
% Get length of atom
?- atom_length(hello, Len).
Len = 5.

?- atom_length('Hello World', Len).
Len = 11.

% Validate input length
validate_username(Username) :-
    atom(Username),
    atom_length(Username, Len),
    Len >= 3,
    Len <= 20,
    !.
validate_username(_) :-
    writeln('Username must be 3-20 characters'),
    fail.

% Pad atom to specific length
pad_atom(Atom, TargetLen, PadChar, Padded) :-
    atom_length(Atom, CurrentLen),
    PadCount is TargetLen - CurrentLen,
    (   PadCount <= 0
    -> Padded = Atom
    ;   create_padding(PadCount, PadChar, Padding),
        atom_concat(Atom, Padding, Padded)
    ).
```

atom_concat/3

Purpose: Concatenates atoms or splits an atom.

When to use: Use for building identifiers, messages, or parsing.

v3.6.0: ISO error terms - all arguments unbound raises `instantiation_error`, a non-atomic argument raises a proper error (`type_error(atom, Culprit)`, `_`) ball naming the culprit (was a plain-text exception). Also accepts string arguments (SWI interop).

```
% Concatenate atoms
?- atom_concat(hello, world, Result).
Result = helloworld.

?- atom_concat('Hello ', 'World', Result).
```

```

Result = 'Hello World'.

% Split atom (finding possible splits)
?- atom_concat(Prefix, Suffix, helloworld).
Prefix = '', Suffix = helloworld ;
Prefix = h, Suffix = elloworld ;
Prefix = he, Suffix = lloworld ;
...

% Find specific prefix/suffix
?- atom_concat(hello, Suffix, helloworld).
Suffix = world.

?- atom_concat(Prefix, world, helloworld).
Prefix = hello.

% Practical example: Build file paths
build_path(Dir, File, Path) :-
    atom_concat(Dir, '/', Temp),
    atom_concat(Temp, File, Path).

?- build_path('/home/user', 'file.txt', Path).
Path = '/home/user/file.txt'.

% Generate unique identifiers
generate_id(Base, Counter, ID) :-
    atom_number(CounterAtom, Counter),
    atom_concat(Base, '_', Temp),
    atom_concat(Temp, CounterAtom, ID).

?- generate_id(user, 42, ID).
ID = user_42.

```

sub_atom/5

v4.5.0 (ISS-2025-0599): positions and lengths count CODE POINTS (as `atom_length/2`), so a character outside the BMP is never split; a negative Before/Length/After FAILS (SWI). `sub_string/5` takes any text (atom, string, number, code list) for String and Sub. **Purpose:** Extracts substrings from atoms.

When to use: Use for parsing, pattern matching, or string manipulation.

v3.10.0, engine v4 (ISS-2025-0453): `sub_atom/5` and `sub_string/5` are lazy generators - one candidate per redo instead of the full $O(n^2)$ cross-product materialised up front, so `once(sub_atom(LongAtom, _, _, _, S))` is cheap. This also fixes a hang: with an **empty** SubAtom (`sub_atom(abc, B, L, A, '')`) the default engine loops forever, because `String.indexOf("", Idx)` stops advancing past the end of the atom; on v4 the four positions are enumerated once.

v4.4.0 (ISS-2025-0509): ISO 8.16.3.3 - `instantiation_error, type_error(atom, A)` for the atom and the sub-atom, `type_error(integer, N)` for Before/Length/After.

```

% Syntax: sub_atom(+Atom, ?Before, ?Length, ?After, ?SubAtom)
% Before: characters before the subatom
% Length: length of the subatom
% After: characters after the subatom

% Extract substring
?- sub_atom(helloworld, 5, 5, 0, Sub).
Sub = world.

?- sub_atom(helloworld, 0, 5, _, Sub).
Sub = hello.

% Find position of substring
?- sub_atom(helloworld, Before, 5, After, world).

```

```

Before = 5, After = 0.

% Check if atom contains substring
contains_substring(Atom, Sub) :-
    sub_atom(Atom, _, _, _, Sub).

?- contains_substring('hello world', world).
true.

% Practical example: Extract file extension
get_extension(Filename, Ext) :-
    sub_atom(Filename, Before, 1, After, '.'),
    After > 0, % Ensure dot is not at the end
    !,
    sub_atom(Filename, _, After, 0, Ext).

?- get_extension('document.pdf', Ext).
Ext = pdf.

% Parse email address
parse_email(Email, User, Domain) :-
    sub_atom(Email, Before, 1, After, '@'),
    sub_atom(Email, 0, Before, _, User),
    sub_atom(Email, _, After, 0, Domain).

?- parse_email('john@example.com', U, D).
U = john, D = 'example.com'.

```

atom_chars/2 and atom_codes/2

Purpose: Convert between atoms and character/code lists.

When to use: Use for character-level processing or encoding conversions.

v3.6.0: ISO error terms - with the atom side unbound, a partial list or a list with an unbound element raises `instantiation_error`; a bad element raises `type_error(character, E)` (`atom_chars`) or `representation_error(character_code)` (`atom_codes`). A number first argument stringifies (`atom_chars(42, L)` gives `L = ['4', '2']`, SWI/GNU behavior). Both also accept a string first argument (SWI interop).

```

% atom_chars/2 - Convert to/from character list
?- atom_chars(hello, Chars).
Chars = [h, e, l, l, o].

?- atom_chars(Atom, [h, e, l, l, o]).
Atom = hello.

% atom_codes/2 - Convert to/from ASCII codes
?- atom_codes(hello, Codes).
Codes = [104, 101, 108, 108, 111].

?- atom_codes(Atom, [72, 69, 76, 76, 79]).
Atom = 'HELLO'.

% Practical example: Reverse an atom
reverse_atom(Atom, Reversed) :-
    atom_chars(Atom, Chars),
    reverse(Chars, RevChars),
    atom_chars(Reversed, RevChars).

?- reverse_atom(hello, Rev).
Rev = olleh.

% Simple encryption (Caesar cipher)

```

```

caesar_cipher(Text, Shift, Encrypted) :-
    atom_codes(Text, Codes),
    maplist(shift_code(Shift), Codes, ShiftedCodes),
    atom_codes(Encrypted, ShiftedCodes).

shift_code(Shift, Code, Shifted) :-
    Shifted is Code + Shift.

?- caesar_cipher(abc, 1, Encrypted).
Encrypted = bcd.

```

number_chars/2 and number_codes/2

Purpose: Convert between numbers and character/code lists.

```

?- number_chars(42, Chars).
Chars = ['4', '2'].

?- number_chars(N, ['3', '.', '1', '4']).
N = 3.14.

?- number_codes(255, Codes).
Codes = [50, 53, 53].

```

v3.5.0: arbitrarily large integers round-trip exactly (values beyond 64-bit precision were previously corrupted silently).

v3.6.0: type-faithful floats - float-syntax text yields a float (number_chars(X, ['1', '.', '0']) gives X = 1.0, not the integer 1) and floats keep float syntax on output (number_codes(1.0, L) gives "1.0"). ISO 0x/0o/0b radix notation and 0'c char-code constants are accepted; Java-only spellings (Infinity, NaN, 1f, '.5', '3.') raise syntax_error(illegal_number). ISO error terms otherwise: both sides unbound raises instantiation_error, a non-number first argument raises type_error(number, N).

atom_number/2

Purpose: Converts between atoms and numbers.

When to use: Use for parsing numeric input or formatting numbers.

v2.8.2: accepts hex (0xFF), binary (0b1010), octal (0o77) prefixes when parsing atom->number.

v3.5.0: arbitrarily large integers convert exactly in both directions (no more silent 64-bit corruption).

v3.6.0: float syntax is preserved in the number->atomdirection - atom_number(A, 123.0) gives A = '123.0' (previously '123'); the atom->numberdirection is type-faithful (atom_number('1.0', X) gives the float 1.0) and rejects Java-only spellings ('Infinity', 'NaN', '1f') with syntax_error(illegal_number).

v4.4.0 (ISS-2025-0506): neither argument bound is instantiation_error; a bound first argument that is not text is type_error(atom, A). An atom that is not a number still FAILS.

```

% Convert atom to number
?- atom_number('42', N).
N = 42.

?- atom_number('3.14', N).
N = 3.14.

% Hex / binary / octal (v2.8.2+)
?- atom_number('0xFF', N).
N = 255.

?- atom_number('0b1010', N).

```

```

N = 10.

?- atom_number('0o77', N).
N = 63.

% Convert number to atom
?- atom_number(A, 42).
A = '42'.

% Practical example: Parse numeric input
parse_number(Input, Number) :-
    atom(Input),
    atom_number(Input, Number),
    !.
parse_number(Input, _) :-
    write('Invalid number: '), write(Input), nl,
    fail.

% Calculate from string expression
calculate_string(Expr, Result) :-
    atom_number(Expr, Result). % Works for simple numbers

?- calculate_string('123', R).
R = 123.

% Format numbers with specific precision (conceptual)
format_currency(Amount, Formatted) :-
    Round is round(Amount * 100) / 100,
    atom_number(Amount, Round),
    atom_concat('$', AtomAmount, Formatted).

?- format_currency(42.3456, F).
F = '$42.35'.

```

String Processing

split_string/4

Purpose: Splits strings by separators with padding removal.

When to use: Use for parsing CSV, processing user input, or tokenization.

v4.4.0 (ISS-2025-0506): `instantiation_error` for an unbound argument and `type_error(string, S)` for a non-string one.

```

% Syntax: split_string(+String, +Separators, +PadChars, -SubStrings)

% Basic splitting
?- split_string("apple,banana,cherry", ",", "", L).
L = [apple, banana, cherry].

% Multiple separators
?- split_string("apple;banana,cherry:date", ";;:", "", L).
L = [apple, banana, cherry, date].

% Remove padding (spaces)
?- split_string(" apple , banana , cherry ", ",", " ", L).
L = [apple, banana, cherry].

% Practical example: Parse CSV line
parse_csv_line(Line, Fields) :-
    split_string(Line, ",", "\t", Fields).

?- parse_csv_line("John, 25, London", Fields).

```

```

Fields = ['John', '25', 'London'].

% Parse configuration line
parse_config(Line, Key, Value) :-
    split_string(Line, "=", " ", [Key, Value]).

?- parse_config("username = john_doe", K, V).
K = username, V = john_doe.

% Tokenize sentence
tokenize(Sentence, Tokens) :-
    split_string(Sentence, " .,!?", " ", Tokens).

?- tokenize("Hello, world! How are you?", T).
T = ['Hello', 'world', 'How', 'are', 'you'].

```

atomic_list_concat/2, atomic_list_concat/3

v4.5.0 (ISS-2025-0597): SWI modes and errors - the split mode also fills the holes of a partial list (`atomic_list_concat([a,B,c], '- ', 'a-x-c')` gives `B = x`); a compound element is `type_error(atomic, E)`, an unbound element in join mode `instantiation_error`, a non-list `type_error(list, L)`, an empty separator in split mode `domain_error(non_empty_atom, '')`; any atomic separator. **Purpose:** Joins atoms (optionally with separator) or splits by separator.

v4.4.0 (ISS-2025-0506): an unbound or partial list, an unbound separator, or neither list nor atom bound is `instantiation_error`; a non-atom separator or atom argument is `type_error(atom, A)`.

- `atomic_list_concat/2` (v2.8.2+): joins without separator
- `atomic_list_concat/3`: joins with separator; reverse mode splits

When to use: Use for building formatted strings or parsing.

```

% Join atoms with separator
?- atomic_list_concat([hello, world], ' ', Result).
Result = 'hello world'.

?- atomic_list_concat([one, two, three], '- ', Result).
Result = 'one-two-three'.

% Split by separator
?- atomic_list_concat(List, '- ', 'one-two-three').
List = [one, two, three].

% Join without separator (use atomic_list_concat/2)
?- atomic_list_concat([hello, world], Result).
Result = helloworld.

% Practical example: Build SQL query
build_insert(Table, Values, Query) :-
    atomic_list_concat(Values, ', ', ValueString),
    atomic_list_concat(['INSERT INTO ', Table, ' VALUES (', ValueString, ')'], Query).

?- build_insert(users, ['John', 25, 'London'], Q).
Q = 'INSERT INTO users VALUES (John, 25, London)'.

% Create formatted message
format_message(Template, Args, Message) :-
    atomic_list_concat(Args, ', ', ArgString),
    atomic_list_concat([Template, ': ', ArgString], Message).

?- format_message('Error', [404, 'Not Found'], Msg).
Msg = 'Error: 404, Not Found'.

```

string_to_atom/2

Purpose: Convert between string and atom representations. Bidirectional.

```
?- string_to_atom(hello, X).
X = hello.

% (-, +) mode binds a string (v3.6.0)
?- string_to_atom(S, foo).
S = "foo".
```

v3.6.0: the (-, +) mode binds a **string** - `string_to_atom(S, foo)` gives `S = "foo"` with `string(S)` true and `atom(S)` false (it used to bind an atom, making the mode a no-op). The (+, -) direction is unchanged.

number_to_atom/2, atom_to_number/2

Purpose: Convert between number and atom representations.

```
?- number_to_atom(42, X).
X = '42'.

?- atom_to_number('3.14', X).
X = 3.14.
```

string_upper/2, string_lower/2

Purpose: Convert text to upper/lower case, giving a string. v4.5.0 (ISS-2025-0598).

```
?- string_upper("hello World", U).
U = "HELLO WORLD".

?- string_lower('ABC', L).
L = "abc".
```

v4.5.0 (ISS-2025-0596): `string_concat/3`, `string_length/2`, `atom_string/2` and `sub_string/5` accept any atomic text (and code/char lists): `string_concat(1, 2, S)` gives `S = "12"`, `string_length(123, L)` gives 3, `atom_string(42, S)` gives `"42"` and `atom_string(A, 42)` gives `'42'`. `string_concat(X, Y, Z)` with nothing bound raises `instantiation_error` (it failed); a compound argument raises `type_error`.

string_code/3

Purpose: Get the character code at a 1-based index in a string/atom.

```
?- string_code(1, hello, X).
X = 104. % ASCII code for 'h'
```

10. Character Processing

Character processing predicates work with individual characters and their properties.

char_code/2

Purpose: Converts between characters and their numeric codes.

When to use: Use for character encoding, ASCII operations, or character arithmetic.

v3.6.0: ISO error terms (8.16.6.3) - both arguments unbound raises `instantiation_error`, a first argument that is not a one-char atom raises `type_error(character, C)`, a non-integer code raises `type_error(integer, Code)`, and an integer outside the Unicode range raises `representation_error(character_code)` (previously these failed silently).

```
% Character to code
?- char_code(a, Code).
Code = 97.

?- char_code('A', Code).
Code = 65.

% Code to character
?- char_code(Char, 65).
Char = 'A'.

?- char_code(Char, 97).
Char = a.

% Practical example: Check character type
is_uppercase(Char) :-
    char_code(Char, Code),
    Code >= 65,
    Code <= 90.

is_lowercase(Char) :-
    char_code(Char, Code),
    Code >= 97,
    Code <= 122.

is_digit(Char) :-
    char_code(Char, Code),
    Code >= 48,
    Code <= 57.

?- is_uppercase('A').
true.

?- is_digit('5').
true.

% Convert case
to_uppercase(Lower, Upper) :-
    is_lowercase(Lower),
    char_code(Lower, LowerCode),
    UpperCode is LowerCode - 32,
    char_code(Upper, UpperCode).

?- to_uppercase(a, U).
U = 'A'.
```

char_type/2

v4.5.0 (ISS-2025-0601): space is SWI's iswspace set (9-13, 32 and the Unicode spaces; no longer 28-31), white is space and tab; new xdigit(Weight), prolog_var_start, prolog_atom_start, prolog_identifier_continue, prolog_symbol; code_type(-1, end_of_file); characters outside the BMP are accepted; an unknown class is domain_error(char_type, T). **Purpose:** Classifies characters into categories.

When to use: Use for parsing, validation, or text processing.

```
% Check character type
?- char_type(a, alpha).
true.

?- char_type('5', digit).
true.

?- char_type(' ', space).
true.

% Find all types of a character
?- char_type('A', Type).
Type = alnum ;    % Alphanumeric
Type = alpha ;   % Alphabetic
Type = ascii ;   % ASCII character
Type = upper ;   % Uppercase
...

% Character type categories:
% - alnum: alphanumeric
% - alpha: alphabetic
% - ascii: ASCII character
% - cntrl: control character
% - digit: decimal digit
% - graph: graphical character
% - lower: lowercase letter
% - print: printable character
% - punct: punctuation
% - space: whitespace (white and layout are synonyms)
% - upper: uppercase letter
% - xdigit: hexadecimal digit
% - csym: letter, digit or underscore      % since 4.3.0
% - csymf: letter or underscore            % since 4.3.0
% - period: . ! ?                          % since 4.3.0
% - quote: " ' `                          % since 4.3.0
% - paren: ( )                            % since 4.3.0
% - newline, end_of_line, end_of_file, layout, meta, solo, symbol

% Parametric forms (since 4.3.0) - they work in EVERY mode: with the argument bound
% they test, with it unbound they bind, and with the character unbound they generate.
% char_type/2 gives a CHARACTER, code_type/2 a CODE; digit(Weight) gives an integer
% in both.
?- char_type('7', digit(W)).
W = 7.

?- char_type('A', upper(L)).      % 'A' is uppercase, with lowercase L
L = a.

?- char_type(a, lower(U)).        % a is lowercase, with uppercase U
U = 'A'.

?- char_type(a, to_upper(U)).     % U is the uppercase of a (any character)
U = 'A'.

?- char_type('.', to_lower(L)).
```

```

L = '.'.

?- char_type(X, to_upper('A')).    % generate: which characters uppercase to 'A'?
X = 'A' ;
X = a.

?- code_type(0'a, lower(U)).
U = 65.

% Practical example: Validate password
validate_password(Password) :-
    atom_chars(Password, Chars),
    length(Chars, Len),
    Len >= 8,
    member(Upper, Chars), char_type(Upper, upper),
    member(Lower, Chars), char_type(Lower, lower),
    member(Digit, Chars), char_type(Digit, digit),
    !.

% Extract specific character types
extract_digits(Text, Digits) :-
    atom_chars(Text, Chars),
    findall(D, (member(D, Chars), char_type(D, digit)), Digits).

?- extract_digits('abc123def456', D).
D = ['1', '2', '3', '4', '5', '6'].

% Remove non-alphabetic characters
clean_text(Input, Cleaned) :-
    atom_chars(Input, Chars),
    findall(C, (member(C, Chars), char_type(C, alpha)), CleanChars),
    atom_chars(Cleaned, CleanChars).

?- clean_text('hello123world!', C).
C = helloworld.

```

upcase_atom/2 and downcase_atom/2

Purpose: Convert atom case.

When to use: Use for normalization, case-insensitive comparisons, or formatting.

v4.4.0 (ISS-2025-0506): `instantiation_error` for an unbound argument and `type_error(atom, A)` for a non-atom, in place of a message atom.

```

% Convert to uppercase
?- upcase_atom(hello, Upper).
Upper = 'HELLO'.

?- upcase_atom('Hello World', Upper).
Upper = 'HELLO WORLD'.

% Convert to lowercase
?- downcase_atom('HELLO', Lower).
Lower = hello.

?- downcase_atom('Hello World', Lower).
Lower = 'hello world'.

% Practical example: Case-insensitive comparison
equal_ignore_case(Atom1, Atom2) :-
    downcase_atom(Atom1, Lower1),
    downcase_atom(Atom2, Lower2),
    Lower1 = Lower2.

```

```
?- equal_ignore_case('Hello', 'HELLO').
true.

?- equal_ignore_case('Hello', 'World').
false.

% Normalize input
normalize_command(Input, Command) :-
    downcase_atom(Input, Lower),
    atom_string(Lower, TrimmedString),
    atom_string(Command, TrimmedString).

% Create formatted headers
format_header(Text, Header) :-
    upcase_atom(Text, Upper),
    atom_concat('=== ', Upper, Temp),
    atom_concat(Temp, ' ===', Header).

?- format_header('section title', H).
H = '=== SECTION TITLE ==='.
```

11. DCG (Grammar) Predicates

Definite Clause Grammars provide a high-level notation for parsing and generating sequences.

Understanding DCGs

DCGs are a syntactic extension for writing parsers and generators:

- DCG rules use `-->` instead of `:-`
- Automatically handle list processing with difference lists
- Can embed Prolog goals using `{}`

```
% Simple DCG rule
noun --> [cat].
noun --> [dog].

% Is translated internally to:
% noun([cat|Rest], Rest).
% noun([dog|Rest], Rest).
```

phrase/2 and phrase/3

Purpose: Executes DCG rules for parsing or generation. `phrase(Body, List)` parses the whole list; `phrase(Body, List, Rest)` unifies `Rest` with the unparsed remainder.

When to use: Use to parse input with grammar rules or generate valid sequences.

v3.5.0: the first argument may be any DCG body, not just a non-terminal - `(A, B)`, `(A ; B)`, `(A -> B)`, `\+ A`, `!`, `{Goal}`, terminal lists `[a, b]` and `[]` are all translated correctly; a non-list second/third argument raises `type_error(list, _)` and a non-callable body raises `type_error(callable, _)` (previously these failed silently).

v3.10.0, engine v4 (ISS-2025-0451): `phrase/2, 3` is native - the grammar body is translated and the resulting goal is pushed onto the machine's goal stack instead of being solved by a nested recursive sub-solver. A DCG over a **1 000 000-token list** parses at the default JVM stack (the default engine raises `resource_error(stack_overflow)` on the same query even with `-Xss4m`), and the inference budget and the Stop interrupt now fire *inside* a parse. The ISO 13211-3 error clauses are unchanged.

v3.6.0: `phrase/3` with two free variables (e.g. `phrase(nt, [a|T], R)`) no longer raises a spurious `representation_error(cyclic_term)` on the default engine - it now answers with the expected var-var binding; real cyclic-term (rational-tree) protection is unaffected.

```
% Define a simple grammar
article --> [the].
article --> [a].
noun --> [cat].
noun --> [dog].
verb --> [chases].
verb --> [sees].

noun_phrase --> article, noun.
verb_phrase --> verb, noun_phrase.
sentence --> noun_phrase, verb_phrase.

% Parse with phrase/2
?- phrase(noun, [cat]).
true.

?- phrase(noun_phrase, [the, dog]).
true.
```

```

?- phrase(sentence, [the, cat, chases, a, dog]).
true.

% Generate valid sentences
?- phrase(sentence, S).
S = [the, cat, chases, the, cat] ;
S = [the, cat, chases, the, dog] ;
S = [the, cat, chases, a, cat] ;
...

% Practical example: Number parser
digit(D) --> [D], { char_type(D, digit) }.

digits([D|Ds]) --> digit(D), digits(Ds).
digits([D]) --> digit(D).

number(N) -->
    digits(Ds),
    { atom_chars(Atom, Ds), atom_number(Atom, N) }.

?- phrase(number(N), ['1', '2', '3']).
N = 123.

% phrase/3 returns the remainder; the body can be any DCG body (v3.5.0+)
?- phrase((article, noun), [the, dog, runs], Rest).
Rest = [runs].

```

enhanced_phrase/2 and enhanced_phrase/3

v4.6.0 (ISS-2025-0797, wave Q7): `enhanced_phrase/2,3` are `phrase/2,3` (v4 natives). The registry implementation never ran the grammar - it expanded the body and answered `true` - so an unbound or undefined body "succeeded"; now the grammar runs and the errors are `phrase/2,3`'s. **Purpose:** Enhanced DCG parsing with ISO/IEC DTS 13211-3 compliance.

When to use: Use for advanced parsing with complex grammars.

```

% Enhanced parsing with complex control structures
expression --> term, expression_rest.
expression_rest --> ['+'], term, expression_rest.
expression_rest --> ['-'], term, expression_rest.
expression_rest --> [].

term --> factor, term_rest.
term_rest --> ['*'], factor, term_rest.
term_rest --> ['/'], factor, term_rest.
term_rest --> [].

factor --> ['('], expression, [')'].
factor --> number.

% Parse arithmetic expression
?- enhanced_phrase(expression, ['2', '*', '(', '3', '+', '4', ')']).
true.

% With remainder extraction
?- enhanced_phrase(expression, ['2', '+', '3', ';', 'rest'], Rest).
Rest = [';', 'rest'].

% Practical example: Configuration parser
config_line -->
    key(K),
    spaces,
    ['='],
    spaces,

```

```

    value(V),
    { assertz(config(K, V)) }.

key(K) --> identifier(K).
value(V) --> quoted_string(V).
value(V) --> identifier(V).
value(V) --> number(V).

spaces --> [' '], spaces.
spaces --> [].

identifier(ID) -->
    [C],
    { char_type(C, alpha) },
    id_rest(Cs),
    { atom_chars(ID, [C|Cs]) }.

id_rest([C|Cs]) -->
    [C],
    { char_type(C, alnum) },
    id_rest(Cs).
id_rest([]) --> [].

```

phrase_with_options/4

Purpose: DCG parsing with advanced control options.

When to use: Use when you need parsing with error handling, debugging, or depth limits.

```

% Parse with options
Options = [
    syntax_errors(fail),    % How to handle syntax errors
    max_depth(100),        % Maximum recursion depth
    debug(true),           % Enable debug output
    trace(false)           % Disable tracing
].

% Safe parsing with depth limit
safe_parse(Rule, Input, Result) :-
    phrase_with_options(
        Rule,
        Input,
        [],
        [max_depth(1000), syntax_errors(fail)]
    ),
    Result = success.
safe_parse(_, _, failed).

% Practical example: JSON parser with error handling
json_value --> json_object.
json_value --> json_array.
json_value --> json_string.
json_value --> json_number.
json_value --> json_boolean.
json_value --> json_null.

parse_json(Input, Result) :-
    phrase_with_options(
        json_value,
        Input,
        [],
        [syntax_errors(error), max_depth(100)]
    ),
    Result = valid_json.

```

```
% Debug parsing issues
debug_parse(Grammar, Input) :-
    phrase_with_options(
        Grammar,
        Input,
        Rest,
        [debug(true), trace(true)]
    ),
    write('Remaining input: '), writeln(Rest).
```

call_dcg/3

Purpose: Calls DCG rules with explicit difference lists.

v4.5.0 (ISS-2025-0670): the body is translated by the v2 DCG translator (as `phrase/2, 3` and `consult` are), so `\+`, `!`, `{}`, `call//N` and pushback mean what they mean in a grammar rule.

When to use: Use for meta-programming with DCGs or custom parsing control.

```
% Direct DCG call with explicit lists
?- call_dcg(noun, [cat, runs], Rest).
Rest = [runs].

% Compose DCG rules dynamically
parse_sequence([], Input, Input).
parse_sequence([Rule|Rules], Input, Output) :-
    call_dcg(Rule, Input, Temp),
    parse_sequence(Rules, Temp, Output).

?- parse_sequence([article, noun, verb], [the, cat, runs], Rest).
Rest = [].

% Practical example: Dynamic grammar selection
parse_by_type(number, Input, Rest) :-
    call_dcg(number_parser, Input, Rest).
parse_by_type(word, Input, Rest) :-
    call_dcg(word_parser, Input, Rest).
parse_by_type(symbol, Input, Rest) :-
    call_dcg(symbol_parser, Input, Rest).

% Meta-DCG: Apply DCG rule multiple times
repeat_dcg(_, 0, Input, Input) :- !.
repeat_dcg(Rule, N, Input, Output) :-
    N > 0,
    call_dcg(Rule, Input, Temp),
    N1 is N - 1,
    repeat_dcg(Rule, N1, Temp, Output).

?- repeat_dcg(digit, 3, ['1', '2', '3', '4'], Rest).
Rest = ['4'].
```

dcg_translate_rule/2

Purpose: Translates DCG rules to standard Prolog clauses.

v4.5.0 (ISS-2025-0670): native, on the v2 translator; `dcg_translate_rule/4` (SWI's position-carrying form) is not provided - it is an unknown procedure (it used to raise a message atom).

When to use: Use for understanding DCG transformation or meta-programming.

```
% See how DCG rules are translated
?- dcg_translate_rule((noun --> [cat]), Clause).
Clause = (noun([cat|S], S)).
```

```

?- dcg_translate_rule((noun_phrase --> article, noun), Clause).
Clause = (noun_phrase(S0, S) :- article(S0, S1), noun(S1, S)).

% Complex translation with embedded goals
?- dcg_translate_rule(
    (number(N) --> digits(D), { atom_number(D, N) })),
    Clause
).
Clause = (number(N, S0, S) :-
    digits(D, S0, S),
    atom_number(D, N)
).

% Practical example: Generate parser predicates
generate_parser(Grammar, Predicates) :-
    findall(Pred,
        (member(Rule, Grammar),
         dcg_translate_rule(Rule, Pred)),
        Predicates).

% Analyze grammar complexity
analyze_dcg(Rule) :-
    dcg_translate_rule(Rule, Clause),
    Clause = (Head :- Body),
    functor(Head, Name, Arity),
    write('DCG rule: '), write(Name),
    write(' with arity '), write(Arity), nl,
    write('Translates to: '), write(Clause), nl.

```

v4.5.0 (ISS-2025-0571, wave P3.11): now the SAME translator consult uses (core.dcg.v2). The legacy one produced `\+(b, S0, S1), !(S0, S1)`, a wrong push-back and disconnected variables. A non-DCG argument is `type_error(dcg_rule, T)`.

```

?- dcg_translate_rule((a --> \+ b, !, [c]), C).
C = (a(S0, S) :- (\+ b(S0, _), S0 = S1), (!, S1 = S2), S2 = [c|S]).

```

expand_term/2

Purpose: `expand_term(+Term, -Expanded)`: if the user defines `term_expansion/2` and it succeeds on `Term`, its result; otherwise a `-->` rule is translated as consult does, and any other term is returned unchanged. (added v4.5.0, ISS-2025-0571)

Consult applies `term_expansion/2` (v4.5.0): when the user has `term_expansion/2` clauses, every clause read by consult (and every clause of a `.jpc` load) is passed through it first; the result may be one term or a list of terms (`[]` drops the clause).

```

term_expansion(double(X), [X, X]).
double(fact(1)).           % loads fact(1) twice

```

12. Exception Handling

Exception handling provides robust error management and recovery mechanisms.

throw/1

Purpose: Throws an exception to be caught by an enclosing catch/3.

When to use: Use to signal errors or exceptional conditions.

```
% Throw simple exception
divide(_, 0, _) :-
    throw(division_by_zero).
divide(X, Y, Result) :-
    Result is X / Y.

?- divide(10, 0, R).
ERROR: Uncaught exception: division_by_zero

% Throw ISO error terms
validate_positive(X) :-
    (   var(X)
    -> throw(error(instantiation_error, validate_positive/1))
    ;   \+ number(X)
    -> throw(error(type_error(number, X), validate_positive/1))
    ;   X <= 0
    -> throw(error(domain_error(positive_number, X), validate_positive/1))
    ;   true
    ).

% Practical example: File operations with error handling
read_file_safe(Filename, Content) :-
    (   \+ exists_file(Filename)
    -> throw(error(existence_error(file, Filename), read_file_safe/2))
    ;   \+ access_file(Filename, read)
    -> throw(error(permission_error(read, file, Filename), read_file_safe/2))
    ;   read_file_to_string(Filename, Content)
    ).

% Custom error types
process_data(Data) :-
    (   validate_format(Data)
    -> transform_data(Data)
    ;   throw(custom_error(invalid_format, Data))
    ).
```

catch/3

Purpose: Catches exceptions thrown by goals.

v4.5.0 (ISS-2025-0516): Recovery runs as `call(Recovery)` - a ! inside it is local and no longer cuts the clause that called catch/3.

When to use: Use to handle errors gracefully and implement recovery strategies.

```
% Syntax: catch(+Goal, +Catcher, +Recovery)

% Basic exception catching
safe_divide(X, Y, Result) :-
    catch(
        Result is X / Y,
        error(evaluation_error(zero_divisor), _),
        (write('Division by zero!'), nl, fail)
```

```

    ).

?- safe_divide(10, 0, R).
Division by zero!
false.

% Catch specific exceptions
process_with_retry(Data, Result) :-
    catch(
        process_data(Data, Result),
        network_error(timeout),
        (   writeln('Network timeout, retrying...'),
            sleep(1),
            process_with_retry(Data, Result)
        )
    ).

% Multiple exception types
robust_operation(Input, Output) :-
    catch(
        dangerous_operation(Input, Output),
        Error,
        handle_error(Error)
    ).

handle_error(error(type_error(_, _), _)) :-
    writeln('Type error: Invalid input type').
handle_error(error(domain_error(_, _), _)) :-
    writeln('Domain error: Value out of range').
handle_error(error(existence_error(_, _), _)) :-
    writeln('Existence error: Resource not found').
handle_error(Error) :-
    write('Unexpected error: '), writeln(Error).

% Practical example: Transaction with rollback
transaction(Actions, Result) :-
    catch(
        (   begin_transaction,
            perform_actions(Actions),
            commit_transaction,
            Result = success
        ),
        Error,
        (   rollback_transaction,
            write('Transaction failed: '), writeln(Error),
            Result = failed(Error)
        )
    ).

% Cleanup with exception safety
with_resource(Resource, Goal) :-
    acquire_resource(Resource),
    catch(
        call(Goal),
        Error,
        (release_resource(Resource), throw(Error))
    ),
    release_resource(Resource).

```

halt/0 and halt/1

Purpose: Terminates the Prolog processor. `halt/0` is equivalent to `halt(0)`; `halt(Code)` exits with the given exit code.

When to use: Use to exit the program, typically after fatal errors or completion.

v3.5.0: `halt/0`, `halt/1` actually terminate the processor - the CLI process exits with the given exit code, `:- halt` aborts a consult, and the IDE ends the run session gracefully (previously `halt` was effectively a no-op).

v4.5.0 (ISS-2025-0672): the engine never calls `System.exit` itself - an embedder receives a `PrologException` whose `isHalt()` is true (`getExitCode()` gives the code), and `catch/3` cannot intercept it. In **safe mode** (`Prolog.enableSafeMode()`) `halt/0,1` raise `permission_error(call, sandboxed, halt)` / `permission_error(call, sandboxed, halt(N))` instead (SWI's sandbox answer), so untrusted code cannot end the embedder's query; `enableSafeMode(new SafeModeOptions().allowHalt())` keeps them. The CLI's `--safe` allows them (the process belongs to the user). `initialization(G, main)` halts after `G` (ISS-2025-0636).

```
% Exit with success
?- halt(0).
% Program terminates with exit code 0

% Exit with error code
fatal_error(Message) :-
    write('FATAL ERROR: '), writeln(Message),
    halt(1).

% Practical example: Command-line application
main :-
    catch(
        run_application,
        Error,
        (   print_error(Error),
            halt(1)
        )
    ),
    halt(0).

run_application :-
    parse_arguments(Args),
    validate_arguments(Args),
    process_command(Args).

% Conditional exit
check_requirements :-
    (   missing_requirement(Req)
    -> write('Missing requirement: '), writeln(Req),
        halt(2)
    ;   true
    ).
```



```
% Check debug mode
is_debug_mode :-
    current_prolog_flag(debug, on).

debug_print(Message) :-
    ( is_debug_mode
    -> write('[DEBUG] '), writeln(Message)
    ; true
    ).
```

set_prolog_flag/2

Purpose: Sets system flags (where allowed).

When to use: Use to configure system behavior.

v3.5.0: the `unknown` flag is enforced - calling an undefined procedure raises `existence_error(procedure, Name/Arity)` when the flag is `error` (the default), prints a warning and fails when `warning`, and fails silently when `fail`. Procedures declared `dynamic` (via the `:- dynamic` directive, the `dynamic/1` goal, or implied by `assert/retractall`) fail silently instead of raising the error.

v3.9.0 (ISS-2025-0441): `occurs_check` accepts a **third** value, `error`, as ISO 7.11.2.4 requires (`true`, `false`, `error`) - it was rejected before. `error` makes a unification that would build a cyclic term raise `representation_error(cyclic_term)` instead of succeeding with a rational tree.

v3.8.0 (ISS-2025-0437): flags are **per engine**. `set_prolog_flag/2` now changes only the Prolog instance that runs the goal - previously the flag store was a process-wide static, so `set_prolog_flag(unknown, fail)` (or `double_quotes`, or `occurs_check`) in one engine silently reconfigured every other engine in the JVM. The same applies to `trace/0` / `notrace/0`. Embedders can reach a specific engine's store with `Prolog.getFlags()` and toggle tracing from another thread with `Prolog.setTracing(boolean)`.

v3.14.0 (ISS-2025-0472/0474/0477): the last of the process-global state follows. The **stream table**, the **operator store** (`op/3` / `current_op/3`), the **spy points** and the **profiler counters** belong to the Prolog instance, reached with `Prolog.getStreams()`, `Prolog.getOps()` and `Prolog.getEngineState()`. Two engines in one JVM no longer see each other's streams, aliases, operators, spy points or profile numbers. LIM-034 is closed.

v4.4.0 (ISS-2025-0508): ISO 8.17.1.3 - `instantiation_error`, `type_error(atom, F)`, `domain_error(prolog_flag, F)` for an unknown flag (setting an unknown flag no longer creates it), `permission_error(modify, flag, F)` for a read-only one and `domain_error(flag_value, F+V)` for a value the flag does not accept.

```
% Enable debug mode
?- set_prolog_flag(debug, on).
true.

% Set unknown predicate behavior
?- set_prolog_flag(unknown, fail). % Fail on undefined predicates
true.

% Practical example: Configure application
initialize_app :-
    set_prolog_flag(debug, off),
    set_prolog_flag(unknown, error),
    set_prolog_flag(double_quotes, codes),
    writeln('Application configured').

% Toggle debug mode
toggle_debug :-
    current_prolog_flag(debug, Current),
    ( Current = on
```

```

-> set_prolog_flag(debug, off),
    writeln('Debug mode disabled')
;   set_prolog_flag(debug, on),
    writeln('Debug mode enabled')
).

% Set optimization level
set_optimization(Level) :-
    (   member(Level, [0, 1, 2, 3])
    -> set_prolog_flag(optimize, Level)
    ;   writeln('Invalid optimization level')
    ).

```

statistics/2

v4.5.0 (ISS-2025-0608): native, SWI keys and shapes: runtime, walltime, real_time, system_time = [Total, SinceLast] (ms; seconds for real_time); cputime, thread_cputime, process_cputime, epoch = float seconds; inferences = the engine's resolution-step count; stack, stack_limit, localused, globalused, heapused, local, global, heap, trail, threads, garbage_collection, predicates, clauses, modules. The value is UNIFIED (statistics(runtime, [T|_]) works); an unknown key is domain_error(statistics_key, K). statistics/0 prints a summary to user_error. **Purpose:** Queries system statistics and performance metrics.

When to use: Use for performance monitoring and optimization.

```

% Query runtime
?- statistics(runtime, [Total, Since]).
Total = 1234,   % Total milliseconds since start
Since = 10.    % Milliseconds since last call

% Memory usage
?- statistics(memory, [Used, Free]).
Used = 5242880,   % Bytes used
Free = 10485760. % Bytes free

% Practical example: Measure execution time
time_goal(Goal, Time) :-
    statistics(runtime, [Start, _]),
    call(Goal),
    statistics(runtime, [End, _]),
    Time is End - Start.

?- time_goal(sleep(1), T).
T = 1000. % Milliseconds

% Performance profiling
profile(Goal) :-
    statistics(runtime, [T0, _]),
    statistics(memory, [M0, _]),
    call(Goal),
    statistics(runtime, [T1, _]),
    statistics(memory, [M1, _]),
    Time is T1 - T0,
    Memory is M1 - M0,
    format('Execution time: ~w ms~n', [Time]),
    format('Memory used: ~w bytes~n', [Memory]).

% Monitor resource usage
monitor_resources :-
    repeat,
    statistics(memory, [Used, Free]),
    Total is Used + Free,
    Percentage is (Used * 100) / Total,

```

```
format('Memory usage: ~2f%~n', [Percentage]),
sleep(5),
fail.
```

```
<!-- START_CHANGE: ISS-2025-0179 - Add leash/1 debug predicate -->
```

leash/1

Purpose: Controls which ports (call, exit, redo, fail) the debugger pauses at during tracing.

When to use: Use to fine-tune debugger behavior so it only stops at ports you are interested in.

```
% Syntax: leash(+Ports)
% Ports is a list of port names: call, exit, redo, fail
?- leash([call, fail]).
true.
% Debugger will now only pause at call and fail ports
```

```
<!-- END_CHANGE: ISS-2025-0179 -->
```

consult/1, [File|Files]

Purpose: Load Prolog source files from a program. (*added v4.5.0, ISS-2025-0574, wave P3.1*)

`consult(File)` accepts an atom or string, a `Dir/File` path term, `library(Name)` or a list of those; `[F1, F2]` as a goal is the same as consulting each. A relative name resolves against the directory of the file being loaded (else the working directory) and `.pl` is added when the name has no extension (then `.prolog`, then the bare name). Consulting a file that is already loaded **reloads** it: the user predicates it defined are wiped first (SWI). Directives run once, in the module being loaded, on the calling query's machine (same inference budget). Clause errors are printed as warnings on `user_error` and loading continues; a missing file is `existence_error(source_sink, F)`. `library(X)` loads a prelude module (lists, apply, pairs, coroutining, `clpfd`) or is a no-op for a library JProlog implements natively; *v4.6.0* (ISS-2025-0737): any other library, and any `Alias(Path)`, is searched through `file_search_path/2` (see below) - an existence error only when no readable file is found.

v4.6.0 (ISS-2025-0730): a reload removes only the clauses THIS file added (see `multifile/1`); a file that loads itself (directly or through its directives) is not loaded again. A clause with a module-qualified head (`user:portray(X) :- ..., m:foo(1).`) belongs to that module (ISS-2025-0733): `user:H` is an ordinary user clause, a clause for the module being loaded is an ordinary clause of it, and any other module is created on demand (`current_module/1` sees it) and holds the clause as the `M:H` clause `assertz(M:H)` also stores. Its body runs in the module the clause was written in (SWI).

The loader keeps a **load context**: when a file that declares `:- module(M, Exports)` has been loaded, the current (type-in) module is the one the load started in again and `M`'s exports are imported into it (*v4.5.0*, ISS-2025-0573: consulting two module files, or a module file then a plain file, no longer leaves the later clauses inside the first module).

These predicates live in `builtin.filesystem`, so `Prolog.enableSafeMode()` removes them (and `:- include/1` raises `permission_error` in safe mode). *v4.5.0* (ISS-2025-0625): `enableSafeMode(new SafeModeOptions().allowFileRead(dir))` keeps them, restricted to files inside `dir` (`permission_error(open, source_sink, F)` otherwise). *v4.5.0* (ISS-2025-0639): a thread started by a directive may consult while the loading thread is blocked in `thread_join/1, 2` on it (it used to deadlock on the engine's load lock). *v4.6.0* (ISS-2025-0739, wave Q3.8): the load lock is per FILE: loads of different files proceed in parallel on different threads (so a directive may also wait with `thread_get_message/1, 2` or `concurrent_*` for threads that load other files), the same file loaded concurrently waits for the first load, and a load that could never proceed - the file's loader waits (through file locks, joins or `concurrent_*` workers) for the thread that asks - raises `permission_error(load, source_sink, File)` instead of hanging. *v4.5.0* (ISS-2025-0636): `:- initialization(G, main)` in a file given to the CLI runs `G` after loading and ends the process (0; 1 on failure or error; N on `halt(N)`); an embedder's load runs `G` after the

load without halting.

```
?- consult('lib/utils'), [helpers, 'more/stuff.pl'].
?- consult(library(lists)).      % a prelude module
```

ensure_loaded/1

Purpose: Load a file unless it is already loaded (by absolute path); for an already loaded module file its exports are imported into the current module. `:- ensure_loaded(F)` in a file is a real load now (it was a silent no-op). (v4.5.0, ISS-2025-0574)

load_files/1, load_files/2

Purpose: `load_files(+Files, +Options)`: options `if(true)` (always, the default), `if(changed)` (unless loaded and unmodified), `if(not_loaded)` (as `ensure_loaded/1`) and `must_be_module(true)` (`domain_error(module_file, F)` for a non-module file); other options are accepted and ignored. (added v4.5.0, ISS-2025-0574)

make/0

Purpose: Reload every loaded source file modified since it was loaded. (added v4.5.0) v4.6.0 (ISS-2025-0736): a file is also reloaded when a file it `:- includes` changed.

use_module/1, use_module/2

Purpose: Load a module file once (like `ensure_loaded/1`) and import its exports into the current module. (as goals: v4.6.0, ISS-2025-0735; the directive forms since v4.5.0)

`use_module(File, Imports)` imports only what `Imports` names: a list of `Name/Arity`, `Name//Arity` and `PI as NewName` (the predicate is visible under the new name only), or `except(List)` (everything but the listed indicators; an `as` item there renames). Two restricted imports of one module accumulate; `use_module(File)` imports everything again. `library(X)` is a prelude module or a library JProlog implements natively; any other `library(X)` (and any `Alias(Path)`) is searched through `file_search_path/2`.

```
:- use_module(geometry, [area/2, perimeter/2 as perim]).
:- use_module(library(mylib), except([internal/1])).
```

multifile/1, discontinuous/1

Purpose: Declarations (also callable as goals). (v4.6.0, ISS-2025-0730, wave Q3.1)

`:- multifile p/1.` (or `M:p/1`, a comma sequence or a list) makes `p/1` **defined**: a call with no clause fails instead of raising `existence_error(SWI)`, and `predicate_property(p(_), multifile)` holds. Every consulted clause remembers the file it came from; **reconsulting a file removes only the clauses that file added** - clauses other files added to the same predicate always stay, and so do `assertz/1` -added clauses of a multifile predicate (those of any other predicate the file defines are reset with it). `:- discontinuous p/1.` silences the SWI warning `Clauses of p/1 are not together in the source-file` that a load now prints on `user_error` when the clauses of a predicate are interrupted by another predicate's (a warning, never an error).

Goal expansion at load time

v4.6.0 (ISS-2025-0731, wave Q3.2): when `user` (or the module being loaded) defines `goal_expansion/2`, every goal of a loaded clause body and of a directive is expanded - through `,/;->/*->/|, \+, call/1, once/1, ignore/1, forall/2, findall/3, 4, bagof/3, setof/3` (under $\wedge$), `aggregate_all/3`,

catch/3, setup_call_cleanup/3 and M:G - repeatedly until goal_expansion/2 fails or returns a variant of its input (the module's own hook first, then user's). 100 successive rewrites of one goal are a load error (resource_error(goal_expansion_depth)). Variables shared with the rest of the clause stay shared.

```
goal_expansion(dbl(X, Y), Y is X * 2).
p(A, B) :- dbl(A, B).           % stored as p(A, B) :- B is A*2.
```

file_search_path/2

Purpose: file_search_path(Alias, Dir): the user's search path for Alias(Path) file specifications (library(X), mine(foo)). Dynamic and multifile, without clauses in a fresh engine. Its entries are tried first; then the built-in defaults: swi = the JProlog home (-Djprolog.home, else ~/.jprolog), library = swi(library) and the directory of the bundled prelude when it is on disk, foreign = swi(lib). Dir may itself be Alias2(Sub). (v4.6.0, ISS-2025-0737, wave Q3.3)

```
?- assertz(file_search_path(mine, '/opt/prolog/lib')),
   consult(mine(utils)).           % /opt/prolog/lib/utils.pl
```

Prolog.runMain() (Java API)

v4.6.0 (ISS-2025-0736): an embedder that wants SWI's initialization(G, main) calls setDeferInitializationMain(true) before loading and then runMain(): the collected goals run in load order and the method returns the exit code the program would halt with (0 success, 1 failure or an uncaught error - reported on user_error -, N for halt(N), -1 when there is no main goal). Nothing halts the JVM.

include/1 (directive)

Purpose: :- include(File). reads the clauses of File in place, as if they were written there: same module, same load, relative to the including file. (added v4.5.0, ISS-2025-0575)

source_file/1, source_file/2

Purpose: source_file(?File) enumerates the files loaded by consult/1 & co. (absolute paths); source_file(?Head, ?File) relates a user predicate to the file that defined it. (added v4.5.0, ISS-2025-0576)

prolog_load_context/2

Purpose: prolog_load_context(?Key, ?Value) while a file is being loaded: module (the current module), file (the file being read - an included file is itself), source (the file the load started from), directory and dialect (swi). Fails outside a load. (added v4.5.0, ISS-2025-0576)

Compiled files (.jpc)

v4.5.0 (ISS-2025-0577, wave P3.6): Prolog.compile/compileFile read with the same v2 reader as consult and record every clause as read - directives included, DCG rules untranslated - and a .jpc load runs them through the consult clause handler, so a compiled program behaves exactly like the consulted source (r1(a ==> b), 'a' 'b', backquotes, e({}), :- dynamic, :- table, :- module, :- initialization all survive). Only op/3 directives are executed at compile time. A test consults every examples/*.pl both ways and compares the listings.

14. Cryptographic Predicates

Cryptographic predicates provide hashing, encoding, and secure random generation capabilities for data integrity, authentication, and security use cases.

md5_hash/2

Purpose: Computes the MD5 hash of an atom, returning a hex string.

When to use: Use for checksums and non-security fingerprinting of data.

```
% Syntax: md5_hash(+Input, -Hash)
?- md5_hash('hello world', Hash).
Hash = '5eb63bbbe01eeed093cb22bb8f5acdc3'.

% Verify data integrity
verify_integrity(Data, ExpectedHash) :-
    md5_hash(Data, ActualHash),
    ActualHash = ExpectedHash.
```

sha256_hash/2

Purpose: Computes the SHA-256 hash of an atom, returning a hex string.

When to use: Use for secure hashing where collision resistance is needed.

```
% Syntax: sha256_hash(+Input, -Hash)
?- sha256_hash('hello world', Hash).
Hash = 'b94d27b9934d3e08a52e52d7da7dabfac484efe37a5380ee9088f7ace2efcde9'.

% Hash a password (for demonstration; use proper KDF in production)
store_password(User, Password) :-
    sha256_hash(Password, Hash),
    assertz(user_hash(User, Hash)).
```

sha512_hash/2

Purpose: Computes the SHA-512 hash of an atom, returning a hex string.

When to use: Use when maximum hash length and security margin are required.

```
% Syntax: sha512_hash(+Input, -Hash)
?- sha512_hash('hello world', Hash).
Hash = '309ecc489c12d6eb4cc40f50c902f2b4d0ed77ee511a7c7a9bcd3ca86d4cd86f989dd35bc5ff499670da34255b45b0cfd830e81f605dcf7dc5542e93ae9cd76f'.
```

crypto_hash/3

Purpose: Computes a hash using a specified algorithm.

When to use: Use when you need flexibility in choosing the hash algorithm at runtime.

```
% Syntax: crypto_hash(+Algorithm, +Input, -Hash)
?- crypto_hash(sha256, 'test data', Hash).
Hash = '916f0027a575074ce72a331777c3478d6513f786a591bd892da1a577bf2335f9'.

?- crypto_hash(md5, 'test data', Hash).
Hash = 'eb733a00c0c9d336e65691a37ab54293'.

% Choose algorithm based on requirements
```

```
secure_hash(Input, Hash) :-
    crypto_hash(sha256, Input, Hash).
```

hmac/4

Purpose: Computes an HMAC (Hash-based Message Authentication Code) using a key and algorithm.

When to use: Use for message authentication and integrity verification with a shared secret.

```
% Syntax: hmac(+Algorithm, +Key, +Message, -MAC)
?- hmac(sha256, 'secret_key', 'my message', MAC).
MAC = 'a1b2c3d4e5f6...'.

% Verify a received message
verify_message(Key, Message, ExpectedMAC) :-
    hmac(sha256, Key, Message, ComputedMAC),
    ComputedMAC = ExpectedMAC.
```

base64_encode/2

Purpose: Encodes an atom to its Base64 representation.

When to use: Use for encoding binary or text data for safe transmission over text-based protocols.

```
% Syntax: base64_encode(+Plain, -Encoded)
?- base64_encode('Hello, World!', Encoded).
Encoded = 'SGVsbG8sIFdvcmxkIQ=='.

% Encode credentials for HTTP basic auth
basic_auth_header(User, Password, Header) :-
    atom_concat(User, ':', Temp),
    atom_concat(Temp, Password, Credentials),
    base64_encode(Credentials, Encoded),
    atom_concat('Basic ', Encoded, Header).
```

base64_decode/2

Purpose: Decodes a Base64-encoded atom back to its original form.

When to use: Use to decode data received in Base64 format.

```
% Syntax: base64_decode(+Encoded, -Decoded)
?- base64_decode('SGVsbG8sIFdvcmxkIQ==', Decoded).
Decoded = 'Hello, World!'.

% Round-trip encoding
?- base64_encode('test', E), base64_decode(E, D).
E = 'dGVzdA==', D = 'test'.
```

uuid/1

Purpose: Generates a new UUID (version 4, random).

When to use: Use to create unique identifiers for records, sessions, or transactions.

```
% Syntax: uuid(-UUID)
?- uuid(Id).
Id = '550e8400-e29b-41d4-a716-446655440000'.

% Create a unique record
create_record(Name, Record) :-
    uuid(Id),
```

```
Record = record(Id, Name).
```

random_token/2

Purpose: Generates a cryptographically secure random hex token of the specified byte length.

When to use: Use for generating session tokens, API keys, or nonces.

```
% Syntax: random_token(+ByteLength, -Token)
?- random_token(16, Token).
Token = 'a3f2b8c1d4e5f6a7b8c9d0e1f2a3b4c5'.

% Generate an API key
generate_api_key(Key) :-
    random_token(32, Key).
```

crypto_random_int/3

Purpose: Generates a cryptographically secure random integer within a range.

When to use: Use when you need unbiased, unpredictable random numbers for security-sensitive applications.

```
% Syntax: crypto_random_int(+Low, +High, -Value)
% Generates a random integer N where Low <= N < High
?- crypto_random_int(1, 100, N).
N = 42.

% Generate a 6-digit OTP
generate_otp(OTP) :-
    crypto_random_int(100000, 1000000, OTP).
```

```
<!-- START_CHANGE: ISS-2025-0179 - Add AES and password hashing predicates -->
```

crypto_aes_encrypt/4

Purpose: Encrypts plaintext using AES symmetric encryption.

```
% Syntax: crypto_aes_encrypt(+Key, +Plaintext, +Options, -Ciphertext)
?- crypto_aes_encrypt('my_secret_key', 'hello world', [], Cipher).
```

crypto_aes_decrypt/4

Purpose: Decrypts AES-encrypted ciphertext back to plaintext.

```
% Syntax: crypto_aes_decrypt(+Key, +Ciphertext, +Options, -Plaintext)
?- crypto_aes_decrypt('my_secret_key', Cipher, [], Plain).
```

crypto_hash_password/2

Purpose: Hashes a password using PBKDF2 with a random salt for secure storage.

```
% Syntax: crypto_hash_password(+Password, -Hash)
?- crypto_hash_password('my_password', Hash).
Hash = 'pbkdf2:sha256:...'.

```

crypto_verify_password/2

Purpose: Verifies a password against a previously hashed value.

```
% Syntax: crypto_verify_password(+Password, +Hash)
?- crypto_hash_password('secret', H), crypto_verify_password('secret', H).
true.
```

<!-- END_CHANGE: ISS-2025-0179 -->

15. JSON Predicates

JSON predicates provide parsing, serialization, and manipulation of JSON data, enabling integration with web services and configuration files.

json_parse/2

Purpose: Parses a JSON string into a Prolog term representation.

When to use: Use to convert JSON data received from external sources into Prolog terms for processing.

```
% Syntax: json_parse(+JsonAtom, -Term)
?- json_parse('{"name":"John","age":30}', Term).
Term = json([name='John', age=30]).

?- json_parse('[1, 2, 3]', Term).
Term = [1, 2, 3].

% Parse and extract data
process_json(JsonString, Name) :-
    json_parse(JsonString, json(Pairs)),
    member(name=Name, Pairs).
```

json_serialize/2

Purpose: Serializes a Prolog term into a JSON string.

When to use: Use to produce JSON output for APIs, files, or inter-process communication.

```
% Syntax: json_serialize(+Term, -JsonAtom)
?- json_serialize(json([name='John', age=30]), Json).
Json = '{"name":"John","age":30}'.

?- json_serialize([1, 2, 3], Json).
Json = '[1,2,3]'.

% Build and serialize a response
build_response(Status, Message, Json) :-
    json_serialize(json([status=Status, message=Message]), Json).

?- build_response(ok, 'Operation completed', Json).
Json = '{"status":"ok","message":"Operation completed"}'.
```

json_get/3

Purpose: Retrieves a value from a parsed JSON term by key.

When to use: Use to access specific fields from a JSON object without manual member lookup.

```
% Syntax: json_get(+JsonTerm, +Key, -Value)
?- json_parse('{"name":"John","age":30}', J), json_get(J, name, V).
V = 'John'.

% Nested access with dot notation
?- json_parse('{"user":{"name":"John"}}', J), json_get(J, 'user.name', V).
V = 'John'.
```

json_set/4

Purpose: Sets or updates a key-value pair in a JSON term, producing a new term.

When to use: Use to modify JSON data immutably before serialization.

```
% Syntax: json_set(+JsonTerm, +Key, +Value, -NewJsonTerm)
?- json_parse('{ "name": "John", "age": 30}', J),
   json_set(J, age, 31, J2),
   json_serialize(J2, Out).
Out = '{ "name": "John", "age": 31}'.
```

json_keys/2

Purpose: Retrieves all keys from a JSON object as a list.

When to use: Use to enumerate the fields of a JSON object.

```
% Syntax: json_keys(+JsonTerm, -Keys)
?- json_parse('{ "name": "John", "age": 30, "city": "NYC"}', J), json_keys(J, Keys).
Keys = [name, age, city].
```

json_member/3

Purpose: Non-deterministically enumerates key-value pairs in a JSON object.

When to use: Use to iterate over all fields of a JSON object via backtracking.

```
% Syntax: json_member(+JsonTerm, ?Key, ?Value)
?- json_parse('{ "a": 1, "b": 2}', J), json_member(J, K, V).
K = a, V = 1 ;
K = b, V = 2.

% Find keys with numeric values
numeric_keys(Json, Keys) :-
    findall(K, (json_member(Json, K, V), number(V)), Keys).
```

16. Date/Time Predicates

Date and time predicates provide access to the system clock and operations for date arithmetic, formatting, and parsing.

get_time/1

v4.5.0 (ISS-2025-0609): a FLOAT number of seconds since the epoch (SWI; it was integer milliseconds). `format_time/3` and `parse_time/3` use seconds too; SWI's `format_time(+Out, +Format, +Stamp)` with % strftime directives is accepted next to the historical `format_time(+Pattern, +Stamp, -Atom)`; `stamp_date_time(+Stamp, -date(Y,M,D,H,Mn,S,Off,TZ,DST), +TimeZone)` (local, 'UTC' or an offset in seconds west of UTC) and `date_time_stamp(+Date, -Stamp)` are new. **Purpose** : Returns the current time as a Unix timestamp (seconds since epoch) as a float.

When to use: Use for precise timing, benchmarking, or recording event timestamps.

```
% Syntax: get_time(-Timestamp)
?- get_time(T).
T = 1.7110592e+09.

% Measure elapsed time
elapsed(Goal, Seconds) :-
    get_time(T0),
    call(Goal),
    get_time(T1),
    Seconds is T1 - T0.
```

now/1

Purpose: Returns the current date and time as a `datetime(Y,M,D,H,Min,S)` term.

When to use: Use when you need the current date and time as structured components.

```
% Syntax: now(-DateTime)
?- now(DT).
DT = datetime(2026, 3, 21, 14, 30, 45).
```

today/1

Purpose: Returns the current date as a `date(Y,M,D)` term.

When to use: Use when you only need the current date without time components.

```
% Syntax: today(-Date)
?- today(D).
D = date(2026, 3, 21).
```

date_add/4

Purpose: Adds a duration to a date, producing a new date.

When to use: Use for date arithmetic such as computing deadlines or future dates.

```
% Syntax: date_add(+Date, +Duration, +Unit, -NewDate)
% Unit is one of: days, months, years
?- date_add(date(2026, 3, 21), 10, days, NewDate).
NewDate = date(2026, 3, 31).

?- date_add(date(2026, 3, 21), 2, months, NewDate).
```

```
NewDate = date(2026, 5, 21).

% Calculate a deadline
deadline(StartDate, DaysAllowed, Deadline) :-
    date_add(StartDate, DaysAllowed, days, Deadline).
```

date_diff/4

Purpose: Computes the difference between two dates in the specified unit.

When to use: Use to determine the number of days, months, or years between two dates.

```
% Syntax: date_diff(+Date1, +Date2, +Unit, -Difference)
?- date_diff(date(2026, 1, 1), date(2026, 3, 21), days, Diff).
Diff = 79.

% Calculate age in years
age(BirthDate, Age) :-
    today(Today),
    date_diff(BirthDate, Today, years, Age).
```

day_of_week/2

Purpose: Determines the day of the week for a given date.

When to use: Use for scheduling or display purposes.

```
% Syntax: day_of_week(+Date, -DayName)
?- day_of_week(date(2026, 3, 21), Day).
Day = saturday.

% Check if a date is a weekend
is_weekend(Date) :-
    day_of_week(Date, Day),
    member(Day, [saturday, sunday]).
```

date_parts/4

Purpose: Decomposes a date term into its year, month, and day components.

When to use: Use to extract individual components from a date term.

```
% Syntax: date_parts(+Date, -Year, -Month, -Day)
?- date_parts(date(2026, 3, 21), Y, M, D).
Y = 2026, M = 3, D = 21.

% Extract just the year
get_year(Date, Year) :-
    date_parts(Date, Year, _, _).
```

time_parts/4

Purpose: Decomposes a datetime term into its hour, minute, and second components.

When to use: Use to extract time components from a datetime term.

```
% Syntax: time_parts(+DateTime, -Hour, -Minute, -Second)
?- now(DT), time_parts(DT, H, M, S).
H = 14, M = 30, S = 45.
```

Formatting and parsing dates

`format_date/3` and `parse_date/3` used to be documented here but were never implemented (the 4.6 documentation sweep, ISS-2025-0718, removed them). Use `format_time/3` - strftime-style directives, e.g. `format_time(atom(A), '%Y-%m-%d', Stamp)` - and `parse_time/3` instead.

17. Filesystem Predicates

Filesystem predicates provide operations for querying and manipulating files and directories on the host system.

file_exists/1

Purpose: Checks whether a file exists at the given path.

When to use: Use to verify a file is present before attempting to read or process it.

```
% Syntax: file_exists(+Path)
?- file_exists('data.pl').
true.

?- file_exists('nonexistent.pl').
false.

% Safe file loading
safe_consult(File) :-
    (   file_exists(File)
    -> consult(File)
    ;   write('File not found: '), writeln(File)
    ).
```

directory_exists/1

Purpose: Checks whether a directory exists at the given path.

When to use: Use to verify a directory is present before listing or writing into it.

```
% Syntax: directory_exists(+Path)
?- directory_exists('/tmp').
true.

?- directory_exists('/nonexistent').
false.
```

exists_file/1, exists_directory/1

v4.6.0 (ISS-2025-0745). SWI's names for `file_exists/1` and `directory_exists/1`. Relative paths (here and in every filesystem predicate) resolve against the engine's working directory (`working_directory/2`).

delete_file/1

Purpose: Deletes a file from the filesystem.

When to use: Use to clean up temporary files or remove obsolete data.

```
% Syntax: delete_file(+Path)
?- delete_file('temp_output.txt').
true.
```

rename_file/2

Purpose: Renames or moves a file.

When to use: Use to rename files or move them between directories.

```
% Syntax: rename_file(+OldPath, +NewPath)
?- rename_file('old_name.pl', 'new_name.pl').
true.
```

copy_file/2

Purpose: Copies a file to a new location.

When to use: Use to create backups or duplicate files.

```
% Syntax: copy_file(+Source, +Destination)
?- copy_file('data.pl', 'data_backup.pl').
true.

% Backup before modifying
safe_modify(File) :-
    atom_concat(File, '.bak', Backup),
    copy_file(File, Backup).
```

file_size/2

Purpose: Returns the size of a file in bytes.

When to use: Use to check file sizes before processing or to report storage usage.

```
% Syntax: file_size(+Path, -Size)
?- file_size('data.pl', Size).
Size = 4096.
```

read_file_to_atom/2

Purpose: Reads the entire contents of a file into an atom.

When to use: Use for reading small files or configuration data in one operation.

```
% Syntax: read_file_to_atom(+Path, -Content)
?- read_file_to_atom('config.txt', Content).
Content = 'key1=value1\nkey2=value2\n'.
```

write_atom_to_file/2

Purpose: Writes an atom as the entire contents of a file, overwriting any existing content.

When to use: Use for writing small files or saving generated output.

```
% Syntax: write_atom_to_file(+Content, +Path)
?- write_atom_to_file('Hello, World!\n', 'output.txt').
true.
```

make_directory/1

Purpose: Creates a single directory.

When to use: Use when you need to create exactly one directory level.

```
% Syntax: make_directory(+Path)
?- make_directory('output').
true.
```

make_directory_path/1

Purpose: Creates a directory and all necessary parent directories.

When to use: Use when the full path may not exist and intermediate directories need to be created.

```
% Syntax: make_directory_path(+Path)
?- make_directory_path('output/reports/2026').
true.
```

absolute_file_name/2, absolute_file_name/3

Purpose: Resolves a relative or aliased file path to an absolute path.

v4.6.0 (ISS-2025-0737, wave Q3.3): `absolute_file_name(+Spec, -Abs, +Options)` (SWI). `Spec` is a name (relative to `relative_to(Dir)`, else the directory of the file being loaded, else the working directory) or `Alias(Path)` searched through `file_search_path/2`. Options: `extensions(List)`, `file_type(txt|prolog|source|directory|regular|executable|qlf)` (prolog and source try `.pl`, `.prolog`, then no extension), `access(read|write|append|execute|exist|none)`, `relative_to(FileOrDir)`, `solutions(first|all)` and `file_errors(error|fail)`. With no access, type or extension condition the first candidate is the answer; otherwise the first candidate that meets them, or `existence_error(source_sink, Spec)` (fail with `file_errors(fail)`).

```
?- absolute_file_name(library(lists), F, [file_type(prolog), access(read)]).
F = '../prelude/lists.pl'.
```

When to use: Use to normalize file paths for consistent handling.

```
% Syntax: absolute_file_name(+Relative, -Absolute)
?- absolute_file_name('data.pl', Abs).
Abs = '/home/user/project/data.pl'.
```

directory_files/2

Purpose: Lists all files and directories in a given directory.

When to use: Use to enumerate directory contents for batch processing.

```
% Syntax: directory_files(+Directory, -Files)
?- directory_files('.', Files).
Files = ['file1.pl', 'file2.pl', 'subdir'].

% Process all Prolog files in a directory
consult_all(Dir) :-
    directory_files(Dir, Files),
    member(F, Files),
    file_extension(F, pl),
    atom_concat(Dir, '/', Temp),
    atom_concat(Temp, F, Path),
    consult(Path),
    fail ; true.
```

file_extension/2

Purpose: Extracts the file extension from a filename.

v4.6 (ISS-2025-0718): `file_extension/2`, `file_base_name/2` and `file_directory_name/2` exist now (they were documented but missing). They are pure text operations on the path (no file-system access, so they stay available in safe mode); an atom gives an atom, a string a string; a file without an extension gives `''`; `file_directory_name('data.pl', D)` gives `'.'`.

When to use: Use to filter files by type or determine how to process a file.

```
% Syntax: file_extension(+FileName, -Extension)
?- file_extension('report.pdf', Ext).
Ext = pdf.

?- file_extension('archive.tar.gz', Ext).
Ext = gz.
```

file_base_name/2

Purpose: Extracts the base filename (without directory path) from a full path.

When to use: Use to get just the filename from a full path.

```
% Syntax: file_base_name(+Path, -BaseName)
?- file_base_name('/home/user/data.pl', Base).
Base = 'data.pl'.
```

file_directory_name/2

Purpose: Extracts the directory portion from a full file path.

When to use: Use to determine which directory a file resides in.

```
% Syntax: file_directory_name(+Path, -Directory)
?- file_directory_name('/home/user/data.pl', Dir).
Dir = '/home/user'.
```

18. OS Predicates

OS predicates provide access to the operating system environment, process information, and shell command execution.

getenv/2

Purpose: Retrieves the value of an environment variable.

When to use: Use to read configuration from environment variables.

```
% Syntax: getenv(+VarName, -Value)
?- getenv('HOME', Home).
Home = '/home/user'.

?- getenv('PATH', Path).
Path = '/usr/local/bin:/usr/bin:/bin'.

% Configuration from environment
db_host(Host) :-
    (   getenv('DB_HOST', Host)
    -> true
    ;   Host = localhost
    ).
```

setenv/2

Purpose: Sets the value of an environment variable.

When to use: Use to configure the environment for child processes or shell commands.

v4.6 (ISS-2025-0718): implemented (it was documented but missing). A JVM cannot change its own process environment, so the value is a process-wide overlay: `getenv/2` sees it and the child processes of `shell/1,2` and `shell_output/3` inherit it. Removed in safe mode.

```
% Syntax: setenv(+VarName, +Value)
?- setenv('MY_APP_MODE', 'production').
true.

?- setenv('MY_APP_MODE', production), getenv('MY_APP_MODE', V), shell('test "$MY_APP_MODE"
    = production', C).
V = production, C = 0.
```

unsetenv/1

Purpose: `unsetenv(+VarName)` removes a variable from the overlay of `setenv/2` (and hides the JVM's own value from `getenv/2` and from child processes). (v4.6, ISS-2025-0718)

```
?- unsetenv('MY_APP_MODE'), \+ getenv('MY_APP_MODE', _).
true.
```

system_time/1

Purpose: Returns the system time in milliseconds since epoch as an integer.

When to use: Use for timestamps, profiling, or unique identifier generation.

```
% Syntax: system_time(-Millis)
?- system_time(T).
T = 1711059200000.
```

os_name/1

Purpose: Returns the name of the operating system.

When to use: Use to adapt behavior based on the host platform.

```
% Syntax: os_name(-Name)
?- os_name(OS).
OS = 'Linux'.

% Platform-specific path separator
path_separator(Sep) :-
    os_name(OS),
    ( OS = 'Windows' -> Sep = '\\\'; Sep = '/' ).
```

cpu_count/1

Purpose: Returns the number of available CPU cores.

When to use: Use to size thread pools or determine parallelism levels.

```
% Syntax: cpu_count(-Count)
?- cpu_count(N).
N = 8.
```

free_memory/1

Purpose: Returns the amount of free memory available to the JVM in bytes.

When to use: Use to monitor resource consumption or decide whether to proceed with memory-intensive tasks.

```
% Syntax: free_memory(-Bytes)
?- free_memory(M).
M = 268435456.
```

total_memory/1

Purpose: Returns the total memory allocated to the JVM in bytes.

When to use: Use together with free_memory/1 for memory utilization reporting.

```
% Syntax: total_memory(-Bytes)
?- total_memory(T), free_memory(F), Used is T - F.
T = 536870912, F = 268435456, Used = 268435456.
```

pid/1

Purpose: Returns the process ID of the current JVM process.

When to use: Use for logging, lock files, or process identification.

```
% Syntax: pid(-PID)
?- pid(P).
P = 12345.
```

hostname/1

Purpose: Returns the hostname of the machine.

When to use: Use for logging or identifying which host a program is running on.

```
% Syntax: hostname(-Name)
?- hostname(H).
H = 'myserver.example.com'.
```

shell/1

Purpose: Executes a shell command and succeeds if the command exits with status 0.

When to use: Use to run system commands where you only need to know success or failure.

```
% Syntax: shell(+Command)
?- shell('ls /tmp').
true.

?- shell('test -f config.ini').
true. % File exists
```

shell/2

Purpose: Executes a shell command and unifies the second argument with the exit code.

v4.6 (ISS-2025-0718): it was registered only under the name shell2/2; shell/2 now exists (shell2/2 is kept as an alias).

When to use: Use when you need to inspect the exit status of a command.

```
% Syntax: shell(+Command, -ExitCode)
?- shell('grep -q pattern file.txt', Code).
Code = 0. % Pattern found

?- shell('false', Code).
Code = 1.
```

shell_output/3

Purpose: Executes a shell command, capturing its standard output and exit code.

When to use: Use when you need to process the output of a system command in Prolog.

```
% Syntax: shell_output(+Command, -Output, -ExitCode)
?- shell_output('date +%Y', Output, Code).
Output = '2026\n', Code = 0.

?- shell_output('whoami', User, _).
User = 'prolog_user\n'.
```

19. Regex Predicates

Regex predicates provide regular expression matching, substitution, and splitting, powered by Java's regex engine.

re_match/2

Purpose: Tests whether a string matches a regular expression.

When to use: Use for pattern validation such as checking email formats, identifiers, or input constraints.

```
% Syntax: re_match(+Pattern, +String)
?- re_match('[0-9]+', '12345').
true.

?- re_match('^[a-z]+$', 'Hello').
false.

% Validate an email address (simplified)
valid_email(Email) :-
    re_match('^[^@]+@[^@]+\.[^@]+$', Email).

?- valid_email('user@example.com').
true.
```

re_matchsub/3

Purpose: Matches a regex with capturing groups and returns a list of captured substrings.

When to use: Use to extract parts of a string that match specific sub-patterns.

```
% Syntax: re_matchsub(+Pattern, +String, -Captures)
?- re_matchsub('(\d{4})-(\d{2})-(\d{2})', '2026-03-21', Caps).
Caps = ['2026-03-21', '2026', '03', '21'].

% Parse a log line
parse_log(Line, Level, Message) :-
    re_matchsub('\s*([^\s]+)\s+(.*)', Line, [_, Level, Message]).

?- parse_log('[ERROR] Connection failed', Level, Msg).
Level = 'ERROR', Msg = 'Connection failed'.
```

re_replace/4

Purpose: Replaces occurrences of a pattern in a string with a replacement.

When to use: Use for text transformation, sanitization, or templating.

```
% Syntax: re_replace(+Pattern, +Replacement, +Input, -Output)
?- re_replace('[0-9]+', 'NUM', 'Order 123 has 4 items', Out).
Out = 'Order NUM has NUM items'.

% Sanitize input
sanitize(Input, Clean) :-
    re_replace('[^a-zA-Z0-9 ]', '', Input, Clean).

?- sanitize('Hello <World>!', C).
C = 'Hello World'.
```

re_split/3

Purpose: Splits a string by a regex pattern into a list of substrings.

When to use: Use to tokenize or break apart strings using flexible delimiters.

```
% Syntax: re_split(+Pattern, +String, -Parts)
?- re_split('\\\\s+', 'a, b, c, d', Parts).
Parts = ['a', 'b', 'c', 'd'].

?- re_split('\\\\s+', 'Hello World Test', Words).
Words = ['Hello', 'World', 'Test'].
```

re_findall/3

Purpose: Finds all non-overlapping matches of a pattern in a string.

When to use: Use to extract all occurrences of a pattern from text.

```
% Syntax: re_findall(+Pattern, +String, -Matches)
?- re_findall('[0-9]+', 'Order 123 has 4 items at $56', Nums).
Nums = ['123', '4', '56'].

% Extract all email addresses from text
extract_emails(Text, Emails) :-
    re_findall('[\\w.]+@[\\w.]+', Text, Emails).
```

<!-- START_CHANGE: ISS-2025-0179 - Add re_escape/2 -->

re_escape/2

Purpose: Escapes special regex characters in a string so it can be used as a literal pattern.

```
% Syntax: re_escape(+Input, -Escaped)
?- re_escape('hello.world', E).
E = 'hello\\.world'.

% Use to safely match user-provided literal text
safe_match(Literal, Text) :-
    re_escape(Literal, Pattern),
    re_match(Pattern, Text).
```

<!-- END_CHANGE: ISS-2025-0179 -->

20. XML Predicates

XML predicates provide parsing, serialization, and querying of XML documents.

xml_parse/2

Purpose: Parses an XML string into a Prolog term representation.

When to use: Use to process XML data from files, APIs, or configuration sources.

```
% Syntax: xml_parse(+XmlAtom, -Term)
?- xml_parse('<person name="John"><age>30</age></person>', Term).
Term = element(person, [name='John'], [element(age, [], ['30'])]).

% Parse and inspect structure
get_root_tag(Xml, Tag) :-
    xml_parse(Xml, element(Tag, _, _)).
```

xml_serialize/2

Purpose: Serializes a Prolog XML term into an XML string.

When to use: Use to generate XML output from Prolog data structures.

```
% Syntax: xml_serialize(+Term, -XmlAtom)
?- xml_serialize(element(item, [id='1'], ['Hello']), Xml).
Xml = '<item id="1">Hello</item>'.

% Build an XML document
build_person_xml(Name, Age, Xml) :-
    AgeAtom = element(age, [], [Age]),
    Person = element(person, [name=Name], [AgeAtom]),
    xml_serialize(Person, Xml).
```

xpath/3

Purpose: Queries an XML term using an XPath-like expression, returning matching elements.

When to use: Use to extract specific elements from parsed XML data.

```
% Syntax: xpath(+XmlTerm, +Path, -Result)
?- xml_parse('<root><item>A</item><item>B</item></root>', Doc),
    xpath(Doc, '//item', Item).
Item = element(item, [], ['A']) ;
Item = element(item, [], ['B']).

% Extract attribute values
get_attribute(Xml, Path, Attr, Value) :-
    xpath(Xml, Path, element(_, Attrs, _)),
    member(Attr=Value, Attrs).
```

21. Threading Predicates

Threading predicates enable concurrent execution with message passing, built on Java's threading model. Use these for parallelism and background task processing.

Since v4.0.0 (ISS-2025-0479) a thread really runs its goal, on its own resolution machine over the same engine. Before that, `thread_create/2` started a thread that slept briefly and recorded a synthetic status; the goal never ran and it had to be an atom. What a worker shares with its creator: the clause store (so `assertz/retract` are visible both ways), the flags, the operator table and the module system. What it does not share: the goal's variables - the goal is copy_term'd, so **a binding made by a worker never appears in the creator's query** - its current input/output streams, and its global variables (v4.5.0: `nb_setval/2` in a worker is invisible to its creator, as in SWI). v4.5.0: the inference budget is ONE pool for the whole query - workers draw from their creator's budget instead of each getting a full copy. A thread id is an integer; an `alias(Name)` option gives it a name usable wherever an id is. Every thread that is not a `thread_create/2,3` worker is the thread main (id 1).

Sandbox: `Prolog.enableSafeMode()` removes every predicate in this section and in section 40 (concurrent execution) - a JVM thread is a host resource.

Errors (v4.5.0, ISS-2025-0621): ISO terms - `existence_error(thread, Id)`, `existence_error(message_queue, Q)`, `existence_error(mutex, M)`, `permission_error(create, thread, Alias)` for an alias in use, `permission_error(join, thread, Id)` for a detached thread, `oneself` or `main`, plus the usual instantiation and type errors. `thread_join/2` and `thread_get_message/1,2` wait without a time limit (they used to give up after 60 s / 30 s); a Stop still interrupts them.

thread_create/2

Purpose: Creates a new thread that executes a given goal.

When to use: Use to run a goal concurrently in the background.

```
% Syntax: thread_create(:Goal, -ThreadId)
?- thread_create(long_computation(Result), TId).
TId = 1.

% Start a background task
start_worker(Id) :-
    thread_create(worker_loop, Id).

% The goal really runs, and its database writes are visible afterwards
?- thread_create(assertz(done(yes)), T), thread_join(T, true), done(X).
X = yes.

% ... but its BINDINGS are not: the goal is copied
?- thread_create(X = 1, T), thread_join(T, true), var(X).
true.
```

thread_create/3

Purpose: Creates a thread with options.

When to use: Use when the thread needs a name (so other threads can send it messages) or must clean itself up without being joined.

```
% Syntax: thread_create(:Goal, -ThreadId, +Options)
% Options: alias(Name)      - a name usable wherever a thread id is
%           detached(Bool)  - true: the thread frees itself on completion and cannot be jo
ined
% Any other option is accepted and ignored.
```

```
?- thread_create(worker_loop, _, [alias(logger)]).
true.

?- thread_create(cleanup_task, T, [detached(true)]).
T = 4.
```

thread_join/2

Purpose: Waits for a thread to complete and unifies with its exit status.

When to use: Use to synchronize with a thread and retrieve its result.

The status is `true` when the goal succeeded, `false` when it failed, `exception(Ball)` when it threw (including `exception(error(resource_error(inference_limit), _))` when it exhausted the engine's inference budget), and `cancelled` when the thread was interrupted. The thread argument may be an id or an alias. A detached thread cannot be joined.

```
% Syntax: thread_join(+ThreadOrAlias, -Status)
?- thread_create(member(X, [a, b, c]), TId),
   thread_join(TId, Status).
Status = true.

?- thread_create(fail, T), thread_join(T, S).
S = false.

?- thread_create(throw(oops), T), thread_join(T, S).
S = exception(oops).

% Wait for a computation
run_and_wait(Goal, Status) :-
    thread_create(Goal, TId),
    thread_join(TId, Status).
```

thread_detach/1

Purpose: Detaches a thread so its resources are freed automatically upon completion.

When to use: Use for fire-and-forget tasks where you do not need the result.

```
% Syntax: thread_detach(+ThreadId)
?- thread_create(log_event(startup), TId), thread_detach(TId).
true.
```

thread_self/1

Purpose: Returns the identifier of the currently executing thread.

When to use: Use for logging or when threads need to identify themselves.

```
% Syntax: thread_self(-ThreadId)
?- thread_self(Id).
Id = main.
```

A thread that has an **alias** reports the alias, as in SWI-Prolog: the top-level thread is `main`, and a worker started with `thread_create(Goal, Id, [alias(w1)])` reports `w1`. A worker with no alias reports its integer Prolog id. Either form is accepted wherever a thread is named (`thread_join/2`, `thread_send_message/2`, `thread_is_alive/1`, ...).

thread_sleep/1

Purpose: Suspends the current thread for the specified number of seconds.

When to use: Use for delays, polling intervals, or rate limiting.

```
% Syntax: thread_sleep(+Seconds)
?- thread_sleep(2).
true. % Resumes after 2 seconds

% Retry with delay
retry_with_delay(Goal, Retries) :-
    Retries > 0,
    (   call(Goal) -> true
    ;   thread_sleep(1),
        R1 is Retries - 1,
        retry_with_delay(Goal, R1)
    ).
```

thread_is_alive/1

Purpose: Checks whether a thread is still running.

When to use: Use to poll thread status without blocking.

```
% Syntax: thread_is_alive(+ThreadId)
?- thread_create(thread_sleep(10), TId), thread_is_alive(TId).
true.
```

message_queue_create/1

Purpose: Creates a new message queue for inter-thread communication.

When to use: Use to set up communication channels between threads.

Since v4.0.0 a queue carries **arbitrary Prolog terms**, not just atoms, and every message is copied on the way in and on the way out so no variable is shared between the two threads. Every thread created by `thread_create/2,3` also owns a queue: `thread_send_message/2` accepts a queue id, a thread id or a thread alias, and `thread_get_message/1` reads the calling thread's own queue.

```
% Syntax: message_queue_create(-QueueId)
?- message_queue_create(Q).
Q = 1.
```

thread_get_message/1

Purpose: Reads a message from the calling thread's own queue, blocking until one arrives.

When to use: Use inside a worker started with `thread_create/3` and an `alias/1`, so other threads can address it by name.

```
% Syntax: thread_get_message(-Message)
?- thread_create((thread_get_message(M), assertz(got(M))), T, [alias(worker)]),
    thread_send_message(worker, hello(world)),
    thread_join(T, true),
    got(G).
G = hello(world).
```

thread_send_message/2

Purpose: Sends a message (any Prolog term) to a message queue.

When to use: Use to pass data to a consumer thread.

```
% Syntax: thread_send_message(+QueueOrThreadOrAlias, +Message)
?- message_queue_create(Q), thread_send_message(Q, hello).
true.

% Any term, not just an atom
?- message_queue_create(Q), thread_send_message(Q, point(1, [a,b])),
   thread_get_message(Q, M).
M = point(1,[a,b]).

% Producer pattern
produce(Queue, Items) :-
    member(Item, Items),
    thread_send_message(Queue, Item),
    fail ; true.
```

thread_get_message/2

Purpose: Retrieves a message from a queue, blocking until one is available.

When to use: Use in consumer threads to wait for and process incoming messages.

```
% Syntax: thread_get_message(+QueueOrThreadOrAlias, -Message)
?- message_queue_create(Q),
   thread_send_message(Q, world),
   thread_get_message(Q, Msg).
Msg = world.

% Consumer loop
consume_loop(Queue) :-
    thread_get_message(Queue, Msg),
    (   Msg = stop -> true
    ;   process(Msg),
        consume_loop(Queue)
    ).
```

v4.5.0 (ISS-2025-0629): the receive is **selective** - the first queued message that UNIFIES with the pattern is taken, the others stay queued in order, and the call blocks until a matching message arrives: `thread_send_message(Q, a(1)), thread_send_message(Q, b(2)), thread_get_message(Q, b(X))` gives `X = 2` and leaves `a(1)` in the queue.

thread_get_message/3

Purpose: `thread_get_message(+Queue, ?Pattern, +Options)` with `timeout(Seconds)` or `deadline(AbsTime)`: FAILS when no matching message arrives in time (`timeout(0)` polls).

```
?- message_queue_create(Q), thread_get_message(Q, M, [timeout(0.5)]).
false.
```

thread_peek_message/2

Purpose: Checks if a message is available in a queue without removing it.

When to use: Use to inspect the queue non-destructively, or to check for messages without blocking.

```
% Syntax: thread_peek_message(+QueueId, -Message)
?- message_queue_create(Q),
   thread_send_message(Q, test),
   thread_peek_message(Q, Msg).
Msg = test.

% Check if there is work to do
has_work(Queue) :-
   thread_peek_message(Queue, _).
```

`thread_peek_message/1` looks at the calling thread's own queue. Both are selective (v4.5.0).

message_queue_create/2, message_queue_destroy/1

v4.5.0 (ISS-2025-0630). `message_queue_create(-Q, [alias(A)])` names the queue ($Q = A$); `message_queue_destroy(+Q)` removes it - a thread blocked on it gets `existence_error(message_queue, Q)`.

thread_join/1

v4.5.0 (ISS-2025-0631). Succeeds when the thread's status is `true`, otherwise raises `error(thread_error(Id, Status), _)`.

thread_detach/1 (v4.5.0 semantics)

Detaching a live thread (detached already or not) succeeds; a finished thread that was never joined is reclaimed; an unknown or already reclaimed thread is `existence_error(thread, Id)` (ISS-2025-0620).

thread_exit/1

v4.5.0 (ISS-2025-0632). `thread_exit(Term)` ends the calling worker; `thread_join/2` reports `exited(Term)`. `catch/3` cannot intercept it; the cleanups of open `setup_call_cleanup/3` frames run. In `main` it is `permission_error(exit, thread, main)`. `thread_create/3` also accepts `at_exit(Goal)`, run in the thread after its goal.

thread_property/2

v4.5.0 (ISS-2025-0630). `thread_property(?Id, ?Property)` enumerates `id(N)`, `alias(A)`, `status(S)` (`running`, `true`, `false`, `exception(E)`, `exited(T)`) and `detached(Bool)`.

```
?- thread_property(main, status(S)).
S = running.
```

mutex_create/1,2, mutex_destroy/1, mutex_lock/1, mutex_trylock/1, mutex_unlock/1, mutex_unlock_all/0

v4.5.0 (ISS-2025-0631). Recursive mutexes. `mutex_create(-M)` gives an opaque `'$mutex'(N)`; `mutex_create(M, [alias(A)])` or `mutex_create(A)` names it; `mutex_lock/1` and `with_mutex/2` create a mutex named by an atom on first use. `mutex_unlock/1` by a thread that does not hold it is `permission_error(unlock, mutex, M)`. A thread that ends releases the mutexes it still holds.

with_mutex/2

`with_mutex(+Mutex, :Goal)`: `once(Goal)` holding the mutex, released however `Goal` ends.

```
bump :- with_mutex(counter, (retract(n(N)), N1 is N + 1, assertz(n(N1)))).
```

v4.6.0 (ISS-2025-0751): a mutex a thread still holds when it ends is released **and reported** with `print_message(warning, format("Thread ~p exited while holding mutex ~p (released)", [Id, M]))` (SWI warns too; 4.5 released it silently).

thread_signal/2

v4.6.0 (ISS-2025-0749). `thread_signal(+Thread, :Goal)` makes `Thread` run `Goal` at its next inference - pushed on the target's own goal stack, never asynchronously - or at once when the target is blocked in `thread_get_message/1, 2, 3`, `thread_join/1, 2`, `thread_sleep/1`, `mutex_lock/1`, `with_mutex/2`, a full queue's `thread_send_message/2, 3` or a full pool's `thread_create_in_pool/4`. The goal runs as `ignore(\+ \+ Goal)`: once, its bindings undone, a failure ignored; an **exception it raises unwinds the target** as if the interrupted goal had raised it (SWI). Signals run in the order they were sent. Signalling oneself runs the goal immediately. `main` (every non-worker thread: the CLI, an embedder, the IDE) can be signalled too. Unknown or finished thread: `existence_error(thread, T)`.

```
?- thread_create((repeat, fail), T), thread_signal(T, throw(stop)), thread_join(T, S).
S = exception(stop).
```

thread_statistics/3

v4.6.0 (ISS-2025-0750). `thread_statistics(+Thread, +Key, -Value)`: `statistics/2` for another thread. `cputime` (seconds, float), `thread_cputime`, `runtime` ([Ms, 0]) and inferences are measured on the target; every other key is answered by `statistics/2` (process-wide values).

```
?- thread_statistics(main, inferences, N).
```

message_queue_property/2

v4.6.0 (ISS-2025-0750). `message_queue_property(?Queue, ?Property)`: `alias(A)`, `size(N)` (messages waiting), `max_size(N)` (when created with it) and `waiting(N)` (threads blocked receiving). An unbound `Queue` enumerates the queues made by `message_queue_create/1, 2`; a thread id or alias names that thread's queue. `message_queue_create/2` accepts `max_size(N)`: a sender then waits for room.

thread_send_message/3

v4.6.0 (ISS-2025-0750). `thread_send_message(+Queue, +Term, +Options)`: as `/2`; for a queue created with `max_size(N)` that is full, `timeout(Seconds)` / `deadline(Time)` bound the wait and the call **fails** when it expires (SWI).

```
?- message_queue_create(Q, [max_size(1)]), thread_send_message(Q, a),
\+ thread_send_message(Q, b, [timeout(0.1)]).
```

mutex_property/2

v4.6.0 (ISS-2025-0750). `mutex_property(?Mutex, ?Property)`: `alias(A)` and `status(S)`, where `S` is `unlocked` or `locked(Owner, Count)` (the holding thread and its recursion count).

```
?- mutex_create(M, [alias(m)]), mutex_lock(m), mutex_property(m, status(S)).
S = locked(main, 1).
```

thread_pool_create/3

v4.6.0 (ISS-2025-0750, library(thread_pool)). `thread_pool_create(+Pool, +Size, +Options)` creates a named pool of at most `Size` running threads. `backlog(N)` bounds the creators that may wait for a slot (default unbounded); the other options are `thread_create/3` options every member inherits. A second pool with the same name is `permission_error(create, thread_pool, Pool)`.

thread_create_in_pool/4

v4.6.0 (ISS-2025-0750). `thread_create_in_pool(+Pool, :Goal, -Id, +Options)` creates a thread that occupies one slot until it ends. With the pool full, `wait(true)` (the default) blocks until a member ends; `wait(false)`, or a full backlog, raises `resource_error(threads_in_pool(Pool))`. An unknown pool is `existence_error(thread_pool, Pool)`.

```
?- thread_pool_create(workers, 4, []),
   thread_create_in_pool(workers, heavy_job, Id, []), thread_join(Id, S).
```

thread_pool_destroy/1

v4.6.0 (ISS-2025-0750). `thread_pool_destroy(+Pool)` removes the pool; running members go on, waiting creators get `existence_error(thread_pool, Pool)`.

thread_pool_property/2

v4.6.0 (ISS-2025-0750). `thread_pool_property(?Pool, ?Property): size(N), running(N), free(N), backlog(N)` (creators waiting) and `options(List)`.

current_thread_pool/1

v4.6.0 (ISS-2025-0750). `current_thread_pool(?Pool)` enumerates the pools.

22. CSV Predicates

CSV predicates provide parsing, serialization, and file I/O for comma-separated value data, commonly used for data exchange with spreadsheets and databases.

Sandbox (v4.5.0, ISS-2025-0625): `csv_read_file/2,3` and `csv_write_file/2,3` touch the host file system and are removed by `Prolog.enableSafeMode()`; `csv_parse/2` and `csv_serialize/2` stay.

csv_parse/2

Purpose: Parses a CSV string into a list of rows, where each row is a list of fields.

When to use: Use to process CSV data received as a string.

```
% Syntax: csv_parse(+CsvAtom, -Rows)
?- csv_parse('name,age\nJohn,30\nMary,25', Rows).
Rows = [['name', 'age'], ['John', '30'], ['Mary', '25']].

% Parse and skip header
csv_data(CsvString, Header, DataRows) :-
    csv_parse(CsvString, [Header|DataRows]).
```

csv_serialize/2

Purpose: Serializes a list of rows into a CSV string.

When to use: Use to produce CSV output from Prolog data.

```
% Syntax: csv_serialize(+Rows, -CsvAtom)
?- csv_serialize(['name', 'age'], ['John', '30'], Csv).
Csv = 'name,age\nJohn,30\n'.
```

csv_read_file/2

Purpose: Reads a CSV file and returns its contents as a list of rows.

When to use: Use to load tabular data directly from a file.

```
% Syntax: csv_read_file(+Path, -Rows)
?- csv_read_file('data.csv', Rows).
Rows = [['id', 'name', 'score'], ['1', 'Alice', '95'], ['2', 'Bob', '87']].

% Process file data
average_score(File, Avg) :-
    csv_read_file(File, [_Header|Rows]),
    findall(S, (member(Row, Rows), last(Row, SA), atom_number(SA, S)), Scores),
    sumlist(Scores, Total),
    length(Scores, N),
    Avg is Total / N.
```

csv_write_file/2

Purpose: Writes a list of rows to a CSV file.

When to use: Use to export Prolog data to CSV format for use in other tools.

```
% Syntax: csv_write_file(+Path, +Rows)
?- csv_write_file('output.csv', [['name', 'score'], ['Alice', '95'], ['Bob', '87']]).
true.
```

```
% Export query results
export_employees(File) :-
    findall([Name, Dept, Sal],
            employee(Name, Dept, Sal),
            Rows),
    csv_write_file(File, [['Name', 'Department', 'Salary']|Rows]).
```

23. Logging Predicates

Logging predicates provide structured, level-based logging for diagnostics and monitoring during program execution.

v4.5.0 (ISS-2025-0626): the log is **per engine** - its level and its sink (the engine's `user_error`, or the file named by `log_to_file/1`) belong to one `Prolog` instance; it used to be the JVM-wide `java.util.logging` logger, so one engine's `log_to_file/1` redirected every engine's log. A line reads `2026-09-23 17:00:00 INFO: message. log_level(L) takes debug, info, warning, error, off, all (domain_error(log_level, L) otherwise). log_to_file/1 is removed by Prolog.enableSafeMode().`

log_info/1

Purpose: Logs a message at the INFO level.

When to use: Use for general informational messages about normal program operation.

```
% Syntax: log_info(+Message)
?- log_info('Application started').
% Output: [INFO] Application started
true.
```

log_warning/1

Purpose: Logs a message at the WARNING level.

When to use: Use to report unexpected but non-fatal conditions.

```
% Syntax: log_warning(+Message)
?- log_warning('Configuration file missing, using defaults').
% Output: [WARNING] Configuration file missing, using defaults
true.
```

log_error/1

Purpose: Logs a message at the ERROR level.

When to use: Use to report errors that may affect program correctness.

```
% Syntax: log_error(+Message)
?- log_error('Failed to connect to database').
% Output: [ERROR] Failed to connect to database
true.

% Log errors in exception handlers
safe_process(Goal) :-
    catch(
        call(Goal),
        Error,
        ( term_to_atom(Error, Msg),
          log_error(Msg),
          fail
        )
    ).
```

log_debug/1

Purpose: Logs a message at the DEBUG level.

When to use: Use for detailed diagnostic information during development.

```
% Syntax: log_debug(+Message)
?- log_debug('Entering solve/3 with X = 42').
% Output: [DEBUG] Entering solve/3 with X = 42
true.
```

log_level/1

Purpose: Sets the minimum logging level. Messages below this level are suppressed.

When to use: Use to control verbosity at runtime.

```
% Syntax: log_level(+Level)
% Levels (from most to least verbose): debug, info, warning, error
?- log_level(warning).
true.

% Now only warning and error messages are displayed
?- log_info('This will be suppressed').
true.

?- log_warning('This will be shown').
% Output: [WARNING] This will be shown
true.
```

log_to_file/1

Purpose: Redirects log output to a file.

When to use: Use to capture logs for later analysis instead of writing to the console.

```
% Syntax: log_to_file(+FilePath)
?- log_to_file('app.log').
true.

?- log_info('This goes to the file').
true.

% Set up application logging
init_logging :-
    log_to_file('logs/app.log'),
    log_level(info),
    log_info('Logging initialized').
```

24. CLP(FD) Predicates

CLP(FD) (Constraint Logic Programming over Finite Domains) predicates allow you to declare constraints over integer variables and let the solver find valid assignments. JProlog v3.0.0 uses a clean-room interval-domain solver (builtin.clpfd.v2) by default: a per-query trail-backtracked constraint store, bound-consistency propagation (e.g. real $\neq$ propagation, Cmp/Sum/Mul/Abs/AllDifferent/Linear/Reified/Mod), and sound labeling. The legacy AC-3 store remains available via `-Djprolog.clpfd=legacy`.

v4.5.0 (wave P5, SWI-Prolog library(clpfd) semantics):

- **Domains are interval lists** with holes (`X in 1..3 \ 5..7`, $\neq$ removes interior values) and may be unbounded: an unconstrained CLP(FD) variable is `inf..sup`, as in SWI.
- **Plain unification propagates**: `X in 0..9, Y in 0..9, X+Y = 9, X = 4` binds `Y = 5`; every variable whose domain becomes a singleton is bound.
- **Labeling is lazy**: `label/1` and `labeling/2` hand out one solution at a time on the engine's own choice points (the first 20-queens solution no longer costs the whole search tree), with SWI's options and `label/1 = leftmost selection`.
- **Exact integers**: ground expressions are evaluated exactly (`X = 1012*1012`), and a value beyond the 64-bit domain range that a constraint determines is bound exactly (`X = Y+1, Y = 263`). Coefficients must fit 64 bits (else `representation_error(max_integer)`) - lifted in v4.6.0, see below.
- **Cyclic difference constraints fail fast**: `X #> Y, Y #> X` fails at once (negative-cycle check), also over huge finite domains.
- New: `ins/2`, `sum/3`, `scalar_product/4`, reification (`#<==>`, `#==>`, `#<==`, `#\`, `#/\`, `#\`), `//`, `div`, `rem`, `mod`, `^` in expressions, `fd_inf/2`, `fd_sup/2`, `fd_var/1`, `element/3`, `tuples_in/2`, `global_cardinality/2`, `transpose/2`, and a domain-consistent `all_distinct/1`.

clpfd:attribute_goals//1

v4.6.0 (ISS-2025-0760, wave Q5): the SWI hook behind residual goals -

`phrase(clpfd:attribute_goals(X), Gs)` lists the CLP(FD) goals (domain and live constraints, SWI's printed forms) that constrain `X`. `copy_term/3`, the answers of the CLI and Prolog `residualGoals` use it.

```
?- X #> Y, phrase(clpfd:attribute_goals(X), Gs).
Gs = [Y#=<X+ -1].
```

v4.6.0 (wave Q5):

- **Residual constraints**: an answer, `copy_term/3` and `clpfd:attribute_goals//1` show the constraints still alive on the answer's variables in SWI-Prolog's printed forms, after their domains (`X in 1..2\4..5`, shown unless it is `inf..sup` and a constraint is left): `X #> Y` answers `Y#=<X+ -1`, `X #> Y+Z` answers `Y+Z#<X`, `B #<==> X #> 3` answers `B in 0..1, X#>=4#<==>B`; the globals print by name (`all_different([X,Y])`, `circuit(Vs)`, ...). An entailed constraint is not shown (`X #> Y, X = 5` answers `X = 5, Y in inf..4`), an auxiliary variable of the compiled store is a fresh `_A`. `freeze/2`, `dif/2` and `when/2` residuals are shown as before. `X #> Y` between two variables unifies them (SWI).
- **Coefficients beyond 64 bits**: `X*1020 #> Y` is an ordinary linear constraint (an exact, slower propagation path only for such constraints; everything else keeps the 64-bit path).
- **Branch and bound** (`labeling/2` with `min(E)/max(E)`) is one iterative search that tightens the bound at every solution; it no longer recurses once per variable.
- New: `circuit/1`, `cumulative/1,2`, `disjoint2/1`, `lex_chain/1`, `chain/2`, `automaton/3,8`, `zcompare/3`, `fd_degree/2`, `global_cardinality/3`.

in/2

Purpose: Constrains a variable to a finite domain range.

When to use: Use to declare the possible values for a constraint variable before posting constraints.

v3.5.0: unification respects domains - `X in 1..3, X = 5` fails (it used to succeed unsoundly), while `X in 1..3, X = 2` succeeds; constraint posts are undone on backtracking.

v4.5.0: a domain is an integer N, a range `Lo..Hi` whose bounds may be `inf/sup`, or a union `D1 \ D2`. An unbound domain raises `instantiation_error`, anything else `type_error(clpfd_domain, D)`.

```
?- X in 1..3 \ 5..7, fd_dom(X, D).
D = 1..3\5..7.

?- X in 1..3 \ 5..7, X = 4.
false.
```

ins/2

Purpose: `Vars ins Domain` posts `V in Domain` for every element of the list `Vars` (v4.5.0).

```
?- [X,Y] ins 0..5, X + Y #= 9, X = 4.
X = 4, Y = 5.

?- _ ins 1..3.
% throws instantiation_error
```

```
% Syntax: X in +Low..+High
?- X in 1..10.
true.
```

```
% Constrain multiple variables
sudoku_vars(Vars) :-
    length(Vars, 9),
    maplist(clpfd:in_(1..9), Vars).
```

#=/2

Purpose: Posts an equality constraint between two arithmetic expressions.

When to use: Use instead of `is/2` when working with constraint variables that are not yet bound.

v3.5.0: when propagation narrows a variable to a singleton domain the variable is bound - `X #= 2` gives `X = 2` without labeling.

v4.5.0: propagation also runs after plain unification (`X #= Y+Z, Y = 1, Z = 2` binds `X = 3`); expressions accept `//` (truncating), `div` (floored), `rem`, `mod` (sign of the divisor), `^` (power), `abs/1`, `min/2`, `max/2`; ground sub-expressions are exact big integers (`X #= 2^100`), and a division by zero makes the constraint fail.

```
?- X #= -7 // 2, Y #= -7 div 2, Z #= -7 rem 2, W #= -7 mod 2.
X = -3, Y = -4, Z = -1, W = 1.

?- X #= 2^100.
X = 1267650600228229401496703205376.

?- X in -5..5, Y #= abs(X), Y #>= 4, fd_dom(X, D).
D = -5.. -4\4..5.
```

v3.6.0: non-linear constraints propagate instead of failing silently - variable products (`X*Y`) and squares (`X*X #= 16` narrows `X in 0..10` to `X = 4` before labeling) work via interval products, as do `abs/1`, `min/2`,

$\text{max}/2$ and $E \bmod M$ (constant positive M); a genuinely unsupported functor raises `type_error(evaluable, F/N)` and a float coefficient ($2.5 * X$) raises `type_error(integer, 2.5)` instead of truncating.

```
% Syntax: Expr1 #= Expr2
?- X in 1..10, X #= 3 + 4.
X = 7.

% Bidirectional reasoning
?- X in 1..10, Y in 1..10, X + Y #= 10, X #= 3.
X = 3, Y = 7.

% Non-linear products (v3.6.0)
?- X in 0..10, X*X #= 16.
X = 4.

?- A in 1..5, B in 1..5, C in 1..5,
   A*A + B*B #= C*C, label([A, B, C]).
A = 3, B = 4, C = 5 ;
A = 4, B = 3, C = 5.
```

#\=/2

Purpose: Posts a disequality (not-equal) constraint.

When to use: Use to declare that two expressions must have different values.

v3.5.0: multi-variable expressions work ($X \# \backslash = Y + 1$), and $X \# \backslash = X$ fails.

v4.5.0: as soon as all but one variable are fixed the remaining value is removed from the domain, including interior values ($X \text{ in } 1..5, Y \text{ in } 1..5, X \# \backslash = Y, X = 2$ leaves $Y \text{ in } 1 \backslash 3..5$).

```
% Syntax: Expr1 #\= Expr2
?- X in 1..5, Y in 1..5, X #\= Y, X #= 3, label([Y]).
Y = 1 ; Y = 2 ; Y = 4 ; Y = 5.
```

#</2

Purpose: Posts a strict less-than constraint.

```
% Syntax: Expr1 #< Expr2
?- X in 1..10, Y in 1..10, X #< Y, X #= 8, label([Y]).
Y = 9 ; Y = 10.
```

#>/2

Purpose: Posts a strict greater-than constraint.

v4.5.0: a cycle of difference constraints that cannot hold fails immediately, even over unbounded or huge domains ($X \# > Y, Y \# > X$ used to run 29 s and exhaust memory).

```
% Syntax: Expr1 #> Expr2
?- X in 1..10, X #> 7, label([X]).
X = 8 ; X = 9 ; X = 10.
```

#=</2

Purpose: Posts a less-than-or-equal constraint.

```
% Syntax: Expr1 #=< Expr2
?- X in 1..10, X #=< 3, label([X]).
```

```
X = 1 ; X = 2 ; X = 3.
```

#>=/2

Purpose: Posts a greater-than-or-equal constraint.

```
% Syntax: Expr1 #>= Expr2
?- X in 1..5, X #>= 4, label([X]).
X = 4 ; X = 5.
```

all_different/1, all_distinct/1

Purpose: Constrains all variables in a list to take pairwise different values. `all_different/1` prunes a value as soon as another variable is fixed to it (plus a pigeonhole check). *v4.5.0:* `all_distinct/1` is domain-consistent, as in SWI - Régin's matching-based propagator removes every value that cannot be part of any all-different assignment (Hall sets).

When to use: Use for problems like Sudoku, graph coloring, or any assignment problem requiring distinct values; prefer `all_distinct/1` when propagation strength matters more than the per-node cost.

```
% Syntax: all_different(+Vars)
?- X in 1..3, Y in 1..3, Z in 1..3,
   all_different([X, Y, Z]),
   label([X, Y, Z]).
X = 1, Y = 2, Z = 3 ;
X = 1, Y = 3, Z = 2 ;
...

% Hall set {1,2} for X, Y forces Z (v4.5.0)
?- X in 1..2, Y in 1..2, Z in 1..3, all_distinct([X, Y, Z]).
Z = 3.
```

sum/3, scalar_product/4

Purpose: `sum(Vars, Op, Expr)` posts $V_1 + \dots + V_n \text{ Op Expr}$; `scalar_product(Cs, Vs, Op, Expr)` posts $C_1*V_1 + \dots + C_n*V_n \text{ Op Expr}$ (`Op` is one of `#=`, `#\=`, `#<`, `#>`, `#=<`, `#>=`). *v4.5.0.* Library predicates (module `clpfd`): a program that defines its own `sum/3` keeps it. A bad `Op` raises `domain_error(scalar_product_relation, Op)`.

```
?- [A,B,C] ins 0..1, sum([A,B,C], #>=, 3).
A = 1, B = 1, C = 1.

?- [X,Y] ins 0..6, findall(X-Y, (scalar_product([2,3], [X,Y], #=, 12), label([X,Y])), L).
L = [0-4, 3-2, 6-0].
```

#<==>/2, #==>/2, #<==/2, #\!/2, #!/2, #/2, #/1 (reification)

Purpose: Boolean combinations of reifiable constraints (*v4.5.0*). A reifiable expression is a comparison `#=/#\=#</#\=#>/#=</#>=`, `X in Dom`, a 0/1 variable or integer, or a nested connective. `B #<==> C` binds `B` to 1 when `C` is entailed and to 0 when it is disentailed, and posts `C` (or its negation) when `B` becomes known. Anything else raises `domain_error(clpfd_reifiable_expression, E)`. Operator priorities are SWI's (see the operator reference).

```
?- X in 0..10, B #<==> (X #> 5), B = 0, fd_dom(X, D).
D = 0..5.

?- X in 1..3, #\ X #= 2, fd_dom(X, D).
D = 1\3.
```

```
% reification counting: exactly two of three variables equal 1
?- length(Xs, 3), Xs ins 0..2,
   maplist([X,B]>>(B #<==> (X #= 1)), Xs, Bs), sum(Bs, #=, 2),
   findall(Xs, label(Xs), L), length(L, N).
N = 6.
```

element/3, tuples_in/2, global_cardinality/2

Purpose (v4.5.0, library predicates in module `clpfd`):

- `element(I, List, V)`: `V` is the `I`-th (1-based) element of `List`.
- `tuples_in(Tuples, Relation)`: every tuple (a list of variables) is one of the rows of the ground list of lists `Relation`.
- `global_cardinality(Vs, Pairs)`: `Pairs` is a list of Key-Count; every element of `Vs` is one of the keys, and `Count` (an integer or a CLP(FD) variable) is the number of occurrences of `Key` (basic propagation).

```
?- element(I, [10,20,30], V), V #> 15, fd_dom(I, D).
D = 2..3.

?- tuples_in([[X,Y]], [[1,2],[2,3],[5,5]]), X #< 3, findall(X-Y, label([X,Y]), L).
L = [1-2, 2-3].

?- Vs = [A,B,C], global_cardinality(Vs, [1-2, 2-1]), findall(Vs, label(Vs), L).
L = [[1,1,2], [1,2,1], [2,1,1]].
```

circuit/1

Purpose: `circuit(Succs)`: the list `Succs` of 1-based successors forms a single Hamiltonian cycle - element `I` is the node that follows node `I`. v4.6.0: all-different (domain consistent), no self loop, a chain of fixed successors shorter than `N` may not close on itself, and the successor graph must stay strongly connected (no subtours).

```
?- length(L, 3), circuit(L), findall(L, label(L), Ls).
Ls = [[2,3,1], [3,1,2]].

?- circuit([2,1,_]).
false.
```

cumulative/1, cumulative/2

Purpose: `cumulative(Tasks, Options)`: `Tasks` is a list of `task(S, D, E, C, T)` (start, duration, end, resource amount, task id); at every moment the running tasks use at most the limit of the option `limit(L)` (default 1; `cumulative/1` is limit 1). `S + D #= E` is posted for every task. v4.6.0: time-table propagation over the compulsory parts. An unknown option is `domain_error(cumulative_option, 0)`, a malformed task `domain_error(cumulative_task, T)`.

```
?- Ts = [task(S1,3,_,1,a), task(S2,2,_,1,b)], [S1,S2] ins 0..4,
   cumulative(Ts), S1 #< S2, label([S1,S2]).
S1 = 0, S2 = 3 ;
S1 = 0, S2 = 4 ;
S1 = 1, S2 = 4 ;
false.
```

disjoint2/1

Purpose: `disjoint2(Rects)`: the rectangles $F(X, W, Y, H)$ (x , width, y , height; any functor) do not overlap pairwise (v4.6.0, a reified disjunction per pair).

```
?- disjoint2([r(X,2,0,2), r(Y,2,0,2)]), [X,Y] ins 0..3, findall(X-Y, label([X,Y]), L).
L = [0-2, 0-3, 1-3, 2-0, 3-0, 3-1].
```

lex_chain/1

Purpose: `lex_chain(Lists)`: the lists (of integers or CLP(FD) variables) are lexicographically non-decreasing, each with respect to the next (v4.6.0).

```
?- lex_chain([[A],[1]]), A in 0..5.
A in 0..1.
```

chain/2

Purpose: `chain(Zs, Relation)`: `Relation` (`#=`, `#=<`, `#>=`, `#<` or `#>`) holds between every pair of neighbours of `Zs`; another relation is `domain_error(chain_relation, R)` (v4.6.0).

```
?- chain([A,B,C], #<), A = 1, C = 3.
A = 1, B = 2, C = 3.
```

automaton/3, automaton/8

Purpose: `automaton(Sigs, Nodes, Arcs)`: the finite automaton with `Nodes` (`source(N)` and `sink(N)` terms) and `Arcs` (`arc(N0, Label, N1)`) accepts the list of CLP(FD) variables `Sigs`. `automaton(Seqs, Template, Sigs, Nodes, Arcs, Counters, Initials, Finals)` adds counters: an arc `arc(N0, L, N1, Exprs)` updates each counter to its expression, where the `Counters` variables stand for the previous values and the `Template` variables for the current element of `Seqs`; an arc without expressions leaves the counters unchanged; `Initials` are the initial values and `Finals` the values after the last transition (SWI semantics; v4.6.0, decomposed into `tuples_in/2` and reified implications).

```
?- length(Vs, 3), automaton(Vs, [source(a),sink(c)],
    [arc(a,0,a), arc(a,1,b), arc(b,0,a), arc(b,1,c), arc(c,0,c), arc(c,1,c)]),
    findall(Vs, label(Vs), L).
L = [[0,1,1], [1,1,0], [1,1,1]].

?- length(Vs, 4), Vs ins 0..1,
    automaton(Vs, _, Vs, [source(s),sink(s)], [arc(s,0,s), arc(s,1,s,[C+1])], [C], [0], [2]),
    findall(Vs, label(Vs), L), length(L, N).
N = 6.
```

zcompare/3

Purpose: `zcompare(Order, A, B)`: `Order` (`<`, `=` or `>`) is the order of the integers `A` and `B`, reified: it is bound as soon as the domains decide it, and binding it posts the matching constraint (v4.6.0).

```
?- X in 1..3, Y in 5..7, zcompare(0, X, Y).
0 = (<), X in 1..3, Y in 5..7.
```

fd_degree/2

Purpose: `fd_degree(X, D)`: `D` is the number of constraints still alive on the CLP(FD) variable `X` (0 for an integer or a variable with no constraint; v4.6.0).

```
?- X #> Y, Y #> Z, fd_degree(Y, D).
D = 2.
```

global_cardinality/3

Purpose: `global_cardinality(Vs, Pairs, Options)`: `global_cardinality/2` with options: `consistency(value)` (accepted: it is the propagation `global_cardinality/2` does) and `cost(Cost, Matrix)` - `Matrix` has one row per variable and one column per key, and `Cost` is the sum of the entries the assignment selects. Another option is `domain_error(global_cardinality_option, 0)` (v4.6.0).

```
?- global_cardinality([X,Y,Z], [1-_, 2-], [cost(C, [[1,5],[1,5],[2,2]])],
    once(labeling([min(C)], [X,Y,Z]))).
C = 4, X = 1, Y = 1, Z = 1.
```

transpose/2

Purpose: `transpose(Matrix, Transposed)` for a list of equal-length lists (SWI library(`clpfd`), the sudoku idiom). v4.5.0.

```
?- transpose([[1,2,3],[4,5,6]], T).
T = [[1,4], [2,5], [3,6]].
```

label/1

Purpose: Assigns concrete values to constraint variables by exhaustive search.

When to use: Use after posting all constraints to enumerate solutions.

v3.6.0: list elements that are neither variables nor integers raise `type_error(integer, T)` - `label([a])` no longer succeeds silently; ground integers remain legal.

v4.5.0: `label(Vs)` is `labeling([], Vs)`: **leftmost** variable selection (it used to be first-fail), smallest value first, and **lazy** - solutions are produced one at a time on backtracking, so `once(label(Vs))` explores only as far as the first solution. Every variable must have a finite domain, otherwise `instantiation_error` (SWI): `X #> 3, label([X])` raises.

```
% Syntax: label(+Vars)
?- X in 1..3, Y in 1..3, X #< Y, label([X, Y]).
X = 1, Y = 2 ;
X = 1, Y = 3 ;
X = 2, Y = 3.
```

labeling/2

Purpose: Labels variables with options controlling search strategy.

When to use: Use for fine-grained control over variable and value ordering.

v3.6.0: the options are honoured - `leftmost/ff/ffc/min/max` select the branching variable, `up/down` set the value enumeration order (`down` yields the largest value first), `min(Expr)/max(Expr)` order solutions optimum-first, and `step/enum` are accepted. An unknown option raises `domain_error(labeling_option, 0)`, an unbound option raises `instantiation_error`, and a

non-list options argument raises `type_error(list, 0)`. Non-variable, non-integer list elements raise `type_error(integer, T)` (as for `label/1`).

v4.5.0: lazy, and SWI's full option set: variable selection `leftmost` (default), `ff`, `ffc` (first-fail, ties broken by the number of constraints), `min/max` (smallest lower / largest upper bound); value order `up` (default), `down`; branching step (default: `X = V ; X #\= V`), `enum` (one branch per value), `bisect` (`X #=< Mid ; X #> Mid`); and optimisation `min(Expr)/max(Expr)` by branch and bound - solutions come in order of the objective, ties in labeling order; several objectives are lexicographic. At most one option per category: repeating one raises `domain_error(nonrepeating_labeling_options, Opts)`, two different ones `domain_error(consistent_labeling_options, Opts)`.

```
% Syntax: labeling(+Options, +Vars)
% Options: leftmost, ff (first-fail), ffc, min, max, up, down, step, enum, bisect,
%           min(Expr), max(Expr)
?- X in 1..5, Y in 1..5, X + Y #= 6,
    labeling([ff], [X, Y]).
X = 1, Y = 5 ;
X = 2, Y = 4 ;
...

% Value order (v3.6.0)
?- X in 0..5, labeling([down], [X]).
X = 5 ;
X = 4 ;
...

?- X in 1..3, labeling([bogus], [X]).
% throws domain_error(labeling_option, bogus)

% Branch and bound (v4.5.0): best product first
?- X in 1..5, Y in 1..5, X + Y #= 6, labeling([max(X*Y)], [X, Y]).
X = 3, Y = 3 ;
X = 2, Y = 4 ;
X = 4, Y = 2 ;
...
```

indomain/1

Purpose: Nondeterministically assigns a value from the domain of a single variable.

When to use: Use for custom labeling strategies where you want to control variable order manually.

v4.5.0: `indomain(X)` is `label([X])` - lazy, so a huge domain (`X in 1..2000000000`, `indomain(X)`) hands out its first value at once instead of raising `resource_error`.

```
% Syntax: indomain(+Var)
?- X in 1..3, indomain(X).
X = 1 ; X = 2 ; X = 3.
```

fd_dom/2

Purpose: Returns the current domain of a constraint variable.

When to use: Use to inspect the remaining possible values after constraint propagation.

```
% Syntax: fd_dom(+Var, -Domain)
?- X in 1..10, X #> 7, fd_dom(X, Dom).
Dom = 8..10.
```

v4.5.0: unions and infinite bounds print as SWI does (`1..3\5..7`, `4..sup`); an unconstrained variable is `inf..sup`, an integer `N` is `N..N`, and any other term raises `type_error(integer, T)`.

fd_size/2

Purpose: Returns the number of elements in the current domain of a variable.

When to use: Use in custom search heuristics (e.g., first-fail selects the variable with smallest domain).

```
% Syntax: fd_size(+Var, -Size)
?- X in 1..10, X #> 7, fd_size(X, S).
S = 3.
```

v4.5.0: an unbounded domain has size sup.

fd_inf/2, fd_sup/2, fd_var/1

Purpose: `fd_inf(X, Inf)` / `fd_sup(X, Sup)` give the smallest / largest value of the domain (`inf` / `sup` when unbounded); `fd_var(X)` succeeds when `X` is an unbound CLP(FD) variable (v4.5.0).

```
?- X #> 3, fd_inf(X, I), fd_sup(X, S).
I = 4, S = sup.

?- X in 1..3, fd_var(X).
true.
```

25. Tabling Predicates

Tabling (memoization, tabled resolution) caches the answers of a tabled predicate per **call variant**, so a repeated call with the same argument pattern is answered from the table and a left-recursive or cyclic definition terminates instead of looping.

One implementation since 4.1.0

Tabling used to differ per engine. The engine that carried the bounded, sometimes-wrong implementation was deleted in 4.1.0, so what follows is simply how tabling works.

	the pre-4.1.0 fallback engines (deleted)	the engine (v4, an option since 3.12.0, the default since 4.0.0)
Algorithm	bounded re-evaluation: the goal was re-run at most 100 times over name-keyed answer maps	linear tabling with completion (SLD + iterative completion, B-Prolog/DRA style) in the machine's own choice points
Correctness	wrong answers for a left-recursive predicate over a long chain (LIM-038 / design limit L-03)	correct and complete for definite programs, left recursion included
Recursion depth	the pre-4.0.0 recursive solver's 2 000-deep Java cap	none (a tabled call is a choice point)
Inference budget / Stop	not enforced inside the fixpoint	enforced
Four-port trace / debugger	the whole tabled call was opaque	Call/Exit/Redo/Fail like any predicate
current_table/2	not available	available

The classic repro, which **succeeds today and failed on the deleted fallback**:

```

edge(I, J) :- between(1, 3000, I), J is I + 1.
:- table path/2.
path(X, Y) :- edge(X, Y).
path(X, Y) :- path(X, Z), edge(Z, Y).

?- path(1, 3001).                % true   (old fallback: fails)
?- path(1, 51).                 % true   (old fallback: fails)
?- findall(Y, path(1, Y), L), length(L, N). % N = 3000

```

Semantics

- A **variant table** is created per tabled subgoal. `path(1, Y)` and `path(1, 51)` are different variants and each gets its own table; both are answered correctly.
- The first call to a variant is its **generator**: it runs the predicate's clauses against a private copy of the call and records every answer. A call to a variant that is still being evaluated is a **consumer** over the answers found so far. When a generator that leads its strongly connected component exhausts its clauses, it re-runs them until a round produces no new answer anywhere in the component, and the whole component is then marked complete. There is no iteration cap; termination follows from the finite, deduplicated answer set.
- Answers are deduplicated by variant, and returned in the order they were first found.
- A `!` in a tabled clause body is **local to that body**: it prunes the body's own choice points and never truncates the table.
- An evaluation abandoned by an exception, a cut or the resource guard **discards its tables**, so the next call recomputes them rather than reading a partial answer set.
- Tabled calls work inside `findall/3`, `bagof/3`, `\+ /1`, `catch/3`, `once/1` and `forall/2`.
- Invalidation policy**: asserting to or retracting from a *tabled* predicate drops that predicate's tables. A change to a **non-tabled** predicate that a tabled one depends on is **not** tracked - call `abolish_all_tables/0` yourself (this is also what XSB requires). Tables persist across queries; two

safety caps (100 000 tables, 4 000 000 answers) drop the oldest completed tables at a query boundary so a long-lived engine cannot grow the store without bound.

- v4.6.0 (ISS-2025-0755): ``tnot/1``, ``undefined/0`` and ``call_delays/2`` give a minimal well-founded semantics - see `tnot/1` below. Plain `\+` keeps the 4.5 rule that follows.
- v4.6.0 (ISS-2025-0752): **tables across threads**. Every worker thread (`thread_create/2, 3`, the `concurrent_*` pool) evaluates in a **private** table space that ends with it - SWI's default; every non-worker thread (main) shares the engine's main space. A predicate declared `:- table p/1` as `shared` also publishes its **complete** tables engine-wide, where every thread reuses them; an incomplete table always belongs to the thread evaluating it, and a thread that calls a shared variant another thread is still evaluating evaluates it itself instead of waiting (SWI waits) - so tabling can never deadlock two threads. Asserting to / retracting from a tabled predicate, `abolish_table/1` and `abolish_all_tables/0` invalidate the tables of every space.
- v4.5.0 (ISS-2025-0661, LIM-046): **non-stratified negation raises**. `\+ G` inside a tabled evaluation whose `G` reads a table that an ancestor of the negation is still evaluating (e.g. `:- table p/1. p(X) :- \+ p(X).`, or `win(X) :- move(X, Y), \+ win(Y).` over a cycle) raises `error(permission_error(negate, incomplete_table, G), (\+)/1)` - 4.4.0 answered it inconsistently, SWI answers it with the well-founded semantics, which JProlog does not have. Stratified negation (the negated goal's tables complete inside the negation) is unaffected.

table/1

Purpose: Declares a predicate as tabled, enabling automatic memoization.

When to use: Use for predicates with overlapping subproblems (e.g., Fibonacci, transitive closure) or left-recursive definitions. Declare the predicate **before** it is first called.

```
% Syntax: :- table Predicate/Arity.          (also callable as a goal: table(Predicate/Arity)
)
:- table fib/2.
fib(0, 0).
fib(1, 1).
fib(N, F) :-
    N > 1,
    N1 is N - 1, N2 is N - 2,
    fib(N1, F1), fib(N2, F2),
    F is F1 + F2.

% Without tabling: exponential time
% With tabling: linear time
?- fib(30, F).
F = 832040.
?- fib(1000, F).                % 209 digits, ~0.1 s on v4
```

Errors (v4.6.0): `instantiation_error`, `type_error(predicate_indicator, S)`, `domain_error(tabled_mode, M)` (SWI's name; 4.5 said `table_mode`), `domain_error(lattice_arity, N)`, `domain_error(po_arity, N)`, `domain_error(table_option, O)`.

v4.5.0 (ISS-2025-0572, wave P3.3): every form of the directive is accepted - `:- table ev/1, od/1.`, `:- table([a/1, b/1]).`, `Name//Arity` (a DCG non-terminal), `Spec as Options` - and **mode-directed tabling**: `:- table path(_, _, min).` keeps, per distinct combination of the index arguments (`_, index, +`), one aggregated answer. A table directive never fails silently (the old code logged a warning and the program then looped). `predicate_property(P, tabled)` holds for a tabled predicate.

v4.6.0 (ISS-2025-0754, SWI's boot/`tabling.pl` semantics): the modes are `first / -`, `last`, `min` and `max` (standard order of terms), `sum`, `lattice(PI)` - `call(PI, Old, New, Agg)`, `PI` of arity 3 given as `Name/3`, `Name`, `Module:Name/3` or a head - and `po(PI)` - keep `Old` when `call(PI, Old, New)` succeeds, else take `New` (`po(<'/2)` keeps the smallest). **Each** moded argument is aggregated on its own, and the answer is replaced only when the aggregate is not a variant of the old one; a failing lattice join rejects the new answer. A call with a **bound** moded argument is evaluated with it free and then unified with

the aggregate (SWI): with `:- table f(_, first). f(k, 1). f(k, 2)., f(k, 2) fails`. sum over a recursive component may not terminate (every completion round adds again).

v4.6.0 (ISS-2025-0753): Spec as Options - variant (the default), shared and private (see "tables across threads" above) are implemented; subsumptive, incremental, opaque, monotonic, lazy, dynamic, `max_answers/1`, `subgoal_abstract/1` and `answer_abstract/1` are accepted with a Warning: on `user_error` and variant tabling is used; any other option is `domain_error(table_option, 0)`. Options may be one option, a , -sequence or a list.

```
:- table sp(_, _, min).
sp(X, Y, D) :- edge(X, Y, D).
sp(X, Y, D) :- sp(X, Z, D1), edge(Z, Y, D2), D is D1 + D2.
?- sp(a, b, D).                % the shortest distance only, even on a cyclic graph

:- table lp(_, _, lattice(longer/3)).
longer(P1, P2, P) :- length(P1, L1), length(P2, L2), (L1 >= L2 -> P = P1 ; P = P2).
```

tnot/1

v4.6.0 (ISS-2025-0755). `tnot(:Goal)`: tabled negation of a **tabled** Goal under a minimal well-founded semantics (SWI's algorithm on linear tabling). An unconditional answer of Goal makes it fail; a complete table with no answer makes it succeed; otherwise - Goal is still being evaluated by an ancestor (a loop through negation) or has only conditional answers - it succeeds with `tnot(Goal)` **delayed**. The answers derived with delays are **conditional**; when the strongly connected component completes they are simplified (a literal that became true is dropped, one that became false kills the derivation) and what is left is **undefined**. The CLI prints `undefined` for such an answer; an embedder asks `Prolog.currentAnswerDelays()` inside a `solveStream` sink, or uses `call_delays/2`. A non-tabled goal is `permission_error(tnot, non_tabled_procedure, Name/Arity)`.

```
:- table win/1.
win(X) :- move(X, Y), tnot(win(Y)).
move(a, b). move(b, a). move(b, c). move(c, d).
?- win(c).           % true
?- win(d).           % false
?- win(a).           % undefined
```

Not done (LIM-046): answer completion (an unsupported positive loop among conditional answers stays `undefined` instead of becoming `false`), the residual program, delays through mode-directed tables.

undefined/0

v4.6.0 (ISS-2025-0755). Succeeds with the truth value *undefined* (a delay that never resolves), like SWI's `undefined :- tnot(undefined)`.

call_delays/2

v4.6.0 (ISS-2025-0755). `call_delays(:Goal, -Delays)`: runs Goal; Delays is true for an unconditional solution, otherwise the conjunction of the literals Goal delayed (`tnot(G)`, a conditional answer G, `undefined`). The delays are consumed: the enclosing derivation stays unconditional.

```
?- call_delays(win(a), D).
D = win(a).
```

abolish_all_tables/0

Purpose: Clears all cached answers, forcing every tabled predicate to recompute on its next call. The table **declarations are kept**.

When to use: after `assert/retract` on facts a tabled predicate depends on, or between computation phases.

```
% Syntax: abolish_all_tables
?- fib(10, F).
F = 55.

?- abolish_all_tables.
true.

% Next call to fib/2 recomputes from scratch
```

Errors (v4): `permission_error(modify, table, ...)` if called from inside a running tabled evaluation.

`abolish_table/1`

Purpose: Clears the cached answers of one predicate **and un-declares it as tabled**.

When to use: to invalidate one predicate's table selectively, or to switch a predicate back to ordinary evaluation. Call `table(Name/Arity)` again to re-enable tabling for it.

```
% Syntax: abolish_table(+Name/Arity)
?- abolish_table(fib/2).
true.

% Only fib/2 is affected; other tabled predicates keep their tables and their declarations
```

Errors: `instantiation_error` for an unbound argument, `type_error(predicate_indicator, T)` for anything that is not `Name/Arity`, `permission_error(modify, table, ...)` from inside a running tabled evaluation. (On the engines deleted in 4.0.0/4.1.0 a malformed argument made the call fail silently.)

`current_table/2`

Purpose: Enumerates the tables that currently exist. Added with engine v4 (v3.12.0).

When to use: debugging a tabled program, or checking that a table was really discarded.

```
% Syntax: current_table(?Variant, ?Status)      Status = complete | incomplete
?- path(a, c), current_table(V, S).
V = path(a, c), S = complete ;
V = path(a, _), S = complete.
```

`Variant` is unified with each table's call pattern (so a partially instantiated `Variant` enumerates every table it matches, as `current_op/3` does); `Status` is `complete` for a table whose evaluation has finished and `incomplete` for one still being evaluated (only observable from inside a tabled computation).

26. HTTP Predicates

HTTP predicates provide both server-side and client-side HTTP capabilities, enabling JProlog programs to serve web APIs and consume external HTTP services.

http_server/2

Purpose: Starts an HTTP server on the specified port with given options.

When to use: Use to create web services or REST APIs from Prolog.

```
% Syntax: http_server(+Port, +Options)
?- http_server(8080, []).
true.

% Start server with registered handlers
start_api :-
    http_handler('/hello', handle_hello, []),
    http_server(8080, []).
```

http_stop/1

Purpose: Stops a running HTTP server on the specified port.

When to use: Use to gracefully shut down a server.

```
% Syntax: http_stop(+Port)
?- http_stop(8080).
true.
```

http_handler/3

Purpose: Registers a handler predicate for a URL path pattern.

When to use: Use to define routing for your HTTP server before starting it.

```
% Syntax: http_handler(+Path, +Handler, +Options)
?- http_handler('/api/users', handle_users, [method(get)]).
true.

handle_users(Request) :-
    http_reply(Request, 200, 'application/json', '{"users":[]}').
```

http_get_request/2

Purpose: Extracts request details (method, path, headers, body) from an HTTP request object.

When to use: Use inside handler predicates to inspect incoming requests.

```
% Syntax: http_get_request(+Request, -Details)
handle(Request) :-
    http_get_request(Request, Details),
    member(method(Method), Details),
    member(path(Path), Details).
```

http_reply/4

Purpose: Sends an HTTP response with a status code, content type, and body.

When to use: Use inside handler predicates to send responses.

```
% Syntax: http_reply(+Request, +StatusCode, +ContentType, +Body)
handle_hello(Request) :-
    http_reply(Request, 200, 'text/plain', 'Hello, World!').
```

http_reply_json/3

Purpose: Sends a JSON HTTP response (convenience wrapper around http_reply/4).

When to use: Use when building JSON APIs to avoid specifying content type manually.

```
% Syntax: http_reply_json(+Request, +StatusCode, +JsonTerm)
handle_user(Request) :-
    http_reply_json(Request, 200, json([name='Alice', age=30])).
```

http_client_get/2

Purpose: Performs an HTTP GET request and unifies with the response body.

When to use: Use to fetch data from external HTTP services.

```
% Syntax: http_client_get(+URL, -Response)
?- http_client_get('http://example.com/api/data', Response).
Response = '{"key":"value"}'.
```

http_client_post/3

Purpose: Performs an HTTP POST request with a body and unifies with the response.

When to use: Use to send data to external HTTP services.

```
% Syntax: http_client_post(+URL, +Body, -Response)
?- http_client_post('http://example.com/api/data', '{"key":"value"}', Response).
Response = '{"status":"ok"}'.
```

http_open/3

Purpose: Opens an HTTP connection as a stream for reading.

When to use: Use for streaming large HTTP responses or when you need fine-grained control over reading.

```
% Syntax: http_open(+URL, -Stream, +Options)
?- http_open('http://example.com/data.csv', Stream, []),
    read_stream_to_codes(Stream, Codes),
    close(Stream).
```

url_encode/2

Purpose: Percent-encodes a string for use in URLs.

When to use: Use when constructing URLs with user-supplied parameters.

```
% Syntax: url_encode(+Plain, -Encoded)
?- url_encode('hello world', E).
E = 'hello%20world'.
```

url_decode/2

Purpose: Decodes a percent-encoded URL string.

When to use: Use when parsing URL parameters from incoming requests.

```
% Syntax: url_decode(+Encoded, -Plain)
?- url_decode('hello%20world', D).
D = 'hello world'.
```

27. Persistence Predicates

Persistence predicates allow saving and loading the Prolog knowledge base to/from files, including support for JSON interchange format, predicate-level export, and snapshot/restore for transactional workflows.

db_save/1

Purpose: Saves the entire knowledge base to a Prolog-format file.

When to use: Use to persist all facts and rules to disk so they can be reloaded in a later session.

```
% Syntax: db_save(+File)
?- assert(person(alice, 30)),
   assert(person(bob, 25)),
   db_save('mydata.pl').
true.

% The file mydata.pl now contains:
% person(alice, 30).
% person(bob, 25).
```

db_load/1

Purpose: Loads facts and rules from a previously saved Prolog-format file into the knowledge base.

When to use: Use to restore a knowledge base saved with db_save/1 or any standard Prolog source file.

```
% Syntax: db_load(+File)
?- db_load('mydata.pl').
true.

?- person(alice, Age).
Age = 30.
```

db_save_predicate/2

Purpose: Saves only the clauses for a specific predicate to a file.

When to use: Use when you want to export a subset of the knowledge base rather than the entire database.

```
% Syntax: db_save_predicate(+Predicate/Arity, +File)
?- db_save_predicate(person/2, 'people.pl').
true.

% Only person/2 clauses are written to people.pl
```

persist/1

Purpose: Marks a predicate for automatic persistence. Changes to this predicate are automatically saved.

When to use: Use for predicates that should survive across sessions without explicitly calling db_save/1.

```
% Syntax: persist(+Predicate/Arity)
?- persist(config/2).
true.

?- assert(config(theme, dark)).
true.
% config/2 changes are automatically persisted
```

unpersist/1

Purpose: Removes automatic persistence marking from a predicate.

When to use: Use when you no longer need automatic persistence for a predicate.

```
% Syntax: unpersist(+Predicate/Arity)
?- unpersist(config/2).
true.
% config/2 changes are no longer automatically persisted
```

db_export_json/1

Purpose: Exports the knowledge base to a JSON file.

When to use: Use for interoperability with other systems that consume JSON data.

```
% Syntax: db_export_json(+File)
?- assert(employee(john, engineering)),
   db_export_json('kb.json').
true.

% kb.json contains a JSON representation of all facts and rules
```

db_import_json/1

Purpose: Imports facts and rules from a JSON file into the knowledge base.

When to use: Use to load data exported by db_export_json/1 or generated by external tools.

```
% Syntax: db_import_json(+File)
?- db_import_json('kb.json').
true.

?- employee(john, Dept).
Dept = engineering.
```

db_snapshot/1

Purpose: Captures a named snapshot of the current knowledge base state.

When to use: Use before making bulk changes so you can roll back if needed.

```
% Syntax: db_snapshot(+Name)
?- db_snapshot(before_update).
true.

?- retractall(person(_, _)).
true.
% Knowledge base is now empty for person/2, but snapshot is saved
```

db_restore/1

Purpose: Restores the knowledge base to a previously captured snapshot.

When to use: Use to roll back changes when an operation fails or produces incorrect results.

```
% Syntax: db_restore(+Name)
?- db_restore(before_update).
true.
```

```
?- person(alice, Age).
Age = 30.
% Knowledge base restored to its state at snapshot time
```

db_clear/0

Purpose: Removes all user-defined facts and rules from the knowledge base.

When to use: Use to reset the knowledge base to an empty state, for example before loading fresh data.

```
% Syntax: db_clear
?- db_clear.
true.

?- person(_, _).
false.
% All user-defined clauses have been removed
```

```
<!-- START_CHANGE: ISS-2025-0179 - Add db_transaction/1, db_sync/0, db_batch_assert/1 -->
```

db_transaction/1

Purpose: Executes a goal within a transaction; rolls back changes if the goal fails or throws.

```
% Syntax: db_transaction(+Goal)
?- db_transaction((assert(account(alice, 100)), assert(account(bob, 200)))).
true.

% If the goal fails, no changes are committed
?- db_transaction((assert(temp(1)), fail)).
false.
% temp(1) was not asserted
```

db_sync/0

Purpose: Forces all pending persistence writes to be flushed to disk.

```
% Syntax: db_sync
?- db_sync.
true.
```

db_batch_assert/1

Purpose: Asserts a list of clauses in bulk, more efficiently than individual assert calls.

```
% Syntax: db_batch_assert(+ClauseList)
?- db_batch_assert([fact(a), fact(b), fact(c)]).
true.

?- fact(X).
X = a ; X = b ; X = c.
```

```
<!-- END_CHANGE: ISS-2025-0179 -->
```

28. Graph Algorithm Predicates

Graph algorithm predicates provide common graph operations including path finding, shortest path computation, connectivity analysis, topological sorting, minimum spanning trees, and cycle detection. Graphs are represented as lists of edges in the form `edge(From, To)` or `edge(From, To, Weight)` for weighted graphs.

graph_path/4

Purpose: Finds a path between two vertices in a graph.

When to use: Use to determine if two nodes are connected and to retrieve the connecting path.

```
% Syntax: graph_path(+Edges, +Start, +End, -Path)
?- Edges = [edge(a,b), edge(b,c), edge(c,d), edge(b,d)],
   graph_path(Edges, a, d, Path).
Path = [a, b, c, d] ;
Path = [a, b, d].
```

shortest_path/4

Purpose: Finds the shortest (minimum weight) path between two vertices in a weighted graph.

When to use: Use for route planning, network optimization, or any problem requiring minimum-cost paths.

```
% Syntax: shortest_path(+Edges, +Start, +End, -Path)
?- Edges = [edge(a,b,1), edge(b,c,2), edge(a,c,10)],
   shortest_path(Edges, a, c, Path).
Path = [a, b, c].
```

graph_connected/2

Purpose: Checks whether two vertices are connected in the graph (i.e., a path exists between them).

When to use: Use for simple reachability checks without needing the actual path.

```
% Syntax: graph_connected(+Edges, +Vertex-Vertex)
?- Edges = [edge(a,b), edge(b,c)],
   graph_connected(Edges, a-c).
true.

?- graph_connected(Edges, a-d).
false.
```

graph_vertices/2

Purpose: Extracts the set of all vertices from a graph's edge list.

When to use: Use to enumerate all nodes in a graph.

```
% Syntax: graph_vertices(+Edges, -Vertices)
?- graph_vertices([edge(a,b), edge(b,c), edge(c,a)], Vs).
Vs = [a, b, c].
```

graph_edges/2

Purpose: Extracts all edges from a graph representation, normalizing them to a uniform format.

When to use: Use to inspect or iterate over all edges in a graph.

```
% Syntax: graph_edges(+Graph, -Edges)
?- graph_edges([edge(a,b,1), edge(b,c,2)], Es).
Es = [edge(a, b, 1), edge(b, c, 2)].
```

graph_neighbors/3

Purpose: Finds all neighbors (adjacent vertices) of a given vertex.

When to use: Use to explore the local structure of a graph around a specific node.

```
% Syntax: graph_neighbors(+Edges, +Vertex, -Neighbors)
?- Edges = [edge(a,b), edge(a,c), edge(b,d)],
   graph_neighbors(Edges, a, Ns).
Ns = [b, c].
```

topological_sort/2

Purpose: Produces a topological ordering of vertices in a directed acyclic graph (DAG).

When to use: Use for dependency resolution, task scheduling, or build ordering.

```
% Syntax: topological_sort(+Edges, -Sorted)
?- Edges = [edge(a,b), edge(a,c), edge(b,d), edge(c,d)],
   topological_sort(Edges, Sorted).
Sorted = [a, c, b, d].
```

graph_components/2

Purpose: Finds all connected components in an undirected graph.

When to use: Use to identify clusters or disconnected subgraphs.

```
% Syntax: graph_components(+Edges, -Components)
?- Edges = [edge(a,b), edge(c,d)],
   graph_components(Edges, Cs).
Cs = [[a, b], [c, d]].
```

minimum_spanning_tree/2

Purpose: Computes the minimum spanning tree of a weighted undirected graph.

When to use: Use for network design problems where you need to connect all nodes at minimum total cost.

```
% Syntax: minimum_spanning_tree(+Edges, -MST)
?- Edges = [edge(a,b,1), edge(b,c,2), edge(a,c,3)],
   minimum_spanning_tree(Edges, MST).
MST = [edge(a, b, 1), edge(b, c, 2)].
```

graph_degree/3

Purpose: Computes the degree (number of incident edges) of a vertex.

When to use: Use to analyze the connectivity of individual nodes, find hubs, or identify leaf nodes.

```
% Syntax: graph_degree(+Edges, +Vertex, -Degree)
?- Edges = [edge(a,b), edge(a,c), edge(a,d)],
   graph_degree(Edges, a, D).
D = 3.
```

graph_has_cycle/1

Purpose: Checks whether a directed graph contains a cycle.

When to use: Use to validate that a dependency graph is a DAG before performing topological sort.

```
% Syntax: graph_has_cycle(+Edges)
?- graph_has_cycle([edge(a,b), edge(b,c), edge(c,a)]).
true.

?- graph_has_cycle([edge(a,b), edge(b,c)]).
false.
```

graph_reachable/3

Purpose: Finds all vertices reachable from a given starting vertex.

When to use: Use to compute the transitive closure from a single source node.

```
% Syntax: graph_reachable(+Edges, +Start, -Reachable)
?- Edges = [edge(a,b), edge(b,c), edge(b,d)],
   graph_reachable(Edges, a, Rs).
Rs = [b, c, d].
```

<!-- START_CHANGE: ISS-2025-0179 - Add graph_scc/2 -->

graph_scc/2

Purpose: Computes the strongly connected components (SCCs) of a directed graph using Tarjan's algorithm.

```
% Syntax: graph_scc(+Edges, -SCCs)
?- Edges = [edge(a,b), edge(b,c), edge(c,a), edge(d,e)],
   graph_scc(Edges, SCCs).
SCCs = [[a, b, c], [d], [e]].
```

<!-- END_CHANGE: ISS-2025-0179 -->

29. Java FFI Predicates

The Java Foreign Function Interface (FFI) predicates allow Prolog programs to interact with Java classes, objects, methods, fields, and arrays. This enables seamless interoperability between Prolog logic and Java libraries.

v4.6.0 (ISS-2025-0796, wave Q7): argument faults RAISE (they failed silently): an unbound argument is `instantiation_error`, a wrong type `type_error(atom|list|integer|java_object|java_array, Culprit)`, an unknown class / method / field / constructor `existence_error(class|method|field|constructor, Name)` (`static_field` for an instance field reached through a class), assigning a final field `permission_error(modify, final_field, Name)`, a negative array length `domain_error(not_less_than_zero, N)`, and an exception thrown by the invoked Java code `error(java_exception(ExceptionClassName), context(Name/Arity, Message))`. An array index outside the array still fails (as `arg/3` does), and `java_instanceof(null, C)` fails. Every context is `context(Name/Arity, Message)`.

Understanding Java FFI

The FFI bridges Prolog and Java by:

- Creating Java objects from Prolog
- Calling Java methods and accessing fields
- Working with Java arrays
- Converting between Java objects and Prolog terms

Java objects are referenced by opaque handles (atoms) that can be passed between FFI predicates.

java_new/3

Purpose: Create a new Java object by calling a constructor.

```
% java_new(+ClassName, +ArgsList, -ObjectRef)
?- java_new('java.util.ArrayList', [], Ref).
Ref = java_obj_1.

?- java_new('java.lang.StringBuilder', ['Hello'], Ref).
Ref = java_obj_2.
```

java_call/4

Purpose: Call a method on a Java object or class (for static methods).

```
% java_call(+ObjectRef, +MethodName, +ArgsList, -Result)
?- java_new('java.util.ArrayList', [], List),
   java_call(List, add, ['hello'], _),
   java_call(List, size, [], Size).
Size = 0.

% Static method call
?- java_call('java.lang.Math', max, [3, 7], Result).
Result = 7.
```

java_get_field/3

Purpose: Get the value of a field on a Java object or class.

```
% java_get_field(+ObjectOrClass, +FieldName, -Value)
?- java_get_field('java.lang.Integer', 'MAX_VALUE', Val).
```

```
Val = 2147483647.
```

java_set_field/3

Purpose: Set the value of a field on a Java object.

```
% java_set_field(+ObjectRef, +FieldName, +Value)
?- java_new('MyClass', [], Obj),
   java_set_field(Obj, count, 42).
```

java_instanceof/2

Purpose: Check if a Java object is an instance of a given class.

```
% java_instanceof(+ObjectRef, +ClassName)
?- java_new('java.util.ArrayList', [], Obj),
   java_instanceof(Obj, 'java.util.List').
true.
```

java_class/2

Purpose: Get the class name of a Java object.

```
% java_class(+ObjectRef, -ClassName)
?- java_new('java.util.HashMap', [], Obj),
   java_class(Obj, Class).
Class = 'java.util.HashMap'.
```

java_array_new/3

Purpose: Create a new Java array of a given type and size.

```
% java_array_new(+ElementType, +Size, -ArrayRef)
?- java_array_new(int, 5, Arr).
Arr = java_arr_1.
```

java_array_get/3

Purpose: Get an element from a Java array by index.

```
% java_array_get(+ArrayRef, +Index, -Value)
?- java_array_get(Arr, 0, Val).
Val = 0.
```

java_array_set/3

Purpose: Set an element in a Java array at a given index.

```
% java_array_set(+ArrayRef, +Index, +Value)
?- java_array_set(Arr, 0, 42).
true.
```

java_array_length/2

Purpose: Get the length of a Java array.

```
% java_array_length(+ArrayRef, -Length)
?- java_array_length(Arr, Len).
Len = 5.
```

java_to_term/2

Purpose: Convert a Java object to a Prolog term representation.

```
% java_to_term(+ObjectRef, -Term)
?- java_new('java.lang.Integer', [42], Obj),
   java_to_term(Obj, Term).
Term = 42.
```

java_from_term/2

Purpose: Convert a Prolog term to a Java object.

```
% java_from_term(+Term, -ObjectRef)
?- java_from_term(hello, Obj),
   java_class(Obj, Class).
Class = 'java.lang.String'.
```

<!-- START_CHANGE: ISS-2025-0179 - Add java_release_ref/1 and java_gc/0 -->

java_release_ref/1

Purpose: Releases a Java object reference, allowing it to be garbage collected.

```
% Syntax: java_release_ref(+ObjectRef)
?- java_new('java.util.ArrayList', [], Ref),
   java_release_ref(Ref).
true.
```

java_gc/0

Purpose: Requests Java garbage collection to free unreferenced objects.

```
% Syntax: java_gc
?- java_gc.
true.
```

<!-- END_CHANGE: ISS-2025-0179 -->

<!-- START_CHANGE: ISS-2025-0466..0471 - engine v4 wave W6: the module system -->

30. Concurrent Execution Predicates

Parallel goal execution using Java threads.

Since v4.0.0 (ISS-2025-0480) every goal here runs on its own resolution machine over the same engine: the clause store is shared and thread-safe, the goal is copied so no variable is shared between workers, and interrupting the parent cancels the workers. Goals still have to be independent - two workers asserting to the same predicate see each other's writes.

v4.5.0 (ISS-2025-0622/0623/0624): the workers draw from the caller's ONE inference budget; what a worker raises reaches the caller unchanged - `catch(concurrent_and([throw(foo)], []), foo, true)` catches, and the budget or a Stop is never turned into a catchable ball; there is no 60-second cap; argument faults are ISO errors; `concurrent/3` and `concurrent_maplist/2..4` unify the bindings the goals made back into the caller (SWI).

concurrent/3

Execute a list of goals using at most N worker threads. All must succeed.

```
concurrent(4, [goal1, goal2, goal3], [])
```

concurrent_maplist/2

Like `maplist/2` but parallel. `call(Goal, Elem)` for each element.

```
concurrent_maplist(is_positive, [1, 2, 3, 4, 5])
```

concurrent_maplist/3

Like `maplist/3` but parallel. `call(Goal, Elem, Result)` collecting results in order.

```
concurrent_maplist(square, [1,2,3,4], [1,4,9,16])
```

(Before v4.0.0 this was registered under the uncallable name `concurrent_maplist3` and `concurrent_maplist/3` raised an arity error - ISS-2025-0480.)

concurrent_maplist/4

Parallel `maplist` with two input lists.

```
concurrent_maplist(add, [1,2,3], [10,20,30], [11,22,33])
```

(Before v4.0.0: `concurrent_maplist4`, likewise uncallable.)

first_solution/3

Run goals in parallel, return bindings from the first to succeed. OR-parallelism.

```
first_solution(X, [search_db1(X), search_db2(X), search_db3(X)], [])
```

concurrent_and/2

AND-parallelism: all goals must succeed.

```
concurrent_and([check1, check2, check3], [])
```

concurrent_or/2

OR-parallelism: returns 1-based index of the first goal to succeed.

```
concurrent_or([strategy1, strategy2], WinnerIndex)
```

concurrent_forall/2, concurrent_forall/3

v4.5.0 (ISS-2025-0634). `concurrent_forall(:Cond, :Action [, [threads(N)])`: runs Action for every solution of Cond on up to N workers (default: the number of processors); succeeds iff every Action succeeds; binds nothing.

```
?- concurrent_forall(between(1, 100, X), check(X), [threads(4)]).
```

31. Module System

Version note. Everything in this section arrived with engine v4 (v3.13.0, the default since 4.0.0). The older `ModuleManager` behaviour - `Module:Goal` reaching only user-defined clauses, no library modules, no `autoload`, `meta_predicate/1` recorded but never consulted, no `current_module/1` - belonged to the engine deleted in 4.1.0.

The three kinds of module

Module	What it holds
user	The default. It is the flat knowledge base - the clauses <code>consult/1</code> , <code>assertz/1</code> and <code>listing/1</code> see.
system	The built-in predicates (both the v4 natives and the ~400 registry ones). It holds no clauses. <code>Prolog.enableSafeMode()</code> removes the host-touching built-ins from it.
library modules	<code>lists</code> , <code>apply</code> , <code>pairs</code> , <code>coroutining</code> - written in Prolog, shipped as classpath resources under <code>prelude/</code> , and autoloaded by predicate indicator : a module is parsed the first time one of its predicates is referenced.

Resolution order

For an unqualified call to `f/n` from module `M`:

```
M's own clauses -> M's imports (in import order, exported predicates only)
                  -> user (the flat knowledge base)
                  -> the autoloaded library modules
                  -> system (the built-ins)
```

Two practical consequences:

- **A definition in the calling context wins over the library.** A program that defines its own `partition/4` gets its own, while `apply:partition/4` stays reachable under its qualified name. A definition inside module `M` overrides only for `M` and the modules that import it.
- **Built-ins are still not redefinable.** `system` is consulted last in the list above, but the machine checks the built-in tables *before* the clause layers for compatibility, and `consult/1` refuses a clause whose head is a registered built-in (Cannot redefine built-in predicate `f/n`). The library layer is the documented way to override a library predicate.

Module:Goal

```
?- lists:append([1], [2], L).      % a library predicate
L = [1, 2].

?- system:atom_length(abc, N).    % a built-in, explicitly
N = 3.

?- user:my_fact(X).               % the flat knowledge base, explicitly
```

- The **innermost** qualification of a nested `a:b:Goal` wins, so `user:lists:append(...)` runs in `lists`.
- **Export enforcement:** a module that *defines* `f/n` answers a qualified call only if it also exports it. `:- module(secret, []). followed by hidden(42).` makes `secret:hidden(X)` fail.
- A module that does **not** define `f/n` falls through to the ordinary resolution in its own context, which is why `lists:length(L, N)` (a native) and `othermodule:my_user_fact(X)` (a user predicate) both work.
- An unknown module name behaves like `user:.`

Directives

Directive	Effect
<code>:- module(Name, [f/1, g/2]).</code>	Declare the current module and its export list. An empty list exports nothing.
<code>:- use_module(Name).</code>	Import every predicate <code>Name</code> exports into the current module.
<code>:- use_module(Name, [f/1]).</code>	Import only the listed predicates (v4.6.0: really - <code>f/1</code> as <code>g</code> renames, except <code>(L)</code> excludes; it imported everything before).
<code>:- use_module(library(X)).</code>	Accepted for the library modules; they autoload anyway, so it is a no-op.
<code>:- meta_predicate(Spec).</code>	See below.

meta_predicate/1 (directive)

`:- meta_predicate(maplist(2, ?, ?)).` says that argument 1 of `maplist/3` is a goal that will be called with 2 extra arguments. When a predicate with such a declaration is called from module `C`, its **module-sensitive** arguments are qualified with `C` before the clause head is unified, so the callee's call/`N` runs them in the caller's context.

Argument specifiers: 0-9 (a goal called with that many extra arguments), `:` (a module-sensitive term), `^`, `//`; `+`, `-`, `?` and anything else are ordinary arguments.

```
:- module(m1, [go/1]).
helper(from_m1).
mk(_, R) :- helper(R).
go(L) :- maplist(mk, [x], L).      % mk/2 resolves in m1, not in apply and not in user
```

current_module/1

```
?- current_module(lists).
true.

?- current_module(M).          % enumerates: user first, then the rest, sorted
M = user ;
M = apply ;
M = coroutining ;
...
```

predicate_property/2 module properties

v4.5.0 (ISS-2025-0610/0611): for a bound head the properties come from the registry, the native table, the KnowledgeBase and the module owner: `built_in`, `system`, `defined`, `visible`, `dynamic` or `static`, `number_of_clauses(N)`, plus the module ones; an undefined predicate FAILS (SWI) instead of answering undefined. A qualified call `M:G` runs `G` in `M` whether or not `M` exports it (export governs import only, SWI).

In addition to `built_in`, `dynamic`, `static` and `defined`:

Property	Meaning
<code>defined_in(Module)</code>	The module whose clauses would answer the call from the current context.
<code>exported</code>	That module exports the predicate.
<code>imported_from(Module)</code>	The predicate is visible here but defined elsewhere.

```
?- predicate_property(append(_, _, _), imported_from(lists)).
true.
```

What the library modules contain

Module	Predicates
lists	member/2, memberchk/2, append/3, select/3, selectchk/3, nth0/3, nth1/3, last/2, reverse/2, plus length/2, msort/2, sort/2, sort/4, sum_list/2, sumlist/2, numlist/3, permutation/2, max_list/2, min_list/2, subtract/3, intersection/3, union/3
apply	maplist/2..7, foldl/4..7, include/3, exclude/3, partition/4, partition/5
pairs	pairs_keys_values/3, pairs_keys/2, pairs_values/2
coroutining	freeze/2, frozen/2, when/2, dif/2, ?=/2
clpfd	sum/3, scalar_product/4, element/3, tuples_in/2, global_cardinality/2, transpose/2 (v4.5.0; the operator-named constraints and the labeling/reflection predicates are natives)

Most of `lists` is implemented as a native generator for speed (on a 1 000 000-element list a Prolog clause walk pushes one choice point per element); `member/2` and `append/3` additionally exist as the two-clause Prolog definitions of module `lists`, which is what the qualified form runs. The two are observationally identical.

Behaviour differences from the default engine

- `append(X, Y, Z)` with all three arguments open **enumerates** on v4 (`X = [], [], [], ...` lazily) where the default engine stops at the single standard solution. A program that relied on that termination will loop.
- `member(X, PartialList)` **extends** the open tail on v4, and `memberchk(a, L)` binds `L = [a|_]`; both fail on the default engine.
- `memberchk/2` and `current_module/1` do not exist at all on the default engine.

<!-- END_CHANGE: ISS-2025-0466..0471 -->

32. String predicates

JProlog strings ("text" with the default `double_quotes=string` flag) are a distinct atomic type. All `string_*` predicates also accept atoms and numbers as text input, and the `atom_*` predicates accept strings, so the two families interoperate freely.

string/1

Purpose: Type check - succeeds if the argument is a string object.

```
?- string("abc").
true.

?- string(abc).
false.
```

string_length/2

Purpose: `string_length(+Text, -Length)` - number of characters in a string (or atom/number).

```
?- string_length("hello", N).
N = 5.

?- string_length(hello, N).
N = 5.

?- string_length(123, N).
N = 3.
```

string_concat/3

Purpose: `string_concat(?S1, ?S2, ?S3)` - concatenation of strings; with `S3` bound and the others unbound it enumerates every split on backtracking. Any atomic text is accepted (`string_concat(1, 2, S)` gives "12"); with too little bound it raises `instantiation_error` (4.5.0).

```
?- string_concat("ab", "cd", S).
S = "abcd".

?- string_concat(X, Y, "ab").
X = "", Y = "ab" ;
X = "a", Y = "b" ;
X = "ab", Y = "".
```

string_chars/2, string_codes/2

Purpose: Convert between a string and its list of characters / character codes; both directions work.

```
?- string_chars("hi", Cs), string_codes("hi", Codes).
Cs = [h, i], Codes = [104, 105].

?- string_chars(S, [o, k]).
S = "ok".
```

string_code/3

Purpose: `string_code(+Index, +String, -Code)` - code of the character at a 1-based index.

```
?- string_code(2, "abc", C).
C = 98.
```

sub_string/5

Purpose: `sub_string(+String, ?Before, ?Length, ?After, ?Sub)` - enumerate or test substrings, exactly like `sub_atom/5` but producing strings. `String` and `Sub` may be any text (atom, string, number); positions count code points; a negative position fails.

```
?- sub_string("hello world", 6, 5, _, S).
S = "world".

?- sub_string("abc", B, 1, A, S).
B = 0, A = 2, S = "a" ;
B = 1, A = 1, S = "b" ;
B = 2, A = 0, S = "c".
```

string_to_atom/2, atom_string/2

Purpose: Convert between strings and atoms in either direction (`string_to_atom(?String, ?Atom)`, `atom_string(?Atom, ?String)`). At least one argument must be bound. `atom_string/2` accepts any atomic text on either side: `atom_string(42, S)` gives `"42"`, `atom_string(A, 42)` gives `'42'`.

```
?- atom_string(hello, S), string_to_atom("world", A).
S = "hello", A = world.

?- atom_string(A, "42").
A = '42'.
```

number_string/2

Purpose: `number_string(?Number, ?String)` - parse a string as a number (leading/trailing whitespace allowed) or render a number as a string. A float keeps its type.

```
?- number_string(N, " 3.0 "), number_string(42, S).
N = 3.0, S = "42".

?- number_string(N, "abc").
false.
```

Unlike `number_codes/2`, which raises `syntax_error(illegal_number)`, `number_string/2` fails silently on text that is not a number (SWI-Prolog behaviour).

```
?- catch(number_codes(N, "abc"), error(E, _), true).
E = syntax_error(illegal_number).
```

split_string/4

Purpose: `split_string(+String, +SepChars, +PadChars, -SubStrings)` - split on any separator character, stripping pad characters from each part.

```
?- split_string("a,b,,c", ",", "", P).
P = ["a", "b", "", "c"].

?- split_string(" key = value ", "=", " ", P).
P = ["key", "value"].
```

number_codes/2, atom_to_number/2, number_to_atom/2, to_codes/2

Purpose: `number_codes(?Number, ?Codes)` is the ISO conversion between a number and its code list (accepting `0x`, `0o`, `0b`, `0'c` notations and raising `syntax_error` on bad text). `atom_to_number(+Atom, -Number)` and `number_to_atom(+Number, -Atom)` are convenience aliases of `atom_number/2`; `to_codes(+Text, -Codes)` turns any atom, string or number into a code list.

```
?- number_codes(N, "0x1F"), number_codes(2.5, C).
N = 31, C = [50, 46, 53].

?- atom_to_number('12', N), number_to_atom(3.5, A), to_codes(hi, Cs).
N = 12, A = '3.5', Cs = [104, 105].
```

upcase_atom/2, downcase_atom/2

Purpose: Case conversion of an atom (or string) into an atom.

```
?- upcase_atom('Hello', U), downcase_atom('Hello', D).
U = 'HELLO', D = hello.
```

33. Stream predicates

Streams are created by `open/3,4` and referred to by the stream term it returns or by an alias (`user_input`, `user_output`, `user_error`, or an alias given with `open/4`). The predicates below complete the file/stream chapter of the ISO core.

current_input/1, current_output/1

Purpose: Unify the argument with the current input / output stream. The current output stream is per thread, so a query running in the IDE or in a worker thread sees its own console.

```
?- current_output(S), write(S, hello), nl(S).
hello
S = user_output.
```

set_input/1, set_output/1

Purpose: Make a stream (or alias) the current input / output for the subsequent `read/1`, `write/1`, `nl/0` ... calls that take no explicit stream.

```
capture(File, Goal) :-
    open(File, write, S),
    current_output(Old),
    set_output(S),
    call(Goal),
    set_output(Old),
    close(S).

?- capture('out.txt', (write(hello), nl)).
true.
```

flush_output/0, flush_output/1

Purpose: Force buffered output to be written (to the current output stream, or to the given one). Useful before a blocking read or a long computation.

```
?- write('Name? '), flush_output, read(Name).
```

close/1, close/2

Purpose: Close a stream. `close(Stream, Options)` accepts `force(true)` to ignore errors while closing. Closing `user_input`/`user_output` is a no-op.

```
?- open('data.txt', read, S), read(S, T), close(S, [force(true)]).
```

stream_property/2

Purpose: `stream_property(?Stream, ?Property)` - enumerate stream properties. Both arguments may be unbound, so the predicate enumerates every open stream and every property of it. The complete ISO set is reported: `file_name(F)`, `mode(M)`, `input`, `output`, `alias(A)`, `position(P)`, `end_of_stream(E)`, `eof_action(A)`, `reposition(B)`, `type(T)`, `text/binary`, `encoding(E)` and `line_count(N)`.

```
?- stream_property(user_error, alias(A)).
A = user_error.
```

```
?- open('data.txt', read, S), stream_property(S, mode(M)), stream_property(S, type(T)), close(S).
M = read, T = text.
```

An unbound first argument is bound to the **canonical** stream term ('\$stream'(N)), never to an alias: `stream_property(S, alias(user_error))` gives `S = '$stream'(2)`. Enumerating never blocks - `end_of_stream` is reported as not for a stream that cannot be repositioned (stdin, a socket), because deciding it there would mean waiting for input.

current_stream/3, set_stream/2

Purpose: `current_stream(?File, ?Mode, ?Stream)` enumerates the open streams with their file name and mode. `set_stream(+Stream, +Property)` changes a property of an open stream; `alias(A)` adds an alias that every stream argument then accepts.

```
?- open('data.txt', read, S), current_stream(F, M, S), close(S).
F = 'data.txt', M = read.

?- open('data.txt', read, S), set_stream(S, alias(input_file)),
   get_char(input_file, C), close(S).
C = h.
```

stream_position/2, set_stream_position/2, seek/4

Purpose: `stream_position(+Stream, -Pos)` reads the byte position of a file stream; `set_stream_position(+Stream, +Pos)` repositions it (permission_error(reposition, stream, S) for streams that cannot be repositioned); `seek(+Stream, +Offset, +Method, -NewPos)` moves relative to bof, current or eof. Since v3.14.0 a **text** stream is decoded one code point at a time through its own decoder, so repositioning really does change what the next `get_char/2` reads - it is no longer binary-only.

```
% data.txt contains the text "hello"
?- open('data.txt', read, S), get_char(S, C1), stream_position(S, P),
   get_char(S, C2), set_stream_position(S, P), get_char(S, C3), close(S).
C1 = h, P = 1, C2 = e, C3 = e.

?- open('data.txt', read, S, [type(binary)]), get_byte(S, B1),
   seek(S, 0, bof, N), get_byte(S, B2), close(S).
B1 = 104, N = 0, B2 = 104.
```

character_count/2, line_count/2, line_position/2, stream_position_data/3

Purpose: `character_count(+Stream, -N)`, `line_count(+Stream, -N)` and `line_position(+Stream, -N)` report how far a stream has been read or written: the number of characters consumed, the current line number (1-based) and the column within it (0-based). `stream_position_data(+Field, +Position, -Value)` pulls the same three fields out of the opaque position term that `stream_property(S, position(P))` yields; Field is `char_count`, `line_count` or `line_position`.

```
% data.txt contains the text "hello"
?- open('data.txt', read, S), get_char(S, _), get_char(S, _),
   character_count(S, N), line_count(S, L), line_position(S, C), close(S).
N = 2, L = 1, C = 2.

?- open('data.txt', read, S), get_char(S, _),
   stream_property(S, position(P)), stream_position_data(char_count, P, N), close(S).
N = 1.
```

peek_char/1,2, peek_code/1,2, peek_byte/1,2

Purpose: Look at the next character / code / byte of the input without consuming it. Return `end_of_file` (or -1 for codes and bytes) at the end of the stream.

```
skip_spaces(S) :- peek_char(S, C), C == ' ', !, get_char(S, _), skip_spaces(S).
skip_spaces(_).
```

put_byte/1,2, get_byte/1,2

Purpose: Binary I/O on streams opened with `type(binary)`.

```
?- open('raw.bin', write, S, [type(binary)]), put_byte(S, 255), put_byte(S, 0), close(S),
    open('raw.bin', read, R, [type(binary)]), get_byte(R, B1), get_byte(R, B2), close(R).
B1 = 255, B2 = 0.
```

at_end_of_stream/0,1

Purpose: Succeeds when the (current or given) input stream has no more data.

```
read_all(S, []) :- at_end_of_stream(S), !.
read_all(S, [C|Cs]) :- get_char(S, C), read_all(S, Cs).
```

34. Global variables

Global variables are a per-engine key/value store. `nb_setval/2` copies the value and keeps it across backtracking; `b_setval/2` stores the value so that it is restored on backtracking.

nb_setval/2, nb_getval/2

Purpose: Set / read a non-backtrackable global variable. Reading an unset name raises `existence_error(variable, Name)`.

```
count_solutions(Goal, N) :-
    nb_setval(cnt, 0),
    ( call(Goal), nb_getval(cnt, C0), C is C0 + 1, nb_setval(cnt, C), fail ; true ),
    nb_getval(cnt, N).

?- count_solutions(member(_, [a, b, c]), N).
N = 3.
```

b_setval/2, b_getval/2

Purpose: Backtrackable global variables: the previous value is restored when execution backtracks over the `b_setval/2` call.

```
?- b_setval(v, 1), ( b_setval(v, 2), fail ; true ), b_getval(v, V).
V = 1.
```

nb_current/2, nb_delete/1

Purpose: `nb_current(?Name, ?Value)` enumerates the defined global variables; `nb_delete(+Name)` removes one (silently succeeds if it does not exist).

```
?- nb_setval(a, 1), nb_setval(b, two), findall(K-V, nb_current(K, V), L), nb_delete(a).
L = [a-1, b-two].
```

35. Attributed variables and coroutining

Attributed variables carry extra data that is consulted when the variable is bound. They are the basis of `freeze/2`, `dif/2`, `when/2` and of the CLP(FD) solver. Attributes survive between top-level queries for the same variable name in the same engine.

`put_attr/3`, `get_attr/3`, `del_attr/2`, `attvar/1`

Purpose: `put_attr(+Var, +Module, +Value)` attaches (or replaces) the attribute `Module` of an unbound variable; `get_attr(+Var, +Module, -Value)` reads it; `del_attr(+Var, +Module)` removes it; `attvar(@Term)` succeeds if the term is a variable with at least one attribute.

```
?- put_attr(X, colour, red), get_attr(X, colour, C), attvar(X).
C = red.

?- put_attr(X, colour, red), del_attr(X, colour), attvar(X).
false.
```

`freeze/2`

Purpose: `freeze(?Var, :Goal)` - delay `Goal` until `Var` is bound; if `Var` is already bound the goal runs immediately. The goal runs at the moment of binding, inside the unification.

```
?- freeze(X, (write(bound(X)), nl)), X = 42.
bound(42)
X = 42.

?- freeze(X, X > 3), member(X, [1, 5, 2, 7]).
X = 5 ;
X = 7.
```

`dif/2`

Purpose: `dif(?X, ?Y)` - constrains two terms to be different; fails at once when they are identical, succeeds when they can never unify, and otherwise suspends until the decision can be made.

```
?- dif(X, a), member(X, [a, b, c]).
X = b ;
X = c.

?- dif(f(X), f(Y)), X = 1, Y = 1.
false.
```

`when/2`

Purpose: `when(+Condition, :Goal)` - run `Goal` as soon as `Condition` holds. Conditions: `nonvar(X)`, `ground(T)`, `?(X, Y)` (decidable comparison), and conjunctions/disjunctions of these.

```
?- when(ground(X-Y), (Z is X + Y, write(sum(Z)), nl)), X = 1, Y = 2.
sum(3)
X = 1, Y = 2, Z = 3.

?- when((nonvar(A) ; nonvar(B)), write(ready)), B = 1.
ready
B = 1.
```

A woken goal runs in the query's own binding context, so the bindings it makes (`Z` above) are ordinary bindings, it is traced like any other goal, it is charged to the inference budget, and it can throw into the

enclosing `catch/3`. (Before 4.0.0 a `when/2`-woken goal's bindings were lost on the then-default engine and only its side effects were visible; that is fixed.)

frozen/2, ?=/2

Purpose: `frozen(?Var, -Goal)` returns the conjunction of goals currently delayed on `Var`, or `true` when there are none. `?=(X, Y)` succeeds when the comparison of `X` and `Y` is already decidable - that is, when they are identical or can never unify.

```
?- freeze(X, writeln(hi)), frozen(X, G).
G = writeln(hi).

?- frozen(Y, G).
G = true.

?- ?=(a, b).
true.

?- ?=(X, Y).
false.
```

term_attvars/2, copy_term/3, unifiable/3

Purpose: `term_attvars(+Term, -Vars)` collects the attributed variables reachable from `Term`. `copy_term(+Term, -Copy, -Attributes)` copies `Term` without its attributes and returns them as a list of goals that would restore them. `unifiable(@X, @Y, -Unifier)` returns the list of `Var = Value` bindings that unifying `X` and `Y` would make, **without** making them. For a CLP(FD) variable the goals are its domain and the constraints still alive on it, in SWI-Prolog's printed forms (v4.6.0; the same goals the toplevel prints after an answer).

```
?- put_attr(X, mymod, 1), term_attvars(f(X, Y), Vs).
Vs = [X].

?- put_attr(X, mymod, 1), copy_term(X, Y, Attrs).
Attrs = [put_attr(Y, mymod, 1)].

?- X #> Y, copy_term(X-Y, C, Gs).
C = _A-_B, Gs = [_B#=<_A+ -1].

?- unifiable(f(X, b), f(a, Y), U).
U = [X = a, Y = b].
```

36. More list and pair predicates

msort/2

Purpose: `msort(+List, -Sorted)` - sort in the standard order of terms **without** removing duplicates (`sort/2` removes them).

```
?- msort([c, a, b, a], L).
L = [a, a, b, c].
```

predsort/3

Purpose: `predsort(:Pred, +List, -Sorted)` - sort with a user comparison called as `call(Pred, Order, A, B)` where `Order` is `<`, `=` or `>`. Elements that compare `=` are merged into one.

```
by_length(0, A, B) :- atom_length(A, LA), atom_length(B, LB), compare(0, LA-A, LB-B).

?- predsort(by_length, [ccc, a, bb], L).
L = [a, bb, ccc].
```

min_list/2, max_list/2, sum_list/2, sumlist/2

Purpose: Minimum, maximum and sum of a list of numbers (`sumlist/2` is an alias of `sum_list/2`). The result keeps the type of the operands: a list containing a float sums to a float. As in SWI-Prolog the elements are evaluated, so a non-number raises `type_error(evaluable, F/N)`, and integer sums are exact (4.5.0).

```
?- min_list([3, 1.5, 2], Min), max_list([3, 1.5, 2], Max), sum_list([1, 2, 3], S), sum_list([1.5, 1.5], F).
Min = 1.5, Max = 3, S = 6, F = 3.0.
```

nth0/3, nth1/3

Purpose: Element at a 0-based / 1-based index; with the index unbound they enumerate. On a partial list a bound index extends the list (`nth0(1, L, X)` gives `L = [_ , x|_]`).

```
?- nth1(2, [a, b, c], X), nth0(I, [a, b, c], c).
X = b, I = 2.
```

partition/4, partition/5

Purpose: `partition(:Pred, +List, ?Included, ?Excluded)` splits a list in one pass; `partition(:Pred, +List, ?Less, ?Equal, ?Greater)` uses `call(Pred, X, Order)` with `Order` one of `<`, `=`, `>`. Both are library predicates written in Prolog, so **a user definition of `partition/4` overrides them** (many textbook quicksorts define their own, taking a pivot rather than a goal).

```
?- partition([X]>>(X > 2), [1, 2, 3, 4], Big, Small).
Big = [3, 4], Small = [1, 2].
```

exclude/3, include/3

Purpose: Filter a list with a goal: `include(:Goal, +List, -Kept)` keeps the elements for which `call(Goal, Elem)` succeeds, `exclude/3` keeps the others. (`partition/4` is deliberately not provided as a built-in so that programs may define their own `partition/4`, as quicksort examples commonly do.)

```
even(X) :- 0 is X mod 2.
```

```
?- exclude(even, [1, 2, 3, 4], Odd), include(even, [1, 2, 3, 4], Even).
Odd = [1, 3], Even = [2, 4].
```

pairs_keys_values/3, pairs_keys/2, pairs_values/2

Purpose: Convert between a list of Key-Value pairs and its keys / values.

```
?- pairs_keys_values(P, [a, b], [1, 2]), pairs_keys(P, K), pairs_values(P, V).
P = [a-1, b-2], K = [a, b], V = [1, 2].
```

37. Term inspection and modification

put_code/1,2

Purpose: `put_code(+Code)` writes the character with the given code to the current output; `put_code(+Stream, +Code)` writes it to `Stream`. The counterpart of `get_code/1, 2`.

```
?- put_code(0'h), put_code(0'i), nl.
hi
true.
```

v4.2.0: the two-argument form actually writes to `Stream`. Before, the arity entry existed but the implementation raised "put_code/1 requires exactly 1 argument".

unify_with_occurs_check/2

Purpose: ISO unification **with** the occurs check, whatever the `occurs_check` flag says: it fails rather than building a cyclic term.

```
?- unify_with_occurs_check(X, f(Y)).
X = f(Y).

?- unify_with_occurs_check(X, f(X)).
false.
```

setarg/3

Purpose: `setarg(+Index, +Term, +Value)` - destructively replace the `Index`-th argument of a compound term; the change is undone on backtracking.

```
?- T = f(a, b), setarg(1, T, z).
T = f(z, b).

?- T = f(a), ( setarg(1, T, z), fail ; true ).
T = f(a).
```

nb_setarg/3

Purpose: `nb_setarg(+N, +Compound, +Value)` replaces the `N`-th argument like `setarg/3`, but the change is **not undone on backtracking**. Used to accumulate a result across a failure-driven loop.

```
?- T = counter(0), forall(member(_, [a,b,c]),
    (arg(1, T, C), C1 is C + 1, nb_setarg(1, T, C1))), arg(1, T, N).
T = counter(3), N = 3.
```

numbervars/3, number_vars/3

Purpose: `numbervars(+Term, +Start, -End)` binds every free variable of `Term` to `'$VAR'(N)` with consecutive `N` from `Start`; `End` is the next free number. `write/1` and `print/1` render `'$VAR'(0)` as `A`, `'$VAR'(1)` as `B`, and so on. `number_vars/3` is an alias.

```
?- T = f(X, Y, X), numbervars(T, 0, End), print(T), nl.
f(A, B, A)
End = 2.
```

must_be/2

Purpose: `must_be(+Type, @Value)` - succeed if `Value` has the type, otherwise throw the appropriate ISO error (`instantiation_error` for an unbound value when the type requires one, `type_error(Type, Value)` or `domain_error(Type, Value)` otherwise). Types: `atom`, `atomic`, `integer`, `float`, `number`, `callable`, `compound`, `var`, `nonvar`, `ground`, `list`, `boolean`, `positive_integer`, `nonneg`.

```
?- must_be(integer, 3).
true.

?- catch(must_be(positive_integer, 0), error(E, _), true).
E = type_error(positive_integer, 0).
```

simple/1, partial_list/1, rational/1, acyclic_term/1

Purpose: Additional type tests: `simple/1` is true for atomic terms (`atom`, `number` or `string`); `partial_list/1` for a list whose tail is unbound; `rational/1` for a rational number term; `acyclic_term/1` succeeds for terms without cycles.

```
?- simple(abc), partial_list([a, b|_]), acyclic_term(f(X)).
true.

?- partial_list([a, b]).
false.
```

38. Clause inspection and operators

clause/2

Purpose: `clause(+Head, ?Body)` - enumerate the clauses of a user predicate; facts have body true. Built-in predicates raise `permission_error(access, private_procedure, PI)`.

```
grandparent(GP, GC) :- parent(GP, P), parent(P, GC).

?- clause(grandparent(A, B), Body).
Body = (parent(A, _P), parent(_P, B)).
```

predicate_property/2

Purpose: `predicate_property(:Head, ?Property)` - properties of a predicate: `built_in`, `dynamic`, `static`, `defined`, `undefined`. With `Property` unbound it enumerates them.

```
?- predicate_property(append(_, _, _), built_in).
true.

?- assertz(fact(1)), predicate_property(fact(_), dynamic).
true.
```

op/3, current_op/3

Purpose: `op(+Priority, +Type, +Name)` defines (priority 1..1200) or removes (priority 0) an operator; `Name` may be a list of atoms. `current_op(?Priority, ?Type, ?Name)` enumerates the active operators. Part III lists the default table.

The operator store belongs to the Prolog instance: two engines in one JVM do not see each other's operators, and an `op/3` inside a module file is local to that module for `current_op/3`. Since 4.5.0 an operator definition is permanent (ISO, SWI): `op(700, xfx, tmp), fail ; true` leaves `tmp` defined (4.1.0 to 4.4.0 undid it on backtracking).

Since 4.4.0 both raise the ISO 13211-1 8.14.3.3 / 8.14.4.3 error terms (ISS-2025-0504): an unbound argument is `instantiation_error`; a non-integer priority `type_error(integer, P)`; a non-atom specifier `type_error(atom, T)`; a name that is neither an atom nor a list of atoms `type_error(list, N)` (a non-atom element is `type_error(atom, E)`); a priority outside 0..1200 `domain_error(operator_priority, P)`; an unknown specifier `domain_error(operator_specifier, T)`; `'` permission `permission_error(modify, operator, ',')`; and `'|'` outside its ISO window (priority 0, or an infix specifier with priority ≥ 1001) `permission_error(create, operator, '|')`. `current_op/3` raises the type and domain errors for a bound argument of the wrong shape instead of failing.

```
?- op(700, xfx, is_bigger), X =.. [is_bigger, elephant, mouse].
X = elephant is_bigger mouse.

?- findall(P-T, current_op(P, T, mod), L).
L = [400-yfx].

?- catch(op(1300, xfx, too_big), E, true).
E = error(domain_error(operator_priority, 1300), op/3).

?- catch(op(700, xfx, '|'), E, true).
E = error(permission_error(modify, operator, '|'), op/3).
```

char_conversion/2, current_char_conversion/2

Purpose: ISO character conversion table applied while reading terms when the `char_conversion` flag is true. `current_char_conversion(?In, ?Out)` enumerates the active mappings - the declared ones first, then the identity mapping of every other printable ASCII character. The table belongs to the Prolog instance, and a conversion declared in a branch that later fails is undone. `char_conversion(C, C)` removes the mapping for C.

Since 4.4.0 both raise ISO 8.14.5.3: an unbound argument is `instantiation_error` and anything that is not a one-character atom is `representation_error(character)` (ISS-2025-0504).

```
?- char_conversion(a, b), current_char_conversion(a, X).
X = b.

?- catch(char_conversion(ab, a), E, true).
E = error(representation_error(character), char_conversion/2).
```

code_type/2, char_type/2

Purpose: Classify a character code (`code_type/2`) or a character (`char_type/2`). Both accept the same classes: `alpha`, `alnum`, `digit`, `xdigit`, `space`, `white`, `layout`, `upper`, `lower`, `punct`, `csym`, `csymf`, `end_of_line`, `newline`, `end_of_file`, `graph`, `print`, `ascii`, `cntrl`, `meta`, `solo`, `symbol`, `period`, `quote`, `paren`, `prolog_var_start`, `prolog_atom_start`, `prolog_identifier_continue`, `prolog_symbol`. `space` is the C `isspace` set (`tab`, `newline`, `vertical tab`, `form feed`, `carriage return`, `space`, and the Unicode spaces); `white` is `space` and `tab`; `code_type(-1, end_of_file)` holds; an unknown class raises `domain_error(char_type, T)` (4.5.0).

Both are nondeterministic: with the character unbound they generate (over the ASCII range), with the type unbound they enumerate every class the character belongs to, and with both unbound they enumerate every pair.

The **parametric forms** - `digit(Weight)`, `upper(Lower)`, `lower(Upper)`, `to_lower(Lower)` and `to_upper(Upper)`, and `xdigit(Weight)` - work in every mode: bound they test, unbound they bind, and with the character unbound they generate. `char_type/2` gives a character where `code_type/2` gives a code; `digit(Weight)` gives an integer weight in both.

```
?- code_type(0'7, digit), char_type('A', upper), char_type(a, lower), char_type(' ', space).
true.

?- char_type('7', digit(W)).
W = 7.

?- char_type('A', upper(L)).
L = a.

?- char_type(a, to_upper(U)).
U = 'A'.

?- code_type(0'a, lower(U)).
U = 65.

?- char_type(X, to_upper('A')).
X = 'A' ;
X = a.
```

39. Debugging and profiling

profile/0, noprofile/0, profile_data/1, reset_profile/0

Purpose: Count predicate calls. `profile/0` starts counting, `noprofile/0` stops it, `profile_data(-Data)` unifies a list of Name/Arity-Count pairs sorted by count, and `reset_profile/0` clears the counters.

Only user-defined predicates are counted.

```
foo(1). foo(2).

?- reset_profile, profile, (foo(_), fail ; true), noprofile, profile_data(D).
D = [foo/1-1].
```

spy/1, nospy/1, spying/1, debugging/0, leash/1

Purpose: Spy points mark predicates whose ports are shown by the tracer. `spy(+PI)` and `nospy(+PI)` add/remove a spy point (Name/Arity), `spying(?PI)` enumerates them, `debugging/0` prints the trace state and spy points, and `leash(+Ports)` selects the ports (call, exit, redo, fail, all, none) at which an interactive tracer stops.

```
?- spy(foo/2), spying(P).
% Spy point set on foo/2
P = foo/2.

?- debugging.
Tracing is OFF
Spy points:
    foo/2
true.
```

trace/0, notrace/0

Purpose: Enable / disable four-port tracing (Call, Exit, Redo, Fail lines are written to the current output for every user predicate and traced built-in).

```
?- trace, member(X, [1, 2]), X > 1, notrace.
% Tracing enabled
Call: (0) member(X, [1, 2])
Exit: (0) member(1, [1, 2])
Redo: (0) member(2, [1, 2])
Exit: (0) member(2, [1, 2])
    Call: (1) notrace
% Tracing disabled
X = 2.
```

cut/0

Purpose: A callable synonym of `!` kept for programs that write the cut as a plain atom goal through `call/1`; inside a clause body always use `!`.

40. System and memory

atom_gc/0, atom_table_size/1

Purpose: `atom_gc/0` releases interned atoms that are no longer referenced; `atom_table_size(-N)` reports the number of interned atoms.

```
?- atom_gc, atom_table_size(N).
N = 75.
```

shell2/2, sleep/1

Purpose: `shell2(+Command, -ExitCode)` runs a shell command and returns its exit status (`shell/1, 2` are the other forms); `sleep(+Seconds)` pauses the current thread (fractions allowed).

```
?- shell2('ls /nonexistent', Code).
Code = 2.

?- sleep(0.5).
true.
```

working_directory/2, file_modified/2, delete_directory/1

Purpose: `working_directory(-Old, +New)` reads and changes the engine's working directory (`working_directory(D, D)` only reads it); `file_modified(+Path, -Millis)` gives the modification time in epoch milliseconds; `delete_directory(+Path)` removes an empty directory.

v4.6.0 (ISS-2025-0745): the working directory belongs to the **engine** (each Prolog instance has its own, starting at the JVM's), never to the JVM: the user `.dir` property is not touched. Every relative file name - `open/3, 4`, `consult/1` and the loaders, `absolute_file_name/2, 3`, `exists_file/1`, the filesystem, CSV, persistence and logging built-ins - is resolved against it. `Old` ends with `/` (SWI); `New` may be relative to the current directory; a `New` that is not a directory is `existence_error(directory, New)`.

```
?- working_directory(D, D).
D = '/home/user/project/'.

?- open('note.txt', write, W), close(W), file_modified('note.txt', T).
T = 1787644907965.
```

format_time/3, parse_time/3, stamp_date_time/3, date_time_stamp/2

Purpose: `format_time(+Pattern, +Timestamp, -Text)` formats a time stamp in SECONDS since the epoch (as `get_time/1` returns it; a float) or a datetime atom with a Java `DateTimeFormatter` pattern; `parse_time(+Pattern, +Text, -Stamp)` is the inverse. SWI-Prolog's argument order `format_time(+Out, +Format, +Stamp)` - `Out` is `atom(A)`, `string(S)`, `codes(C)`, `chars(C)` or a stream, `Format` uses % strftime directives (%Y-%m-%d %H:%M:%S, %a, %b, %F, %T, %Z, ...) - is accepted as well. `stamp_date_time(+Stamp, -date(Y, M, D, H, Mn, S, Off, TZ, DST), +TimeZone)` (`TimeZone` is `local`, `'UTC'` or an offset in seconds west of UTC) and `date_time_stamp(+date(...), -Stamp)` convert between stamps and dates (4.5.0).

```
?- parse_time('yyyy-MM-dd HH:mm:ss', '2026-03-20 14:30:00', S),
   format_time('dd/MM/yyyy', S, Day).
S = 1774013400.0, Day = '20/03/2026'.

?- stamp_date_time(0, D, 'UTC').
D = date(1970, 1, 1, 0, 0, 0.0, 0, 'UTC', -).
```

```
?- get_time(T), format_time(atom(A), '%Y-%m-%d', T).
```

The stamp of a local date and time depends on the local time zone.

41. JDBC database access

The JDBC predicates connect a Prolog program to any database with a JDBC driver on the class path. Connections, prepared statements and callable statements are represented by handle atoms (`'$jdbc_conn_1', ...`). All predicates raise `existence_error` for an unknown handle and a `jdbc_error(Message)` ball for SQL failures. They are removed by `enableSafeMode()`.

`jdbc_driver_load/1`

Purpose: Load a driver class by name (needed only by drivers that do not self-register).

```
?- jdbc_driver_load('org.h2.Driver').
true.
```

`jdbc_connect/2`, `jdbc_connect/4`, `jdbc_disconnect/1`

Purpose: Open a connection from a JDBC URL, optionally with user and password; close it with `jdbc_disconnect/1`.

```
?- jdbc_connect('jdbc:h2:mem:testdb', C).
C = '$jdbc_conn_1'.

?- jdbc_connect('jdbc:postgresql://localhost/mydb', 'user', 'pass', C), jdbc_disconnect(C).
```

`jdbc_query/3`, `jdbc_execute_update/3`

Purpose: `jdbc_query(+Conn, +SQL, -Rows)` runs a SELECT and returns a list of `row(Col1, ...)` terms; `jdbc_execute_update(+Conn, +SQL, -Count)` runs INSERT/UPDATE/DELETE/DDI and returns the number of affected rows.

```
?- jdbc_execute_update(C, 'CREATE TABLE users(id INT, name VARCHAR(50))', _),
   jdbc_execute_update(C, 'INSERT INTO users VALUES (1, 'Alice')', N),
   jdbc_query(C, 'SELECT id, name FROM users', Rows).
N = 1, Rows = [row(1, 'Alice')].
```

`jdbc_prepare/3`, `jdbc_set_param/3`, `jdbc_set_params/2`, `jdbc_execute_prepared_query/2`, `jdbc_execute_prepared_update/2`, `jdbc_close_statement/1`

Purpose: Prepared statements with `?` placeholders: prepare, bind parameters (1-based, or all at once from a list; use the atom `null` for SQL NULL), execute as a query or an update, close.

```
insert_user(C, Id, Name) :-
    jdbc_prepare(C, 'INSERT INTO users(id, name) VALUES (?, ?)', S),
    jdbc_set_params(S, [Id, Name]),
    jdbc_execute_prepared_update(S, 1),
    jdbc_close_statement(S).

older_than(C, Age, Rows) :-
    jdbc_prepare(C, 'SELECT name FROM users WHERE age > ?', S),
    jdbc_set_param(S, 1, Age),
    jdbc_execute_prepared_query(S, Rows),
    jdbc_close_statement(S).
```

jdbc_set_autocommit/2, jdbc_commit/1, jdbc_rollback/1

Purpose: Transaction control. Disable autocommit, run statements, then commit or roll back.

```
transfer(C, From, To, Amount) :-
    jdbc_set_autocommit(C, false),
    catch(( debit(C, From, Amount), credit(C, To, Amount), jdbc_commit(C) ),
        E,
        ( jdbc_rollback(C), throw(E) )),
    jdbc_set_autocommit(C, true).
```

jdbc_tables/2, jdbc_columns/3

Purpose: Metadata: the list of table names, and the columns of a table as column(Name, TypeName, Size) terms.

```
?- jdbc_tables(C, Ts), jdbc_columns(C, 'USERS', Cols).
Ts = ['USERS'], Cols = [column('ID', 'INTEGER', 10), column('NAME', 'VARCHAR', 50)].
```

**jdbc_prepare_call/3, jdbc_call_set_param/3, jdbc_call_register_out/3,
jdbc_call_execute/1, jdbc_call_get_result/3, jdbc_call_get_resultset/2**

Purpose: Stored procedures through callable statements: prepare {call proc(?, ?)}, set IN parameters, register OUT parameters with a type (integer, bigint, double, decimal, varchar, boolean, date, timestamp), execute, then read OUT values or the returned result set as row/N terms.

```
user_count(C, Table, Count) :-
    jdbc_prepare_call(C, '{call get_user_count(?, ?)}', S),
    jdbc_call_set_param(S, 1, Table),
    jdbc_call_register_out(S, 2, integer),
    jdbc_call_execute(S),
    jdbc_call_get_result(S, 2, Count),
    jdbc_close_statement(S).
```

**jdbc_set_clob/3, jdbc_get_clob/3, jdbc_set_blob/3, jdbc_set_blob_bytes/3,
jdbc_get_blob_bytes/3, jdbc_get_blob_to_file/3**

Purpose: Large objects: set a CLOB parameter from text, a BLOB parameter from a file path or a byte list; read a CLOB as an atom, a BLOB as a byte list, or save a BLOB straight to a file. The get variants take a SELECT returning a single LOB column.

```
store_doc(C, Id, Text) :-
    jdbc_prepare(C, 'INSERT INTO docs(id, body) VALUES (?, ?)', S),
    jdbc_set_param(S, 1, Id), jdbc_set_clob(S, 2, Text),
    jdbc_execute_prepared_update(S, _), jdbc_close_statement(S).

?- jdbc_get_clob(C, 'SELECT body FROM docs WHERE id = 1', Text).
Text = 'A long document ...'.

?- jdbc_get_blob_bytes(C, 'SELECT data FROM images WHERE id = 1', Bytes).
Bytes = [137, 80, 78, 71|...].
```

42. TCP, UDP and DNS

Socket predicates exchange UTF-8 text; handles are atoms such as '\$socket_1'. `tcp_accept/2`, `tcp_receive/3` and `udp_receive/4` block until data arrives; a closed peer yields the atom `end_of_stream`. These predicates are removed by `enableSafeMode()`.

tcp_connect/3, tcp_send/2, tcp_receive/3, tcp_close/1

Purpose: TCP client: connect to `Host:Port`, send a string, receive up to `MaxBytes` bytes, close.

```
echo_client(Host, Port, Msg, Reply) :-
    tcp_connect(Host, Port, S),
    tcp_send(S, Msg),
    tcp_receive(S, Reply, 4096),
    tcp_close(S).

?- echo_client(localhost, 7777, 'ping', R).
R = pong.
```

tcp_server_socket/2, tcp_accept/2

Purpose: TCP server: bind a listening socket to a port, then accept client connections one at a time.

```
serve_once(Port) :-
    tcp_server_socket(Port, SS),
    tcp_accept(SS, Client),
    tcp_receive(Client, Data, 4096),
    tcp_send(Client, Data),           % echo
    tcp_close(Client),
    tcp_close(SS).
```

udp_socket/2, udp_send/4, udp_receive/4, udp_close/1

Purpose: UDP datagrams: `udp_socket(+Port, -Sock)` binds a socket (port 0 = any free port), `udp_send(+Sock, +Host, +Port, +Data)`, `udp_receive(+Sock, -Data, -From, +MaxBytes)` with `From = from(IP, Port)`.

```
udp_ping(Host, Port, Reply) :-
    udp_socket(0, S),
    udp_send(S, Host, Port, ping),
    udp_receive(S, Reply, from(_, _), 1024),
    udp_close(S).
```

hostname_address/2, hostname/1

Purpose: `hostname_address(+Name, -IP)` resolves a host name; `hostname(-Name)` returns the local host name.

```
?- hostname_address(localhost, IP).
IP = '127.0.0.1'.
```

http_request/4, http_post/4

Purpose: Minimal HTTP client: `http_request(+Method, +URL, -Status, -Body)` with `Method` one of `get`, `head`, `delete`, `put`, `post`; `http_post(+URL, +RequestBody, -Status, -Body)` sends a form-encoded body. Both use 30-second timeouts. The richer `http_get/3`, `http_open/3` and the JSON-aware `http_client_*` predicates are described in Chapter 26.

```
?- http_request(get, 'http://httpbin.org/get', Status, Body), Status == 200.  
Status = 200, Body = '{"args":{}, ...}'.
```

43. DCG support predicates

dcg_translate_rule/2, dcg_body//2

Purpose: `dcg_translate_rule(+Rule, -Clause)` returns the ordinary clause `a --> rule` is translated to, exposing the translator used by `consult/1` and `assertz/1`. `dcg_body//2` is a reserved grammar non-terminal (two arguments plus the difference list) whose current implementation simply succeeds; it is kept for source compatibility.

```
?- dcg_translate_rule((a --> [x], b), C).  
C = (a(S0, S) :- S0 = [x|S1], b(S1, S)).
```

Appendix A - Prolog flags

`current_prolog_flag(?Flag, ?Value)` reads and `set_prolog_flag(+Flag, +Value)` sets the flags below. Flags marked *read-only* raise `permission_error(modify, flag, Flag)` when set; a flag that is not one of these raises `domain_error(prolog_flag, Flag)` on either predicate, and a value the flag does not accept raises `domain_error(flag_value, Flag+Value)` (ISO 8.17.1.3 / 8.17.2.3, exact since 4.4.0 - ISS-2025-0508).

Flag	Values	Notes
<code>bounded</code>	<code>false</code>	read-only; integers are arbitrary precision (X is 10^{30} is exact)
<code>max_integer, min_integer</code>	<code>integer</code>	read-only; the limits of the fast 64-bit representation an integer uses before it is promoted to a big integer - not a limit on arithmetic. ISO does not require them when <code>bounded</code> is <code>false</code> ; they are kept because programs read them, as SWI-Prolog does.
<code>integer_rounding_function</code>	<code>toward_zero</code>	read-only
<code>max_arity</code>	<code>unbounded</code>	read-only
<code>double_quotes</code>	<code>codes, chars, atom, string</code>	how "text" is read (default string)
<code>unknown</code>	<code>error, fail, warning</code>	what a call to an undefined procedure does (default error)
<code>occurs_check</code>	<code>true, false</code>	occurs check in ordinary unification (default false)
<code>char_conversion</code>	<code>true, false</code>	apply the <code>char_conversion/2</code> table when reading
<code>debug</code>	<code>on, off</code>	debug mode
<code>syntax_errors</code>	<code>error, fail, quiet</code>	behaviour of <code>read/1</code> on a syntax error
<code>character_escapes</code>	<code>true, false</code>	escape processing in quoted atoms
<code>strict_iso</code>	<code>true, false</code>	reject non-ISO extensions in <code>read/1</code>
<code>dialect, prolog_version, version, version_data</code>	<code>atom / term</code>	read-only identification (jprolog)
<code>encoding</code>	<code>atom</code>	default text encoding (utf8)
<code>argv</code>	<code>list</code>	command-line arguments
<code>gc, optimize, traditional, write_strings, trace, initialization, stack_limit, toplevel_print_options</code>	see <code>current_prolog_flag/2</code>	SWI-compatible flags kept for portability

A flag takes effect for the text read after it is set (a query is parsed as a whole before it runs):

```
?- set_prolog_flag(double_quotes, codes).
true.

?- X = "ab".
X = [97, 98].

?- set_prolog_flag(unknown, fail), no_such_pred.
false.
```

Appendix B - Engine configuration and embedding

Subsystem selection (system properties). The clean-room implementations are the defaults. There is **one resolution engine** (`core.engine.v4`, since 4.0.0); the parser, the CLP(FD) solver and the DCG translator each keep an older implementation that a property set before the engine starts can select:

Property	Selects
<code>-Djprolog.parser=legacy</code>	the recursive-descent parser instead of the operator-precedence parser
<code>-Djprolog.clpfd=legacy</code>	the first CLP(FD) implementation (read once, when the engine is created)
<code>-Djprolog.dcg=legacy</code>	the first DCG translator

Only the literal value `legacy` selects a fallback. `-Djprolog.engine` no longer selects anything: the recursive engine it could name as `legacy` was deleted in 4.0.0 and the iterative 3.x machine it could name as `v2` in **4.1.0**, so any value logs a warning and runs the `v4` core. The parser and DCG switches are also available at runtime as `Prolog.setUseV2Parser/Dcg(boolean)`; the CLP(FD) choice must be made before constructing the engine (or upgraded with `prolog.enableV2Clpfd()`).

Java API summary (`it.denzosoft.jprolog.core.engine.Prolog`):

Method	Purpose
<code>consult(String program)</code>	load clauses and run directives (throws on the first error)
<code>consultWithDiagnostics(program, fileName)</code>	per-clause compilation returning every error with its line
<code>solve(String query)</code>	all solutions as <code>List<Map<String, Term>></code>
<code>solveStream(query, sink)</code>	lazy enumeration; the sink returns <code>false</code> to stop
<code>solveStream(query, AnswerSink)</code>	the same, also telling whether more answers may follow (what a toplevel needs to print <code>.</code> or <code>;</code>)
<code>compileFile("p.pl"), consultCompiled("p.jpc"), consultSmart("p.pl")</code>	binary <code>.jpc</code> compilation and cached loading
<code>asserta/assertz/retract(String clause)</code>	modify the database from Java
<code>nbSetval/nbGetval/nbDelete</code>	global variables from Java
<code>enableSafeMode()</code>	sandbox: remove every OS, FFI, filesystem, network, HTTP, JDBC, persistence and threading built-in, <code>open/3,4</code> , the loaders, the CSV file predicates and <code>log_to_file/1</code> - from the legacy registry and the native table (irreversible for the instance)
<code>enableSafeMode(SafeModeOptions)</code>	the same; new <code>SafeModeOptions().allowFileRead(dir)</code> keeps read-only <code>open/3,4</code> and the loaders for files inside <code>dir</code>
<code>setInferenceBudget(long steps)</code>	abort a query with <code>InferenceLimitException</code> after the given number of resolution steps (0 = unlimited); ONE budget shared by the query's meta-calls and worker threads, and natives that walk or build long lists charge per element; the bridged libraries charge per 64 characters of text input and per solution, and a regular expression per 256 characters it reads (4.6)
<code>getFlags(), setTracing(boolean), isTracing()</code>	per-engine ISO flag store and <code>trace/0</code> state (each Prolog instance is isolated)
<code>getPredicateIndicatorAtLine(int line)</code>	map a source line to its clause (IDE breakpoints)
<code>getEngineContext()</code>	the durable per-engine context: <code>setDebugController/getDebugController</code> (IDE debugger wiring) and the running query's <code>ResourceGuard</code>
<code>getEngineState(), getStreams(), getOps()</code>	the engine's own stream table and operator store
<code>residualGoals(solution)</code>	the constraints still attached to an answer's variables

Resource control. A runaway query stops with `InferenceLimitException` when the budget is

exceeded and with `QueryCancelledException` when the solver thread is interrupted (the IDE Stop button). Both are plain Java runtime exceptions that a Prolog `catch/3` cannot intercept, so untrusted programs cannot defeat them. Deep terms or deeply nested input are reported as the catchable ISO errors `resource_error(stack_overflow)` and `resource_error(parser_nesting)`.

Appendix C - ISO error terms

Built-in predicates signal errors by throwing `error(Formal, Context)`. `Context` is `context(Name/Arity, Message)` or a variable. The `Formal` terms used by JProlog:

Formal term	Raised when
<code>instantiation_error</code>	an argument that must be bound is a variable
<code>type_error(Type, Culprit)</code>	an argument has the wrong type (integer, atom, callable, list, evaluable, ...)
<code>domain_error(Domain, Culprit)</code>	the type is right but the value is outside the domain (<code>not_less_than_zero</code> , <code>operator_priority</code> , ...)
<code>existence_error(Kind, What)</code>	a procedure, stream, file, variable or source does not exist
<code>permission_error(Action, Type, Culprit)</code>	e.g. modifying a static procedure, repositioning a stream, accessing a private procedure
<code>representation_error(Flag)</code>	a value cannot be represented (<code>character_code</code> , <code>max_arity</code> , <code>cyclic_term</code>)
<code>evaluation_error(Error)</code>	arithmetic: <code>zero_divisor</code> , <code>undefined</code> , <code>float_overflow</code> , <code>int_overflow</code> , <code>negative_shift</code>
<code>resource_error(Resource)</code>	<code>memory</code> , <code>stack_overflow</code> , <code>parser_nesting</code>
<code>syntax_error(Message)</code>	invalid text in <code>read/1</code> , <code>atom_to_term/3</code> , <code>number_codes/2</code> , ...
<code>system_error</code>	an internal failure inside a built-in

JProlog follows the strict ISO reading: `atom_length(123456, L)` raises `type_error(atom, 123456)` instead of counting the digits as SWI-Prolog does.

```
?- catch(atom_length(123456, L), error(E, _), true).
E = type_error(atom, 123456).

?- catch(atom_length(X, L), error(Err, _), true).
Err = instantiation_error.

?- catch(foo(1), error(existence_error(procedure, PI), _), true).
PI = foo/1.
```

Appendix D - Standard order of terms and text conversion summary

Standard order (used by @<, compare/3, sort/2, msort/2, setof/3, keysort/2): Var < Number < Atom < String < Compound. Numbers are compared by value, and when an integer and a float are equal the float comes first. Compound terms are ordered by arity, then by functor name, then by arguments left to right.

Text conversions at a glance:

From \ To	atom	string	codes	chars	number
atom	-	atom_string/2	atom_codes/2	atom_chars/2	atom_number/2
string	string_to_atom/2	-	string_codes/2	string_chars/2	number_string/2
number	atom_number/2	number_string/2	number_codes/2	number_chars/2	-
any term	term_to_atom/2	with_output_to/2	-	-	-

Appendix E - Libraries not loaded by default

The source tree also documents experimental libraries that are **not** registered in the default engine and are therefore outside the scope of this manual: expert system (`es_*`), NLP, inference engine, AI planner, fuzzy logic, Bayesian networks, genetic algorithms, neural networks, optimisation, simulation, workflow, CLP(R) (`clpr_*`), knowledge graphs, parsing/DSL, Datalog, RDF, model checking, CHR (`chr_*`), BDI agents (`agent_*`), ASP (`asp_*`), explainable AI, type inference, theorem proving, symbolic mathematics, meta-interpretation, event calculus (`ec_*`), probabilistic logic, SAT solving, game playing, term rewriting and description logic (`dl_*`). See `docs/references/BUILTIN_PREDICATES_REFERENCE.md` sections 29-60 for their descriptions and `docs/guides/guide-extension.md` for how to register an extension with `Prolog.registerBuiltInPredicate(name, builtIn)`.

Index of predicates and operators

Every name/arity heading of this manual with the page where it is described.

abolish/1	72	close/2	163	db_transaction/1	148
abolish_all_tables/0	141	code_type/2	174	dcg_translate_rule/2	89
abolish_table/1	142	compare/3	25	debugging/0	175
absolute_file_name/2	111	compound/1	17	del_attr/2	167
absolute_file_name/3	111	concurrent/3	155	delete/3	35
acyclic_term/1	19	concurrent_and/2	155	delete_directory/1	176
aggregate/3	55	concurrent_forall/2	156	delete_file/1	109
aggregate/4	55	concurrent_forall/3	156	dif/2	167
aggregate_all/3	54	concurrent_maplist/2	155	directory_exists/1	109
aggregate_all/4	55	concurrent_maplist/3	155	directory_files/2	111
all_different/1	133	concurrent_maplist/4	155	discontiguous/1	9
all_distinct/1	133	concurrent_or/2	156	disjoint2/1	135
append/2	35	consult/1	97	distinct/1	55
append/3	28	copy_file/2	110	distinct/2	55
arg/3	21	copy_term/2	23	div/2	11
assert/1	70	copy_term/3	168	downcase_atom/2	84
asserta/1	70	cpu_count/1	114	dynamic/1	9
assertz/1	70	crypto_aes_decrypt/4	102	element/3	134
at_end_of_stream/0	63	crypto_aes_encrypt/4	102	enhanced_phrase/2	87
atom/1	15	crypto_hash/3	100	enhanced_phrase/3	87
atom_chars/2	77	crypto_hash_password/2	102	ensure_loaded/1	98
atom_codes/2	77	crypto_random_int/3	102	erase/1	74
atom_concat/3	75	crypto_verify_password/2	103	exclude/3	37
atom_gc/0	176	csv_parse/2	126	exists_directory/1	109
atom_length/2	75	csv_read_file/2	126	exists_file/1	109
atom_number/2	78	csv_serialize/2	126	expand_term/2	90
atom_string/2	161	csv_write_file/2	126	fd_degree/2	136
atom_table_size/1	176	cumulative/1	134	fd_dom/2	137
atom_to_number/2	81	cumulative/2	134	fd_inf/2	138
atom_to_term/3	27	current_char_conversion/2	67	fd_size/2	138
atomic/1	16	current_input/1	163	fd_sup/2	138
atomic_list_concat/2	80	current_key/1	74	fd_var/1	138
atomic_list_concat/3	80	current_module/1	158	file_base_name/2	112
attvar/1	167	current_op/3	173	file_directory_name/2	112
automaton/3	135	current_output/1	163	file_exists/1	109
automaton/8	135	current_predicate/1	73	file_extension/2	111
b_getval/2	166	current_prolog_flag/2	94	file_modified/2	176
b_setval/2	166	current_stream/3	67	file_search_path/2	99
bagof/3	52	current_table/2	142	file_size/2	110
base64_decode/2	101	current_thread_pool/1	125	findall/3	51
base64_encode/2	101	cut/0	175	findall/4	52
between/3	42	cyclic_term/1	19	first_solution/3	155
call/1	50	date_add/4	106	flag/3	74
call_cleanup/2	49	date_diff/4	107	flatten/2	34
call_dcg/3	89	date_parts/4	107	float/1	16
call_delays/2	141	date_time_stamp/2	176	flush_output/0	163
call_nth/2	55	day_of_week/2	107	flush_output/1	163
callable/1	18	db_batch_assert/1	148	foldl/4	37
catch/3	91	db_clear/0	148	foldl/5	37
chain/2	135	db_export_json/1	147	foldl/6	37
char_code/2	82	db_import_json/1	147	foldl/7	37
char_conversion/2	67	db_load/1	146	forall/2	47
char_type/2	83	db_restore/1	147	format/2	59
character_count/2	66	db_save/1	146	format/3	59
circuit/1	134	db_save_predicate/2	146	format_time/3	176
clause/2	173	db_snapshot/1	147	free_memory/1	114
close/1	67	db_sync/0	148	freeze/2	167

frozen/2	168	java_set_field/3	153	maplist/2	37
functor/3	21	java_to_term/2	154	maplist/3	37
garbage_collect/0	94	jdbc_call_execute/1	179	maplist/4	37
garbage_collect_atoms/0	94	jdbc_call_get_result/3	179	maplist/5	37
get_attr/3	167	jdbc_call_get_resultset/2	179	max_list/2	35
get_byte/1	63	jdbc_call_register_out/3	179	max_member/2	36
get_char/1	62	jdbc_call_set_param/3	179	max_member/3	36
get_code/1	63	jdbc_close_statement/1	178	md5_hash/2	100
get_time/1	106	jdbc_columns/3	179	member/2	29
getenv/2	113	jdbc_commit/1	179	message_queue_create/1	121
global_cardinality/2	134	jdbc_connect/2	178	message_queue_create/2	123
global_cardinality/3	136	jdbc_connect/4	178	message_queue_destroy/1	123
graph_components/2	150	jdbc_disconnect/1	178	message_queue_property/2	124
graph_connected/2	149	jdbc_driver_load/1	178	meta_predicate/1	9
graph_degree/3	150	jdbc_execute_prepared_qu..	178	min_list/2	35
graph_edges/2	149	jdbc_execute_prepared_up..	178	min_member/2	36
graph_has_cycle/1	151	jdbc_execute_update/3	178	min_member/3	36
graph_neighbors/3	150	jdbc_get_blob_bytes/3	179	minimum_spanning_tree/2	150
graph_path/4	149	jdbc_get_blob_to_file/3	179	mod/2	11
graph_reachable/3	151	jdbc_get_clob/3	179	module_transparent/1	9
graph_scc/2	151	jdbc_prepare/3	178	msort/2	33
graph_vertices/2	149	jdbc_prepare_call/3	179	multifile/1	9
ground/1	18	jdbc_query/3	178	must_be/2	172
halt/0	92	jdbc_rollback/1	179	mutex_create/1	123
halt/1	92	jdbc_set_autocommit/2	179	mutex_destroy/1	123
hmac/4	101	jdbc_set_blob/3	179	mutex_lock/1	123
hostname/1	115	jdbc_set_blob_bytes/3	179	mutex_property/2	124
hostname_address/2	180	jdbc_set_clob/3	179	mutex_trylock/1	123
http_client_get/2	144	jdbc_set_param/3	178	mutex_unlock/1	123
http_client_post/3	144	jdbc_set_params/2	178	mutex_unlock_all/0	123
http_get_request/2	143	jdbc_tables/2	179	nb_current/2	166
http_handler/3	143	json_get/3	104	nb_delete/1	166
http_open/3	144	json_keys/2	105	nb_getval/2	166
http_post/4	180	json_member/3	105	nb_setarg/3	171
http_reply/4	144	json_parse/2	104	nb_setval/2	166
http_reply_json/3	144	json_serialize/2	104	nextto/3	36
http_request/4	180	json_set/4	105	nl/0	57
http_server/2	143	keysort/2	34	nonvar/1	14
http_stop/1	143	label/1	136	noprofile/0	175
ignore/1	48	labeling/2	136	nospy/1	175
in/2	12	last/2	34	notrace/0	175
include/1	99	leash/1	97	now/1	106
include/3	37	length/2	29	nth0/3	31
indomain/1	137	lex_chain/1	135	nth1/3	31
ins/2	12	limit/2	55	number/1	15
integer/1	15	line_count/2	66	number_chars/2	78
intersection/3	35	line_position/2	66	number_codes/2	78
is/2	10	list_to_ord_set/2	36	number_string/2	161
is_list/1	18	list_to_set/2	36	number_to_atom/2	81
is_ordset/1	36	listing/0	73	number_vars/3	171
java_array_get/3	153	listing/1	73	numbervars/3	27
java_array_length/2	154	load_files/1	98	numlist/3	34
java_array_new/3	153	load_files/2	98	offset/2	55
java_array_set/3	153	log_debug/1	129	once/1	46
java_call/4	152	log_error/1	128	op/3	173
java_class/2	153	log_info/1	128	open/3	67
java_from_term/2	154	log_level/1	129	open_string/2	62
java_gc/0	154	log_to_file/1	129	ord_add_element/3	36
java_get_field/3	152	log_warning/1	128	ord_del_element/3	36
java_instanceof/2	153	make/0	98	ord_disjoint/2	36
java_new/3	152	make_directory/1	110	ord_empty/1	36
java_release_ref/1	154	make_directory_path/1	111	ord_insert/3	36

ord_intersect/2	36	read_term_from_atom/3	61	sum_list/2	34
ord_intersection/2	36	recorda/2	74	sumlist/2	34
ord_intersection/3	36	recorda/3	74	system_time/1	113
ord_memberchk/2	36	recorded/2	74	tab/1	68
ord_selectchk/3	36	recorded/3	74	table/1	9
ord_seteq/2	36	recordz/2	74	tcp_accept/2	180
ord_subset/2	36	recordz/3	74	tcp_close/1	180
ord_subtract/3	36	rem/2	11	tcp_connect/3	180
ord_symdiff/3	36	rename_file/2	109	tcp_receive/3	180
ord_union/2	36	repeat/0	47	tcp_send/2	180
ord_union/3	36	reset_profile/0	175	tcp_server_socket/2	180
order_by/2	55	retract/1	71	term_attvars/2	168
os_name/1	114	retractall/1	71	term_string/2	27
pairs_keys/2	170	reverse/2	30	term_to_atom/2	26
pairs_keys_values/3	170	scalar_product/4	133	term_variables/2	25
pairs_values/2	170	seek/4	67	thread_create/2	119
parse_time/3	176	select/3	32	thread_create/3	119
partial_list/1	172	set_input/1	68	thread_create_in_pool/4	125
partition/4	37	set_output/1	68	thread_detach/1	120
partition/5	37	set_prolog_flag/2	95	thread_exit/1	123
peek_byte/1	63	set_stream/2	66	thread_get_message/1	121
peek_char/1	63	set_stream_position/2	67	thread_get_message/2	122
peek_code/1	63	setarg/3	171	thread_get_message/3	122
permutation/2	31	setenv/2	113	thread_is_alive/1	121
persist/1	146	setof/3	53	thread_join/1	123
phrase/2	86	setup_call_cleanup/3	49	thread_join/2	120
phrase/3	86	sha256_hash/2	100	thread_peek_message/2	122
phrase_with_options/4	88	sha512_hash/2	100	thread_pool_create/3	125
pid/1	114	shell/1	115	thread_pool_destroy/1	125
plus/3	43	shell/2	115	thread_pool_property/2	125
portray_clause/1	65	shell2/2	176	thread_property/2	123
predicate_property/2	158	shell_output/3	115	thread_self/1	120
predsort/3	169	shortest_path/4	149	thread_send_message/2	122
print/1	58	simple/1	172	thread_send_message/3	124
print_message/2	65	sleep/1	176	thread_signal/2	124
profile/0	175	sort/2	33	thread_sleep/1	121
profile_data/1	175	sort/4	33	thread_statistics/3	124
prolog_load_context/2	99	source_file/1	99	throw/1	91
proper_length/2	36	source_file/2	99	time_parts/4	107
proper_list/1	19	split_string/4	79	tnot/1	141
put_attr/3	167	spy/1	175	to_codes/2	162
put_byte/1	63	spying/1	175	today/1	106
put_char/1	62	stamp_date_time/3	176	topological_sort/2	150
put_code/1	171	statistics/2	96	total_memory/1	114
random_token/2	102	stream_position/2	67	trace/0	175
rational/1	16	stream_position_data/3	66	transpose/2	136
rational/3	16	stream_property/2	66	trim_stacks/0	94
rdiv/2	11	string/1	17	tuples_in/2	134
re_escape/2	117	string_chars/2	160	udp_close/1	180
re_findall/3	117	string_code/3	81	udp_receive/4	180
re_match/2	116	string_codes/2	160	udp_send/4	180
re_matchsub/3	116	string_concat/3	160	udp_socket/2	180
re_replace/4	116	string_length/2	160	undefined/0	141
re_split/3	117	string_lower/2	81	unifiable/3	168
read/1	60	string_to_atom/2	81	unify_with_occurs_check/2	171
read/2	60	string_upper/2	81	union/3	35
read_file_to_atom/2	110	sub_atom/5	76	unpersist/1	147
read_line_to_codes/2	62	sub_string/5	161	unsetenv/1	113
read_line_to_string/2	62	subsumes_term/2	26	upcase_atom/2	84
read_string/3	62	subtract/3	35	url_decode/2	145
read_string/5	62	succ/2	42	url_encode/2	144
read_term/2	61	sum/3	133	use_module/1	98

use_module/2	98	\+/1	9
uuid/1	101	\//2	11
var/1	14	\//1	11
when/2	167	\=/2	10
with_input_from/2	62	\==/2	10
with_mutex/2	123	^/2	11
with_output_to/2	68		
working_directory/2	176		
write/1	57		
write_atom_to_file/2	110		
write_canonical/1	64		
write_term/2	64		
writeln/1	57		
writeq/1	58		
xml_parse/2	118		
xml_serialize/2	118		
xor/2	11		
xpath/3	118		
zcompare/3	135		
!/0	10		
#/\//2	133		
#</2	12		
#<==/2	133		
#<==>/2	133		
#=/2	12		
#=</2	12		
#==>/2	133		
#>/2	12		
#>=/2	12		
#\//2	133		
#\//1	133		
#\//2	133		
#\=/2	12		
**/2	11		
*->/2	9		
*/2	11		
+/1	11		
+/2	11		
,/2	9		
-->/2	8		
-/1	8		
-/2	8		
->/2	9		
../2	12		
///2	11		
//1	130		
//2	11		
/\//2	11		
</2	10		
<</2	11		
=../2	10		
=/2	10		
=</2	10		
==/2	10		
=\=/2	10		
>/2	10		
>=/2	10		
>>/2	11		
@</2	10		
@=</2	10		
@>/2	10		
@>=/2	10		